


# Introduction to Machine Learning. Basic Concepts and Learning Paradigms.

Radu Ionescu, Prof. PhD.  
[raducu.ionescu@gmail.com](mailto:raducu.ionescu@gmail.com)

Faculty of Mathematics and Computer Science  
University of Bucharest

# Machine Learning



# Instructors

- Lectures:

➤ Radu Ionescu ([raducu.ionescu@gmail.com](mailto:raducu.ionescu@gmail.com))

- Labs:

➤ Antonio Bărbălău ([antoniobarbalau@gmail.com](mailto:antoniobarbalau@gmail.com))

➤ Mihail Burduja ([warchildmd@gmail.com](mailto:warchildmd@gmail.com))

➤ Marius Todea ([todeamarius2006@gmail.com](mailto:todeamarius2006@gmail.com))

# Grading System

- Your final grade is composed of:
  - 60% for Project 1
  - 40% for Project 2
- Both projects are individual!
- Each project consists of employing machine learning methods on a specific data set
- Project 1 is about participating in a Kaggle competition
- The competition will be launched in a couple of weeks
- Project 2 is about comparing two unsupervised approaches
- There are many datasets out there, so no overlap allowed among students!
- Methods and data sets must be chosen beforehand!

# Grading System

- Project 1 must be presented no later than week 10
- Project 2 must be presented no later than the day of the “exam”
- **There will be no paper exam!**
- **The average grade of projects 1 and 2 must be  $\geq 5$**
- The project consists of the code implementation in Python (any library is allowed) and a PDF report including:
  - a description of the data set (for project 2 only)
  - a description of the implemented machine learning methods
  - figures and/or tables with results
  - comments on the results
  - conclusion

# Grading System

- First project consists of implementing some machine learning method(s) for the proposed Kaggle challenge (TBA)
- The grades will be proportional to your model's accuracy:
  - Top 1-20 => your grade can be up to 10
  - Top 21-50 => your grade can be up to 9
  - Top 51-80 => your grade can be up to 8
  - Top 81-100 => your grade can be up to 7
  - Top 101-120 => your grade can be up to 6
  - Others => your grade can be up to 5
- Submit to: [practical.ml.fmi@gmail.com](mailto:practical.ml.fmi@gmail.com)

# Grading System


- Extra points during lectures / labs
- First to answer a question / solve an exercise gets 0.2 points
- Maximum 0.4 points per lecture / lab for each student
- Up to 1-2 point in total (added to final grade)

# (NO) Collaboration Policy

- Collaboration
  - Each student must write their own code for the project(s)
  - Borrowing code from web sources with copy & paste is not permitted
- No tolerance on plagiarism
  - Neither ethical nor in your best interest
  - **Don't cheat. We will find out (code will be checked!)**

# What is artificial intelligence (AI)?

- The ultimate goal of artificial intelligence is to build systems able to reach human intelligence levels
- Turing test: a computer is said to possess human-level intelligence if a remote human interrogator, within a fixed time frame, cannot distinguish between the computer and a human subject based on their replies to various questions posed by the interrogator



# Perhaps we are going in the right direction?



Alan Turing

1950: Can a computer convince a human that it is not a computer but a real person.

Now: Can a human convince a computer that he is a real person, not a computer


I'm not a robot

reCAPTCHA  
Privacy - Terms


# What is machine learning (ML)?

- Many AI researchers consider the ultimate goal of AI can be achieved by imitating the way humans learn
- **Machine Learning** – is the scientific study of algorithms and statistical models that computer systems use to learn from observations, without being explicitly programmed
- In this context, **learning** refers to:
  - recognizing complex patterns in data
  - making intelligent decisions based on data observations

# Classic Programming



# Machine Learning



# A well-posed machine learning problem

- What problems can be solved\* with machine learning?
- **Well-posed machine learning problem:**

"A computer program is said to learn from experience **E** with respect to some class of tasks **T** and performance measure **P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**." – Tom Mitchell

**(\*) implies a certain degree of accuracy**

# A well-posed machine learning problem

- Arthur Samuel (1959) wrote a program for playing checkers (perhaps the first program based on the concept of learning, as defined by Tom Mitchell)
- The program played 10K games against itself
- The program was designed to find the good and bad positions on the board from the current state, based on the probability of winning or losing
- In this example:
  - E = 10000 games
  - T = play checkers
  - P = win or lose



# Strong AI versus Weak AI


- Strong / generic / true AI  
(see the Turing test and its extensions)
- Weak / narrow AI  
(focuses on a specific well-posed problem)

# When do we use machine learning?

- We use ML when it is hard (impossible) to define a set of rules by hand / to write a program based on explicit rules
- Examples of tasks that be solved through machine learning:
  - face detection
  - speech recognition
  - stock price prediction
  - object recognition

# The essence of machine learning

- A pattern exists
- We cannot express it programmatically
- We have data on it



# What is machine learning?

[Arthur Samuel, 1959] field of study that:

- gives computers the ability to learn without being explicitly programmed

[Kevin Murphy] algorithms that:

- automatically detect patterns in data
- use the uncovered patterns to predict future data or other outcomes of interest

[Tom Mitchell] algorithms that:

- improve their performance (P)
- at some task (T)
- with experience (E)

# Brief history of AI



A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence.

(John McCarthy)



# Brief history of AI


- “We propose that a 2 month, 10 man study of artificial intelligence be carried out during the summer of 1956 at Dartmouth College in Hanover, New Hampshire.”
- The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it.
- An attempt will be made to find how to make machines use language, form abstractions and concepts, solve kinds of problems now reserved for humans, and improve themselves.
- We think that a significant advance can be made in one or more of these problems if a carefully selected group of scientists work on it together for a summer.”

# Brief history of AI

- 1960-1980s: "AI Winter"
- 1990s: Neural networks dominate, essentially because of the discovery of the backpropagation for training neural networks with two or more layers
- 2000s: Kernel methods dominate, essentially because of the instability of training neural networks
- 2010s: The comeback of neural networks, essentially because of the discovery of deep learning

# Why are things working today?

- More compute power
- More data
- Better algorithms / models



# ML in a nutshell

- Tens of thousands of machine learning algorithms
  - Researchers publish hundreds new every year
- Decades of ML research oversimplified:
  - Learn a mapping  $f$  from the input  $X$  to the output  $Y$ ,  
i.e.:
$$f: X \rightarrow Y$$
  - Example:  $X$  are emails,  $Y$ : {spam, not-spam}


# ML in a nutshell

- Input:  $X$  (images, texts, emails...)
- Output:  $Y$  (spam or not-spam...)
- (Unknown) Target Function:  
 $f: X \rightarrow Y$  (the “true” mapping / reality)
- Data  
 $(x_1, y_1), (x_2, y_2), \dots (x_N, y_N)$
- Model / Hypothesis Class  
 $g: X \rightarrow Y$   
 $y = g(x) = sign(w^T x)$

# ML in a nutshell

- Every machine learning algorithm has three components:
  - Representation / Model Class
  - Evaluation / Objective Function
  - Optimization

# Where does ML fit in?



# Learning paradigms

- Standard learning paradigms:

- Supervised learning
- Unsupervised learning
- Semi-supervised learning
- Reinforcement learning

- Non-standard paradigms:

- Active learning
- Transfer learning
- Transductive learning

# Supervised learning

- We have a set of labeled training samples
- Example 1: object recognition in images annotated with corresponding class labels



Car



Car



Person




Person



Dog

# Supervised learning

- Example 2: handwritten digit recognition (on the MNIST data set)




- Images of  $28 \times 28$  pixels
- We can represent each image as a vector  $x$  of 784 components
- We train a classifier  $f(x)$  such that:

$$f : x \rightarrow \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$$

# Supervised learning



- Example 2 (continued): handwritten digit recognition (on the MNIST data set)



- Starting with a training set of about 60K images (about 6000 images per class)
- ... the error rate can go down to 0.23% (using convolutional neural networks)
- Among the first (learning-based) systems used in a large-scale commercial setting for postal code and bank cheque processing

# Supervised learning

- Example 3: face detection



- One approach consists of sliding a window over the image
- The goal is to classify each window into one of the two possible classes: face or not-face
- The original problem is transformed into a classification problem

# Supervised learning


- Example 3: face detection



- We start with a set of face images with different variations such as age, gender, illumination, pose, but no translations
- ... and a larger set of images that do not contain full faces

# Supervised learning

- Example 4: spam detection



The image shows a screenshot of an email inbox. A single email message is selected, indicated by a blue border. The message is from 'rama rama ramaumar002@hotmail.com via yahoo.com' to 'to <redacted>'. The subject line is 'From: Mrs. Rama Umar' followed by three lines of text: 'Groupe Bank of Africa (Annexe) Burkina Faso', 'Foreign Department Operation.', and 'My name is Mrs.Rama Umar. I am working with Bank of Africa here in Burkina Faso. I have a late foreign customer.' Below this, there is a large block of text: 'When I discovered that there had been neither deposits nor withdrawals from this account, I contacted my family members. None of them knew about this account. The late person had no relatives in Burkina Faso. The amount of the transaction was 5 million US dollars. I am sending you the details of the transaction. Kind regards, Mrs. Rama Umar'.

rama rama ramaumar002@hotmail.com via yahoo.com  
to <redacted>

From: Mrs. Rama Umar

Groupe Bank of Africa (Annexe) Burkina Faso

Foreign Department Operation.

My name is Mrs.Rama Umar. I am working with Bank of Africa here in Burkina Faso. I have a late foreign customer.

When I discovered that there had been neither deposits nor withdrawals from this account, I contacted my family members. None of them knew about this account. The late person had no relatives in Burkina Faso. The amount of the transaction was 5 million US dollars. I am sending you the details of the transaction. Kind regards, Mrs. Rama Umar

- The task is to classify an email into spam or not-spam
- The occurrence of the word “Dollars” is a good indicator of spam
- A possible representation is a vector of word frequencies

# We count the words...

obtaining X



rama rama ramaumar002@hotmail.com via yahoo.com  
to ▾

From: Mrs. Rama Umar

Groupe Bank of Africa (Annexe) Burkina Faso

Foreign Department Operation.

My name is Mrs.Rama Umar. I am working with Bank of Africa here in Burkina Faso. I have been working here for 10 years. I am writing to you because I have found a late foreign customer.

When I discovered that there had been neither deposits nor withdrawals from this account, I checked the records of all the accounts in the name of the family member or relations of the late person are aware of this account, I found one account which was opened by a late person in his name. The amount in the account is (Five Million USA Dollars).

free	100
money	2
:	:
account	2
:	:



**Yoshua Bengio** <yoshua.bengio@gmail.com>

to Dong-Hyun, Ian, Dumitru, Pierre, Aaron, Mehdi, Ben, Will, Charlie,

Nice slides!

See you next week,

—Yoshua

free	1
money	1
:	:
account	2
:	:

# The spam detection algorithm



free	100
money	2
:	:
account	2
:	:

Why these words?

$$\begin{pmatrix} 100 \times 0.2 \\ 2 \times 0.3 \\ \vdots \\ 2 \times 0.3 \\ \vdots \end{pmatrix}$$

$$= 3.2$$



Confidence /  
performance  
guarantee?

$$\begin{pmatrix} 100 \times 0.01 \\ 2 \times 0.02 \\ \vdots \\ 2 \times 0.01 \\ \vdots \end{pmatrix}$$


$$= 1.03$$

Why linear  
combination?

Where do the weights  
come from?

# Supervised learning

- Example 5: predicting stock prices on the market



- The goal is to predict the price at a future date, for example in a few days
- This is a regression task, since the output is continuous

# Supervised learning

- Example 6: image difficulty prediction [Ionescu et al. CVPR2016]



2.78



2.82



3.30



3.62



3.80

easy

image difficulty score

hard

2.81



3.15



3.45




3.64




- The goal is to predict the time necessary for a human to solve a visual search task (**data set available for project 2!**)
- This is a regression task, since the output is continuous

# Canonical forms of supervised learning problems

- Classification




- Regression




# Age estimation in images

- Classification?



- Regression?



What age?

# The supervised learning paradigm

Functions  $\mathcal{F}$


$$\textcolor{red}{f} : \mathcal{X} \rightarrow \mathcal{Y}$$

Training data

$$\{(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$$

LEARNING

$$\begin{aligned} &\text{find } \hat{f} \in \mathcal{F} \\ &\text{s.t. } y_i \approx \hat{f}(x_i) \end{aligned}$$



PREDICTION

$$\textcolor{red}{y} = \hat{f}(x)$$

New data



$$x$$

# Supervised learning models

- Naive Bayes (lecture 2)
- k-Nearest Neighbors (lecture 3)
- Decision trees and random forests (lecture 4)
- Support Vector Machines (lecture 5, 6)
- Kernel methods (lecture 5)
- Kernel Ridge Regression (lecture 5)
- Neural networks (lectures 7, 8, 9)
- Many others...


# Unsupervised learning

- We have an unlabeled training set of samples
- Example 1: clustering images based on similarity




# Unsupervised learning

- Example 1: clustering MNIST images based on similarity [Georgescu et al. ICIP2019]




# Unsupervised learning

- Example 2: unsupervised features learning




# Unsupervised learning

- Example 2: unsupervised features learning for abnormal event detection [Ionescu et al. CVPR2019]



# Unsupervised learning


- Example 3: clustering mammals by family, species, etc.




- The task is to generate the phylogenetic tree based on DNA

# Canonical forms of unsupervised learning problems

- Clustering



- Dimensionality Reduction



# Unsupervised learning models


- K-means clustering (lecture 10, 11)
- DBScan (lecture 12)
- Hierarchical clustering (lecture 12)
- Principal Component Analysis (lecture 13)
- t-Distributed Stochastic Neighbor Embedding (lecture 13)
- Hidden Markov Models
- Many others...

# Semi-supervised learning

- We have a training set of samples that are partially annotated with class labels
- Example 1: object recognition in images, some of which are annotated with corresponding class labels



Car



Person




Dog

# Reinforcement learning

- How does it work?
- The system learns intelligent behavior using a reinforcement signal (reward)
- The reward is given after several actions are taken (it does come after every action)
- Time matters (data is sequential, not i.i.d.)
- The actions of the system can influence the data


# Reinforcement learning

- Example 1: learning to play Go
- +/- reward for winning / losing the game





# Reinforcement learning

- Example 2: teaching a robot to ride a bike
- +/- reward for moving forward / falling




# Reinforcement learning


- Example 3: learning to play Pong from image pixels
- +/- reward for increasing
- personal / adversary score




# Reinforcement learning paradigm



# Formalizing as Markov Decision Process



# Formalizing as Markov Decision Process



# Formalizing as Markov Decision Process


- Solution based on dynamic programming (small graphs) or approximation (large graphs)
- Goal: select the actions that maximize the total final reward
- The actions can have long-term consequences
- Sacrificing the immediate reward can lead to higher rewards on the long term

# Formalizing as Markov Decision Process

- AlphaGo example:
  - Narrator 1: “That’s a very strange move”
  - Narrator 2: “I thought it was a mistake”
  - But actually, “the move turned the course of the game. AlphaGo went on to win Game Two, and at the post-game press conference, Lee Sedol was in shock.”
  - <https://www.wired.com/2016/03/two-moves-alphago-lee-sedol-redefined-future/>


# Active learning

- Given a large set of unlabeled samples, we have to choose a small subset for annotation in order to obtain a good classification model




# Transfer learning

- Starting with a model trained for a certain task / domain, use the model for a different task / domain




# Transfer learning

- Adapt the model to specific test samples
- Example 1: facial expression recognition [Georgescu et al. Access2019]




# Transfer learning

- Example 2: zero-shot learning




# Many interesting applications, but...

- What is ethical and what is not?
- Trolley paradox



# Many interesting applications, but...

- What is ethical and what is not?
- Trolley paradox



# Many interesting applications, but...

- What is ethical and what is not?
- Trolley paradox
- <http://moralmachine.mit.edu>

# Bibliography

Springer Series in Statistics

Trevor Hastie  
Robert Tibshirani  
Jerome Friedman

# The Elements of Statistical Learning


Data Mining, Inference, and Prediction

Second Edition

 Springer

Richard O. Duda  
Peter E. Hart  
David G. Stork

# Pattern Classification




Second Edition

O'REILLY®

# Hands-On Machine Learning with Scikit-Learn & TensorFlow

CONCEPTS, TOOLS, AND TECHNIQUES  
TO BUILD INTELLIGENT SYSTEMS



powered by  
 jupyter

Aurélien Géron

Advances in Computer Vision and Pattern Recognition



Radu Tudor Ionescu  
Marius Popescu

# Knowledge Transfer between Computer Vision and Text Mining

Similarity-based Learning Approaches

 Springer