

Computer Vision

Bogdan Alexe

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2020-2021

We



Bogdan



Iuliana

407
(AI)



You

Computer Vision

Nr.	Nume si prenume	Grupa
1	Badea Armeanu Razvan Ilie	405
2	Grigoras Dorin	405
3	Năbirgeac Cezar Ștefan	405
4	Onetiu Radu	405
5	Andrei Tudor Ionescu	406
6	Anei Alexandra Gabriela	406
7	Apostu Robert	406
8	Brînzea Alexandru Eusebiu	406
9	Ciltea Marian	406
10	Costachi T. Ana Maria	406
11	Darii Dan	406
12	Matei Bianca Gabriela	406
13	Munteanu Beatrice Madalina	406
14	Petrasco Sandu	406
15	Pogonaru Stefan	406
16	Stinga Andreea	406
17	Valeria Bolun	406
18	Hodorogea Dragos Cristian	410
19	Acsintoae Andra Maria	411
20	Aslantas Serdal	411
21	Chirila Diana Gabriela	411
22	Chirita Catalina	411
23	Coman Catalina Mihaela	411
24	Cutic Eugen	411
25	Dutu Daniel	411
26	Fulea Andrei	411
27	Menadil Rafael Edy	411
28	Patulea Mihaela Alina	411
29	Deaconu Stefan Eduard	412
30	Iordache Ioan Bogdan	412
31	Manea Andrei Alexandru	412

**MASTER Anul I, 2020 - 2021, Învățământ cu frecvență
„ARTIFICIAL INTELLIGENCE”, Program de 4 semestre**

(Sem.I – 14 săptămâni; Sem.II – 14 săptămâni)

Nr. crt.	Cursuri obligatorii	Semestrul I (14 săpt.)				Semestrul II (14 săpt.)			
		C	S/L	E/V	ECTS	C	S/L	E/V	ECTS
1	Ob.11 Învățare automata aplicată Practical Machine Learning	2	1	V	6	-	-	-	-
2	Ob.12 Programare probabilistă Probabilistic Programming	2	1	E	6	-	-	-	-
3	Ob.13 Reprezentarea cunoștințelor și inferență Knowledge representations and reasoning	2	1	E	6	-	-	-	-
4	Op.14 Curs optional Optional Course	2	1	E	6	-	-	-	-
5	Ob.15 Practică/Practical Training	-	2	V	4	-	-	-	-
6	Ob.16 Deontologie Academică/Academic Deontology	1	-	V	2				
7	Ob.21 Învățare automată Advanced machine learning	-	-	-	-	2	1	E	6
8	Ob.22 Vedere Artificială Computer Vision	-	-	-	-	2	1	E	6
9	Ob.23 Procesarea limbajului natural 1 Natural language processing 1	-	-	-	-	2	1	E	6
10	Op.24 Curs optional Optional Course					2	1	E	6
11	Ob.25 Practică/ Practical Training	-	-	-	-	-	2	V	6
TOTAL		9	6	3E 3V	30	8	6	4E 1V	30

**MASTER Anul I, 2020- 2021, Învățământ cu frecvență
„DATA SCIENCE”, Program de 4 semestre**

(Sem.I – 14 săptămâni; Sem.II – 14 săptămâni)

Nr. crt.	Cursuri obligatorii și optionale	Semestrul I (14 săpt.)				Semestrul II (14 săpt.)			
		C	S/L	E/V	ECTS	C	S/L	E/V	ECTS
1	Ob.11 Învățare automată aplicată Practical Machine Learning	2	1	E	6	-	-	-	-
2	Ob.12 Programare probabilistă Probabilistic Programming	2	1	E	6	-	-	-	-
3	Ob.13 Analiza exploratorie a datelor Exploratory Data Analysis	2	1	E	6	-	-	-	-
4	Ob.14 Curs optional 1 Optional Course 1	2	1	E	6	-	-	-	-
5	Ob.16 Deontologie Academică/Academic Deontology	1	-	V	2				
6	Ob.15 Practică/Practical Training	-	2	V	4	-	-	-	-
7	Ob.21 Învățare automată Advanced Machine Learning	-	-	-	-	2	1	E	6
8	Ob.22 Big Data Big Data	-	-	-	-	2	1	E	6
9	Ob.23 Metode statistice pentru data science Statistics for Data Science	-	-	-	-	2	1	E	6
10	Ob.24 Curs optional 2 Optional Course 2					2	1	E	6
11	Ob.25 Practică/ Practical Training	-	-	-	-	-	2	V	6
TOTAL		9	6	4E 2V	30	8	6	4E 1V	30

ADMINISTRATIVE

- Lecture: Tuesday, 16-18 (Teams)
- Lab: Thursday, 18-20 (Iuliana) or Friday, 8-10 (Bogdan)
- **One lab session every two weeks**
- Bogdan – TEAMS (lecture + lab)
 - channel Lecture: slides + recordings
 - channel Lab: pdfs + recordings + solutions
- Iuliana – Google Meet (lab)
 - group 3: <https://hangouts.google.com/group/DutRAdQKHSKZ2rYP6>
 - group 4: <https://hangouts.google.com/group/bJXFKij6u665dwGK9>

Exam – evaluation in June

Grade = $\min(10, P1 + P2)$

- P1 = project 1 = 5 points (+ some bonus?)
- P2 = project 2 = 5 points (+ some bonus?)
- no constraints, with 4.99 you fail, with 5 you pass

Exam – evaluation in July (restanță)

$$\text{Grade} = \max(2 \times P3, P1 + P3, P2 + P3)$$

- $P3$ = project 3 = 5 points
- no constraints, with 4.99 you fail, with 5 you pass

Exam – evaluation in September (reexaminare)

Grade = $\max(2 \times P4, P1 + P4, P2 + P4, P3 + P4)$

- $P4$ = project 4 = 5 points
- no constraints, with 4.99 you fail, with 5 you pass

Exam – evaluation in September (mărire)

Grade = $2 \times P4$

- $P4 = \text{project 4} = 5 \text{ points}$

Resources

github.com/jbhuang0604/awesome-computer-vision/

Awesome Computer Vision: awesome

A curated list of awesome computer vision resources, inspired by [awesome-php](#).

For a list people in computer vision listed with their academic genealogy, please visit [here](#)

Contributing

Please feel free to send me [pull requests](#) or email (jbhuang@vt.edu) to add links.

Table of Contents

- [Awesome Lists](#)
- [Books](#)
- [Courses](#)
- [Papers](#)
- [Software](#)
- [Datasets](#)
- [Tutorials and Talks](#)
- [Resources for students](#)
- [Blogs](#)
- [Links](#)
- [Songs](#)

Books

Computer Vision

- [Computer Vision: Models, Learning, and Inference](#) - Simon J. D. Prince 2012
- [Computer Vision: Theory and Application](#) - Rick Szeliski 2010
- [Computer Vision: A Modern Approach \(2nd edition\)](#) - David Forsyth and Jean Ponce 2011
- [Multiple View Geometry in Computer Vision](#) - Richard Hartley and Andrew Zisserman 2004
- [Computer Vision](#) - Linda G. Shapiro 2001
- [Vision Science: Photons to Phenomenology](#) - Stephen E. Palmer 1999
- [Visual Object Recognition synthesis lecture](#) - Kristen Grauman and Bastian Leibe 2011
- [Computer Vision for Visual Effects](#) - Richard J. Radke, 2012
- [High dynamic range imaging: acquisition, display, and image-based lighting](#) - Reinhard, E., Heidrich, W., Debevec, P., Pattanaik, S., Ward, G., Myszkowski, K 2010
- [Numerical Algorithms: Methods for Computer Vision, Machine Learning, and Graphics](#) - Justin Solomon 2015

OpenCV Programming

- [Learning OpenCV: Computer Vision with the OpenCV Library](#) - Gary Bradski and Adrian Kaehler
- [Practical Python and OpenCV](#) - Adrian Rosebrock
- [OpenCV Essentials](#) - Oscar Deniz Suarez, M^a del Milagro Fernandez Carrobles, Noelia Vallez Enano, Gloria Bueno Garcia, Ismael Serrano Gracia

Courses

Computer Vision

- [EENG 512 / CSCI 512 - Computer Vision](#) - William Hoff (Colorado School of Mines)
- [Visual Object and Activity Recognition](#) - Alexei A. Efros and Trevor Darrell (UC Berkeley)
- [Computer Vision](#) - Steve Seitz (University of Washington)
- [Visual Recognition Spring 2016, Fall 2016](#) - Kristen Grauman (UT Austin)
- [Language and Vision](#) - Tamara Berg (UNC Chapel Hill)
- [Convolutional Neural Networks for Visual Recognition](#) - Fei-Fei Li and Andrej Karpathy (Stanford University)
- [Computer Vision](#) - Rob Fergus (NYU)
- [Computer Vision](#) - Derek Hoiem (UIUC)
- [Computer Vision: Foundations and Applications](#) - Kalanit Grill-Spector and Fei-Fei Li (Stanford University)
- [High-Level Vision: Behaviors, Neurons and Computational Models](#) - Fei-Fei Li (Stanford University)
- [Advances in Computer Vision](#) - Antonio Torralba and Bill Freeman (MIT)
- [Computer Vision](#) - Bastian Leibe (RWTH Aachen University)
- [Computer Vision 2](#) - Bastian Leibe (RWTH Aachen University)
- [Computer Vision](#) Pascal Fua (EPFL):
- [Computer Vision 1](#) Carsten Rother (TU Dresden):
- [Computer Vision 2](#) Carsten Rother (TU Dresden):
- [Multiple View Geometry](#) Daniel Cremers (TU Munich):

Computer Vision

- [Computer Vision Talks](#) - Lectures, keynotes, panel discussions on computer vision
- [The Three R's of Computer Vision](#) - Jitendra Malik (UC Berkeley) 2013
- [Applications to Machine Vision](#) - Andrew Blake (Microsoft Research) 2008
- [The Future of Image Search](#) - Jitendra Malik (UC Berkeley) 2008
- [Should I do a PhD in Computer Vision?](#) - Fatih Porikli (Australian National University)
- [Graduate Summer School 2013: Computer Vision](#) - IPAM, 2013

Recent Conference Talks

- [CVPR 2015](#) - Jun 2015
- [ECCV 2014](#) - Sep 2014
- [CVPR 2014](#) - Jun 2014
- [ICCV 2013](#) - Dec 2013
- [ICML 2013](#) - Jul 2013
- [CVPR 2013](#) - Jun 2013
- [ECCV 2012](#) - Oct 2012
- [ICML 2012](#) - Jun 2012
- [CVPR 2012](#) - Jun 2012

What is Computer Vision?



Are we done?

What is Computer Vision?

Make computers understand images and video.



What kind of scene?

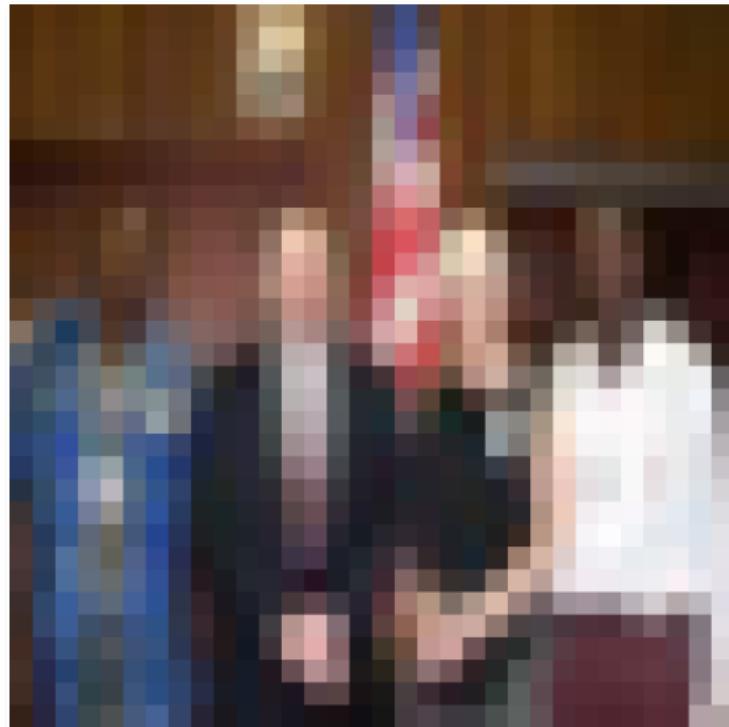
Where are the cars?

How far is the
building?

...

The goal of Computer Vision

- Extract information from pixels



The goal of Computer Vision

- Extract information from pixels



Human vision

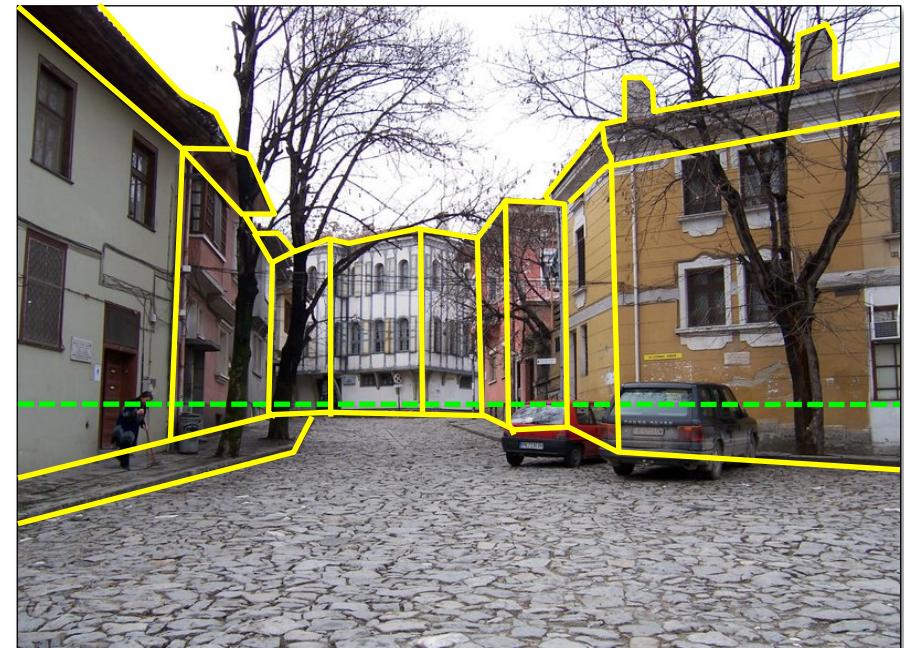
R: 99 G: 95 B: 95	R: 177 G: 172 B: 178	R: 81 G: 79 B: 83	R: 54 G: 49 B: 54	R: 55 G: 46 B: 50	R: 48 G: 34 B: 40	R: 61 G: 45 B: 46	R: 68 G: 53 B: 53	R: 56 G: 43 B: 44	R: 63 G: 49 B: 52	R: 63 G: 54 B: 56
R: 88 G: 84 B: 85	R: 154 G: 148 B: 154	R: 83 G: 81 B: 87	R: 43 G: 42 B: 47	R: 48 G: 42 B: 46	R: 55 G: 44 B: 47	R: 69 G: 53 B: 53	R: 68 G: 53 B: 53	R: 56 G: 43 B: 45	R: 63 G: 49 B: 52	R: 63 G: 54 B: 56
R: 84 G: 79 B: 80	R: 138 G: 133 B: 140	R: 100 G: 98 B: 105	R: 54 G: 51 B: 57	R: 46 G: 41 B: 48	R: 49 G: 41 B: 45	R: 56 G: 43 B: 44	R: 63 G: 49 B: 52	R: 68 G: 53 B: 53	R: 63 G: 49 B: 52	R: 63 G: 54 B: 56
R: 72 G: 66 B: 71	R: 97 G: 92 B: 99	R: 86 G: 84 B: 92	R: 51 G: 50 B: 56	R: 49 G: 46 B: 50	R: 50 G: 43 B: 48	R: 63 G: 49 B: 52	R: 68 G: 53 B: 53	R: 63 G: 49 B: 52	R: 63 G: 54 B: 56	R: 63 G: 54 B: 56
R: 76 G: 72 B: 76	R: 81 G: 79 B: 85	R: 69 G: 69 B: 77	R: 60 G: 59 B: 67	R: 63 G: 59 B: 65	R: 52 G: 45 B: 51	R: 63 G: 49 B: 52	R: 68 G: 53 B: 53	R: 63 G: 49 B: 52	R: 63 G: 54 B: 56	R: 63 G: 54 B: 56

Computer vision

What kind of information?

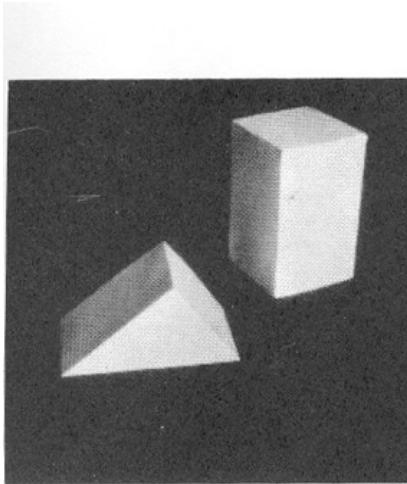


Semantic information

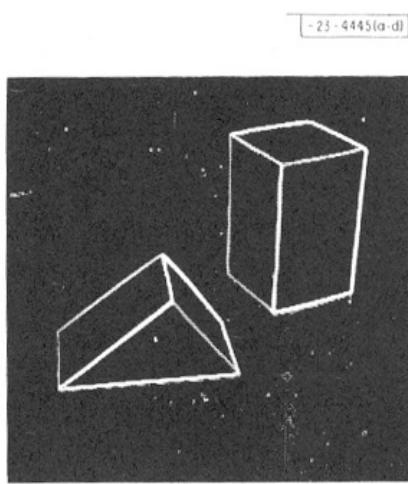


Geometric information (3D)

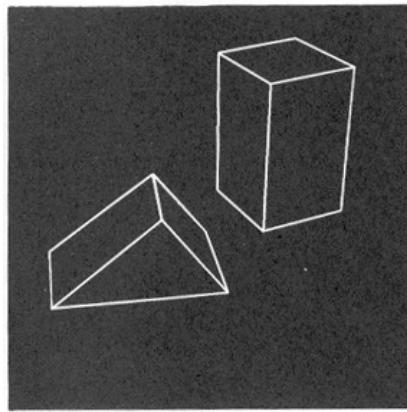
Visual data in 1963



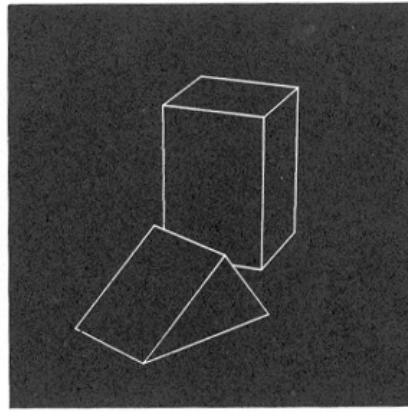
(a) Original picture.



(b) Differentiated picture.



(c) Line drawing.



(d) Rotated view.

- 23 - 4445(a-d)

L. G. Roberts

Machine Perception of Three
Dimensional Solids,

Doctoral thesis, MIT, 1963.

Visual data in 2021



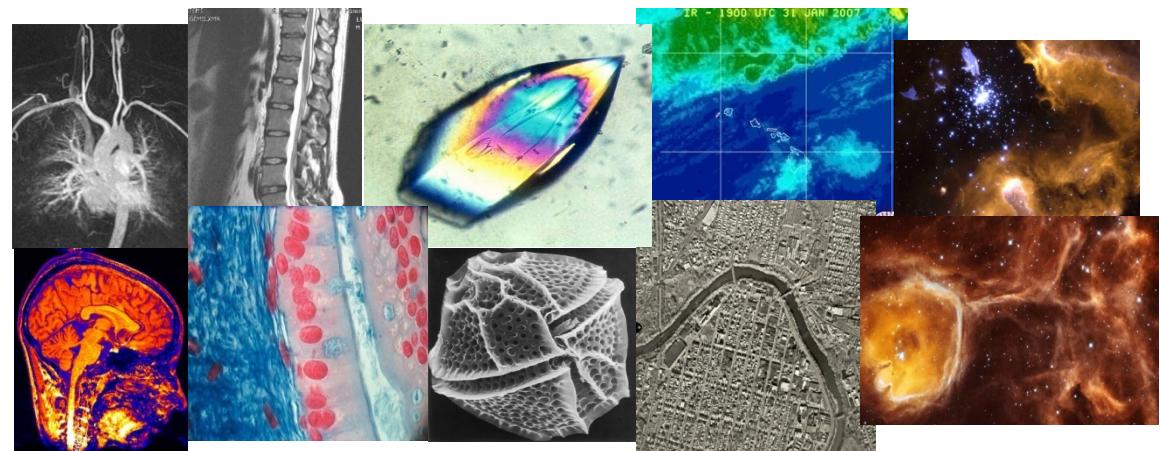
Photo albums



Movies, news, sports



Video surveillance and security



Medical imaging

Slide credit S. Lazebnik

Successful applications in Computer Vision

OCR



Xbox Kinect



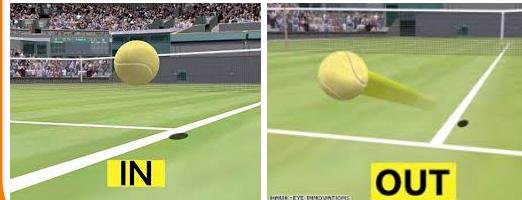
Visual search



Face detection /
recognition



Hawk-eye
decision system



Driving assistance
systems



Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. Object Recognition: high – level vision

Object classification, object detection, part based models, bovw models

5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

Computer Vision projects for undergraduate students

Context aware image resizing

Initial image



Resize the image:
increase/decrease width/height of the image

Main idea: add/remove
irregularly shaped “seams”



How to chose?

Context aware image resizing



Context aware resizing



Traditional resizing

Main idea: Context aware image resizing



Context aware resizing

Intuition:

- Preserve the most “interesting” content
 - Prefer to remove pixels with low gradient energy
- To reduce or increase size in one dimension, remove or add irregularly shaped “seams”
 - Optimal solution via dynamic programming.

Main idea: seam carving



image I

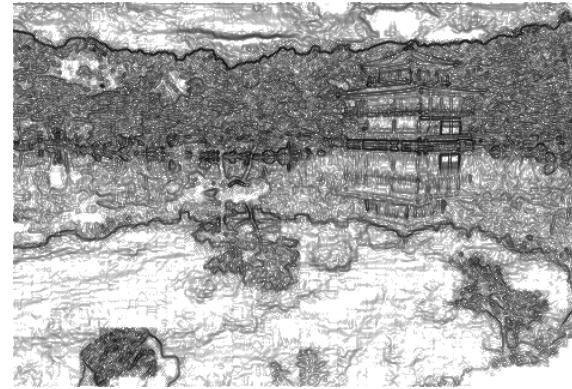


Image gradient

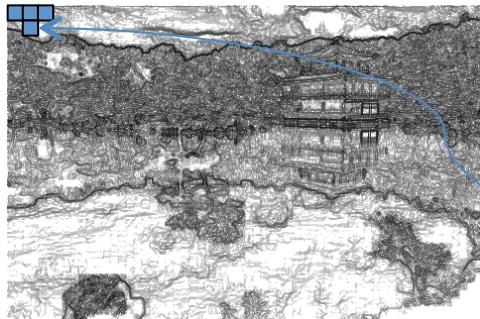
$$\nabla I = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2}$$

- Want to remove seams where they won't be very noticeable:
 - measure “energy” as gradient magnitude
- Choose seam based on **minimum total energy path** across image, subject to 8-connectedness.

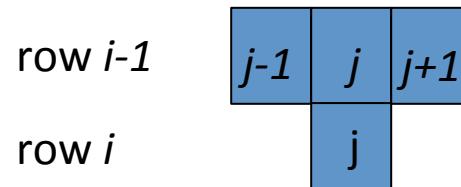
Seam carving: algorithm

- Compute the cumulative minimum energy for all possible connected seams at each entry (i,j) :

$$\mathbf{M}(i, j) = MagGradient(i, j) + \min(\mathbf{M}(i - 1, j - 1), \mathbf{M}(i - 1, j), \mathbf{M}(i - 1, j + 1))$$



Energy matrix
(gradient magnitude)



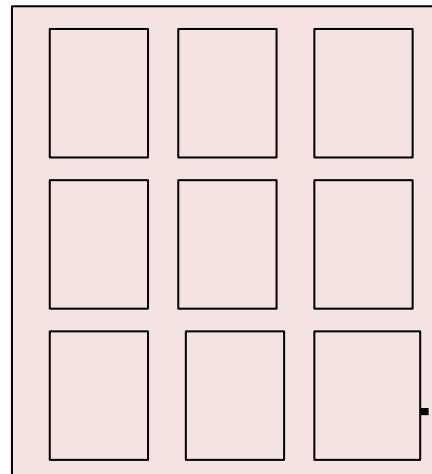
**M matrix: cumulative min energy
(for vertical seams)**

- Then, min value in last row of \mathbf{M} indicates end of the minimal connected vertical seam.
- Backtrack up from there, selecting min of 3 above in \mathbf{M} .

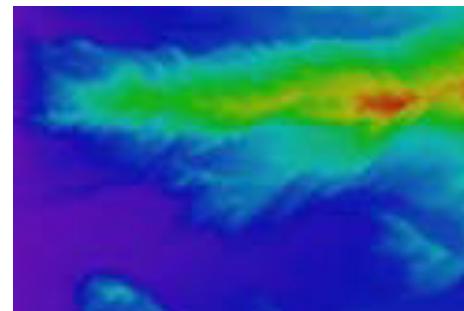
Example

$$\mathbf{M}(i, j) = Energy(i, j) + \min(\mathbf{M}(i - 1, j - 1), \mathbf{M}(i - 1, j), \mathbf{M}(i - 1, j + 1))$$

1	3	0
2	8	9
5	2	6



**Energy matrix
(gradient magnitude)**



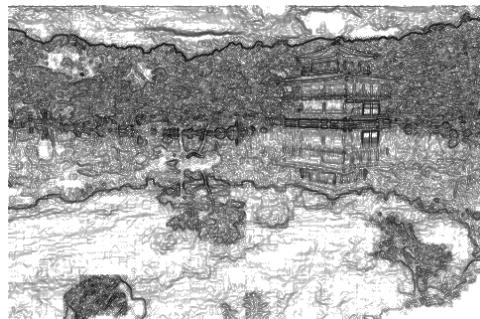
**M matrix
(for vertical seams)** Slide credit K. Grauman

Example

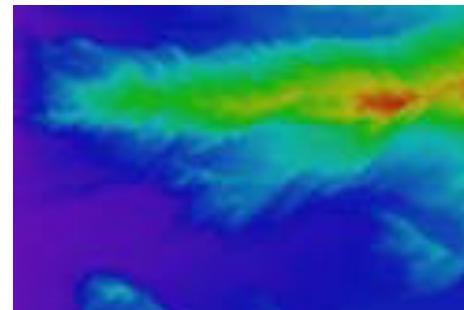
$$\mathbf{M}(i, j) = Energy(i, j) + \min(\mathbf{M}(i-1, j-1), \mathbf{M}(i-1, j), \mathbf{M}(i-1, j+1))$$

1	3	0
2	8	9
5	2	6

1	3	0
3	8	9
8	5	14



**Energy matrix
(gradient magnitude)**



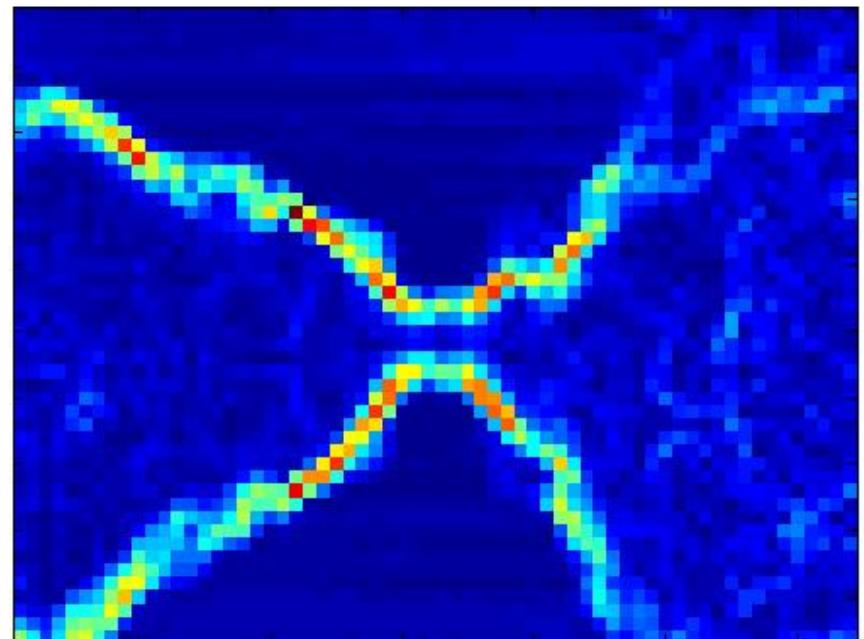
**M matrix
(for vertical seams)** Slide credit K. Grauman

Real image example

Original Image

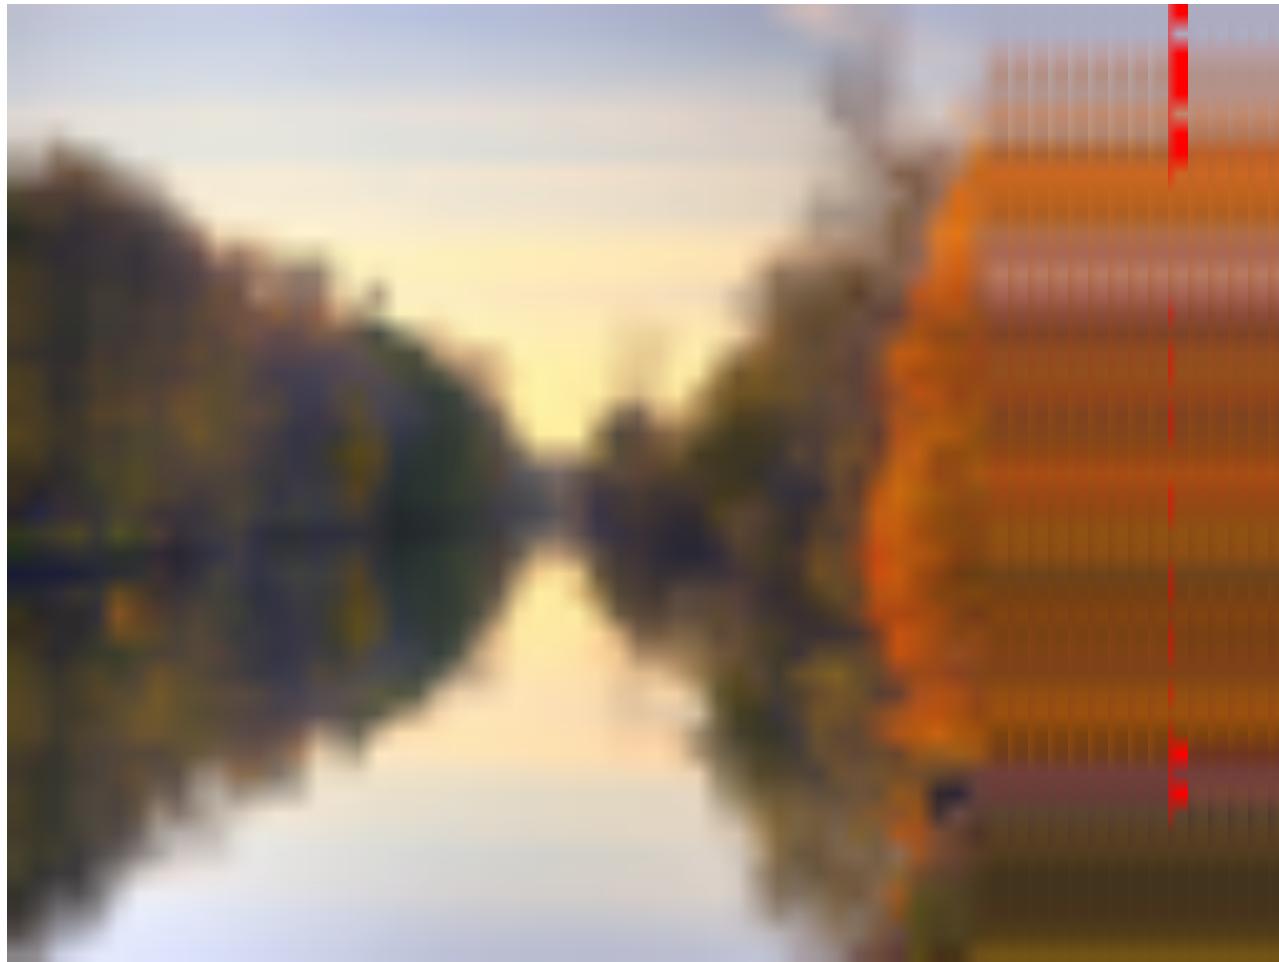


Energy Map



Blue = low energy
Red = high energy

Real image example



Results



**Context aware image
resizing**



Traditional resizing

“Failure cases” with seam carving



Removal of a marked object



Texture synthesis

- Goal: create new samples of a given texture
- Many applications: virtual environments, hole-filling, texturing surfaces

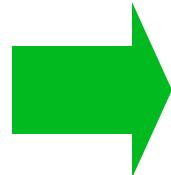
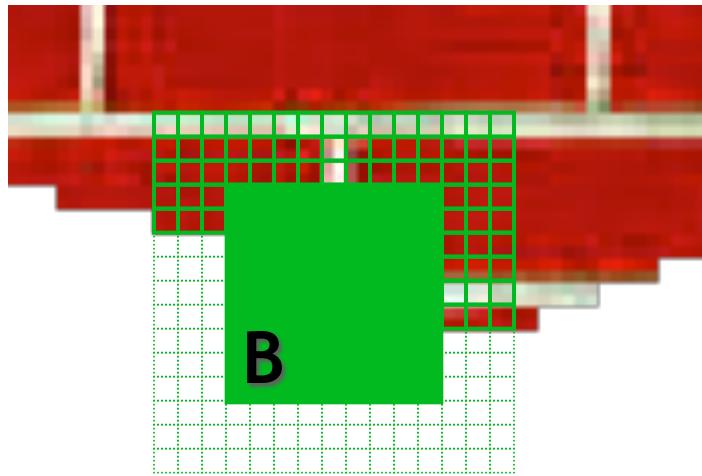
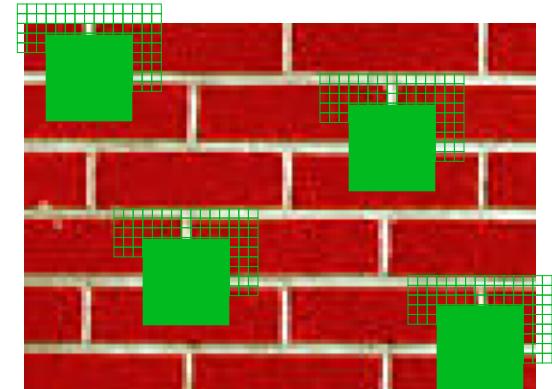


Image Quilting [Efros & Freeman 2001]



Synthesizing a block

non-parametric
sampling

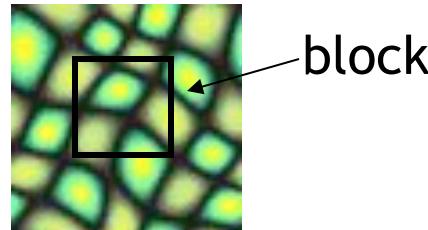


Input image

- Observation: neighbor pixels are highly correlated

Idea: unit of synthesis = block

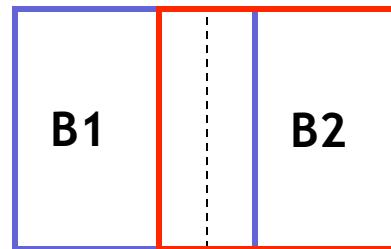
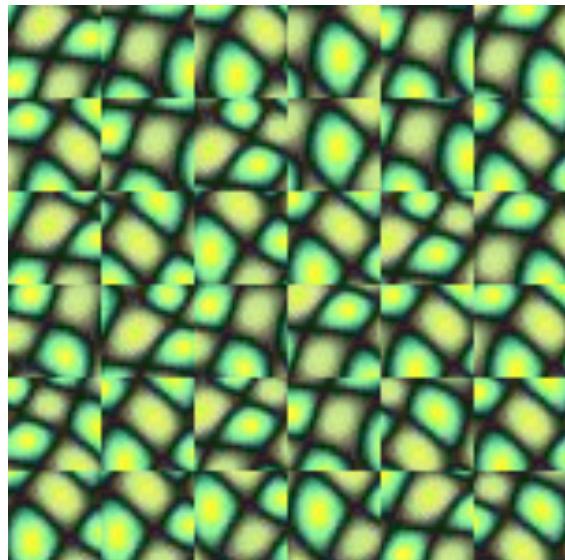
- Exactly the same but now we want $P(B|N(B))$
- Much faster: synthesize all pixels in a block at once



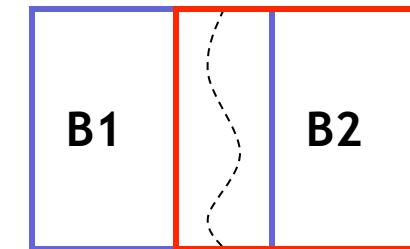
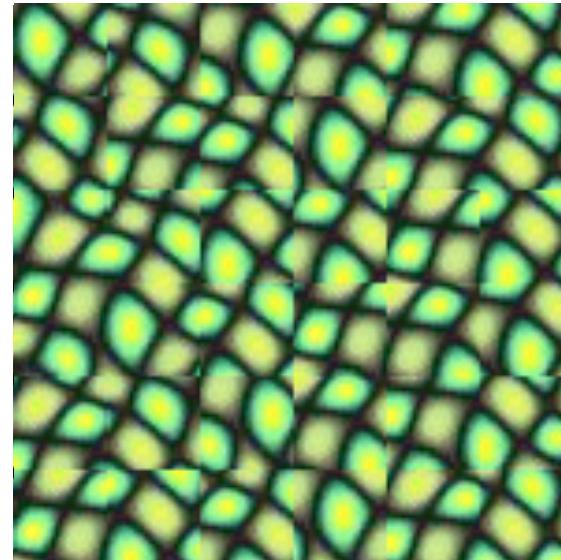
Input texture



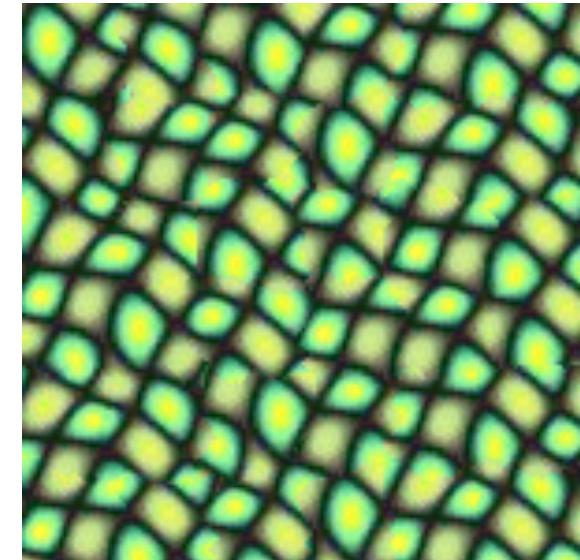
**Random placement
of blocks**



**Neighboring blocks
constrained by overlap**

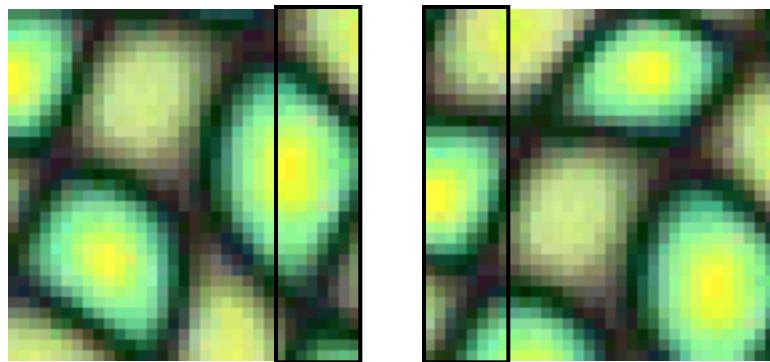


**Minimal error
boundary cut**

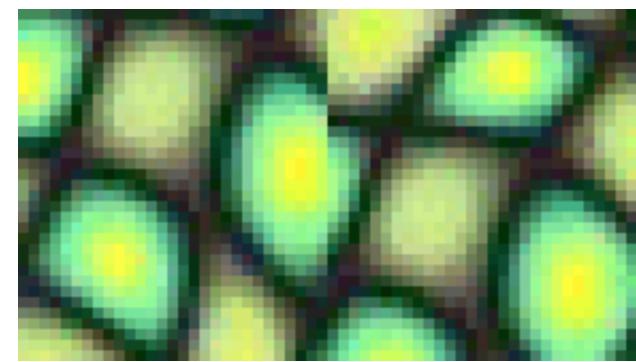


Minimal error boundary

overlapping blocks

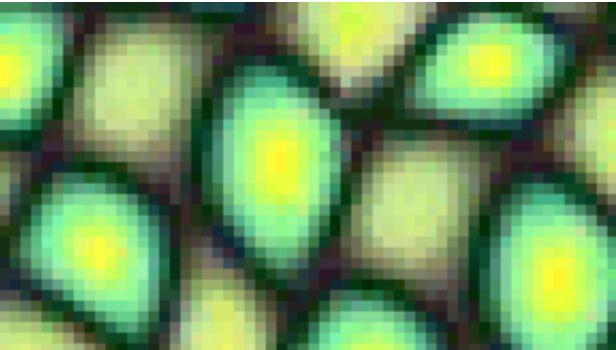


vertical boundary

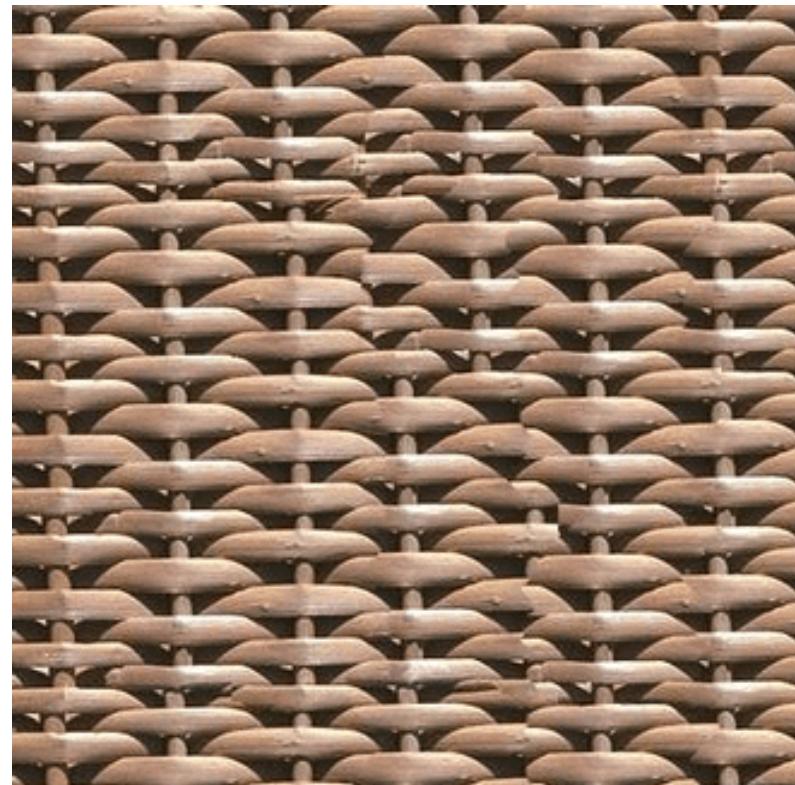
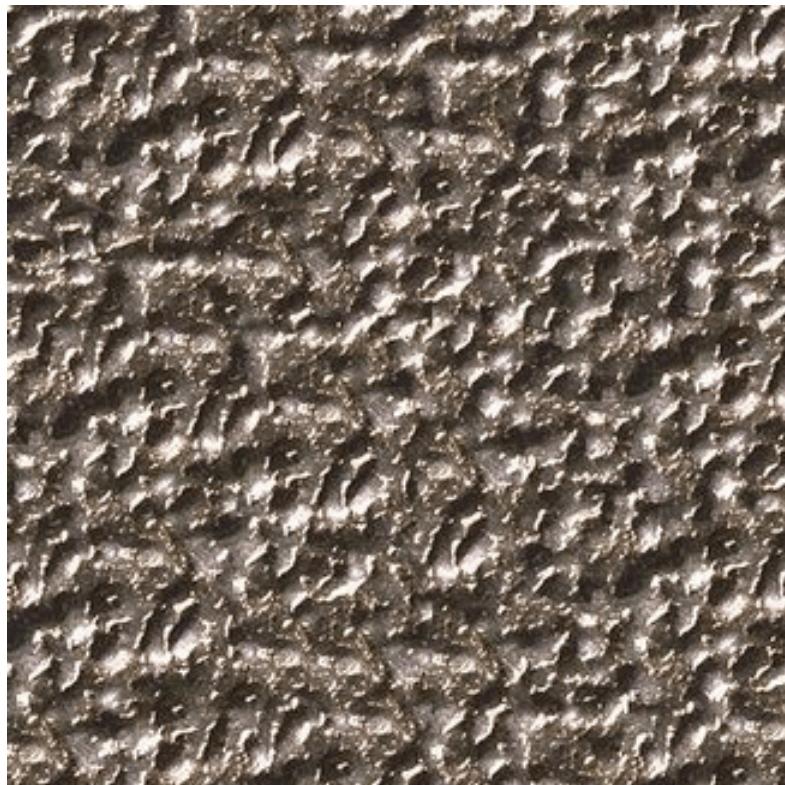
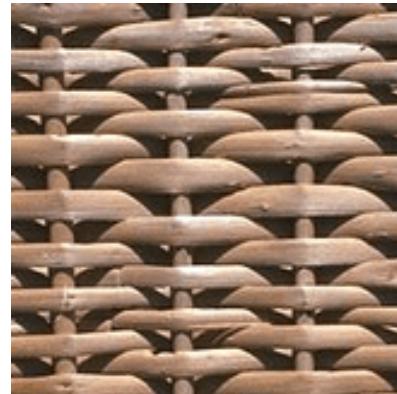


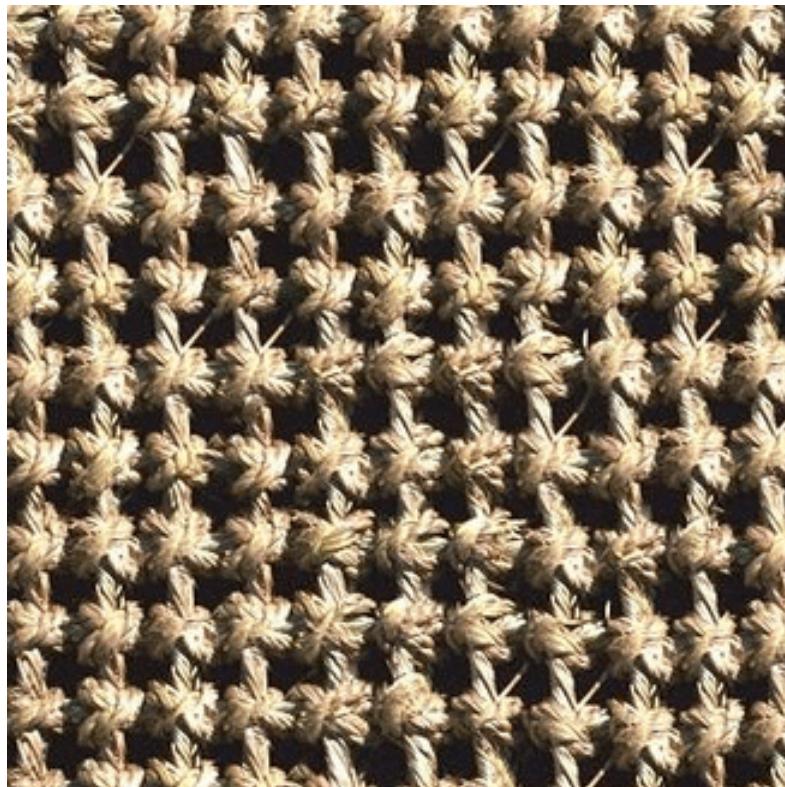
$$\left[\begin{array}{c} \text{overlapping blocks} \\ - \\ \text{vertical boundary} \end{array} \right]^2 = \text{overlap error}$$

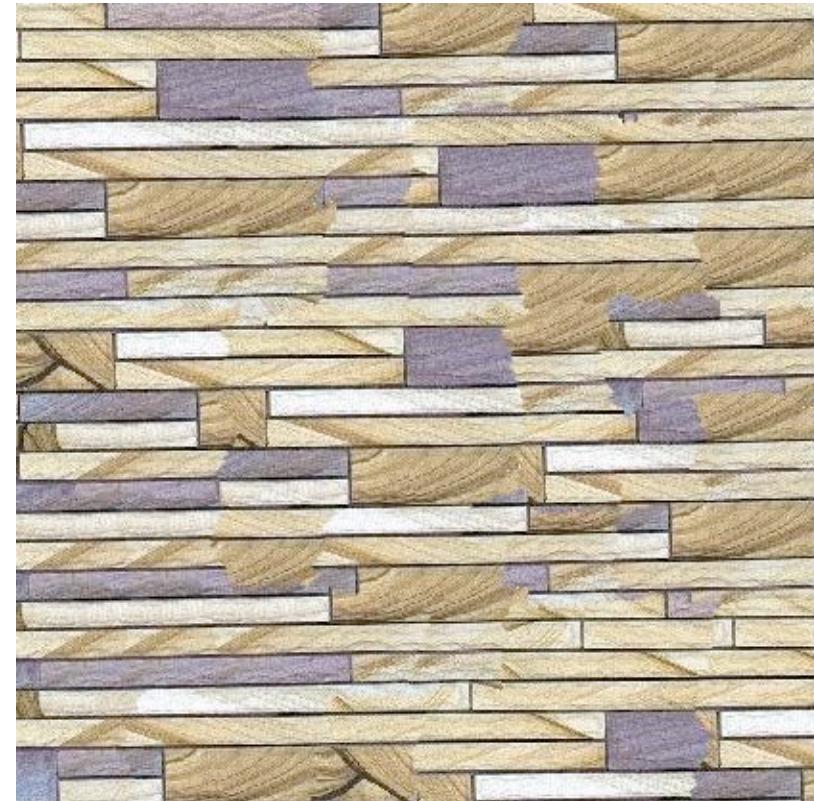
A diagram illustrating the calculation of overlap error. It shows two overlapping blocks of a noisy image (labeled "overlapping blocks") being subtracted from a single image with a vertical boundary (labeled "vertical boundary"). The result is squared (indicated by a large brace and the number 2) to produce a "overlap error" image, which is a binary mask where the boundary pixels are highlighted in red.



min. error boundary

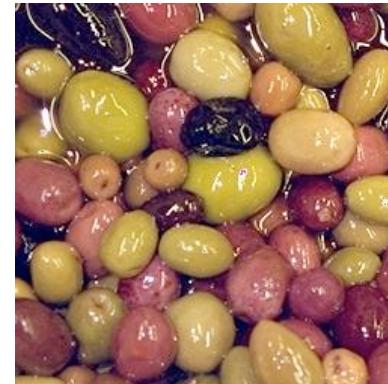






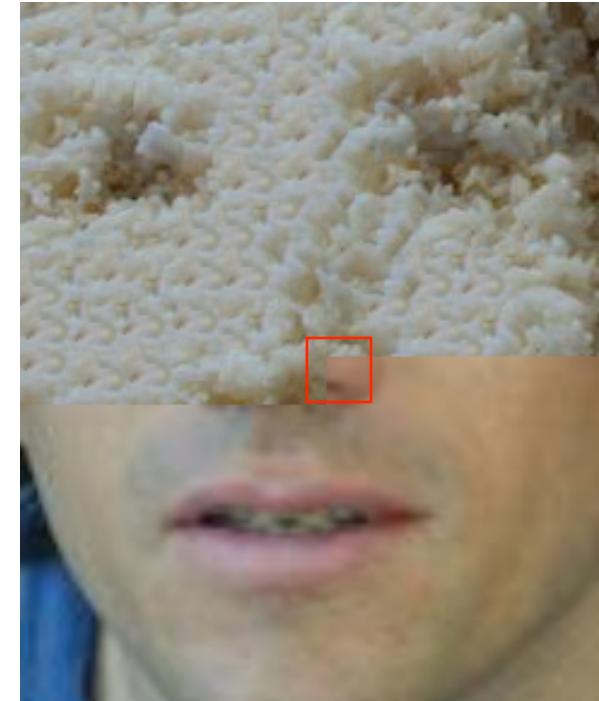


Failures



Texture Transfer

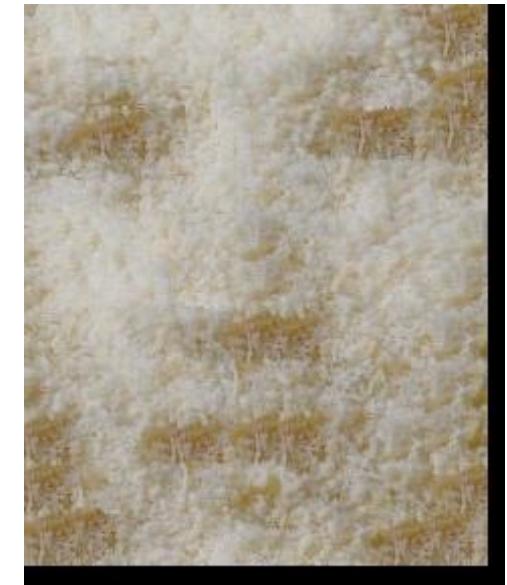
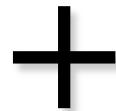
- Take the texture from one object and “paint” it onto another object
 - This requires separating texture and shape
 - That’s HARD, but we can cheat
 - Assume we can capture shape by boundary and rough shading



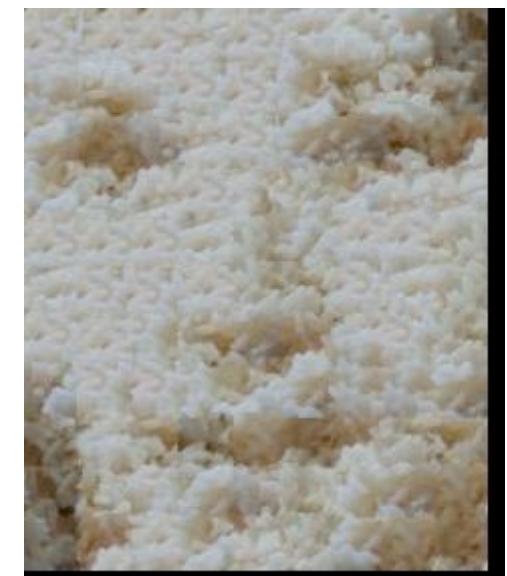
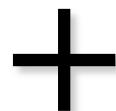
Then, just add another constraint when sampling: similarity to underlying image at that spot



parmesan



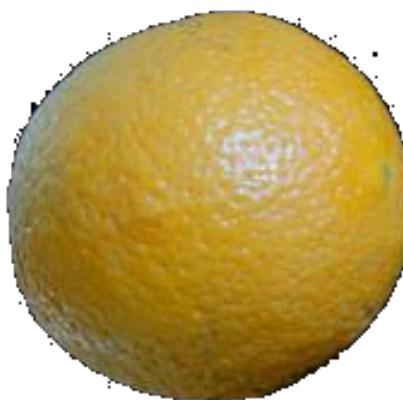
rice



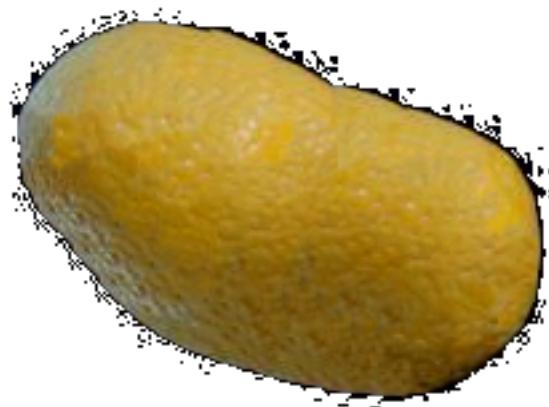
Slide credit A. Efros

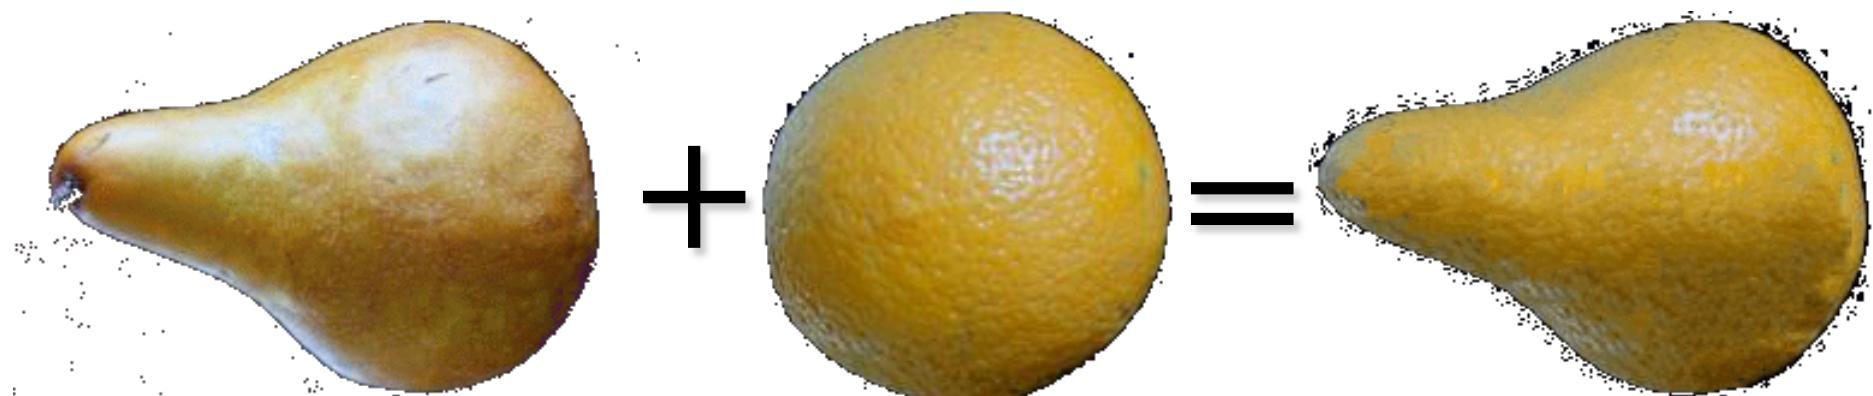


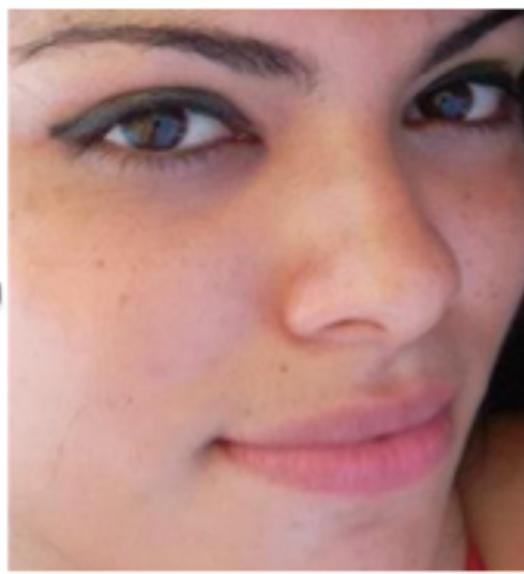
+



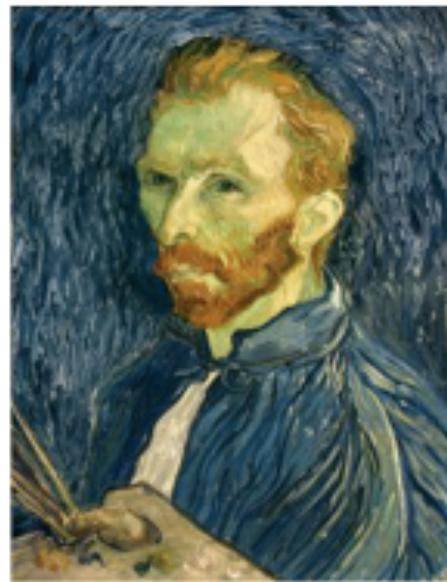
=





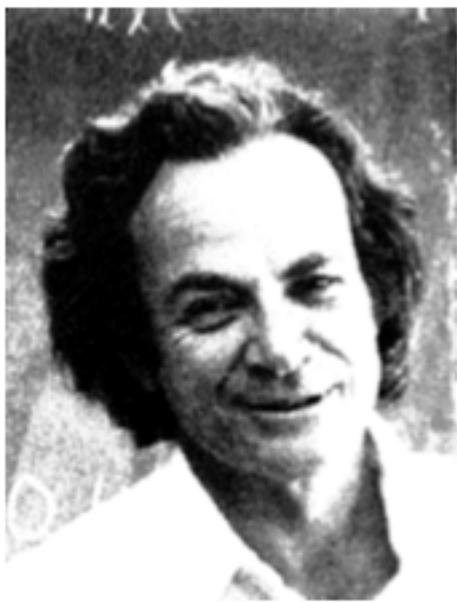


+



=





+

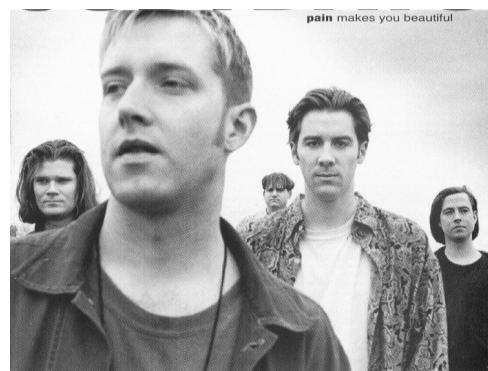


=



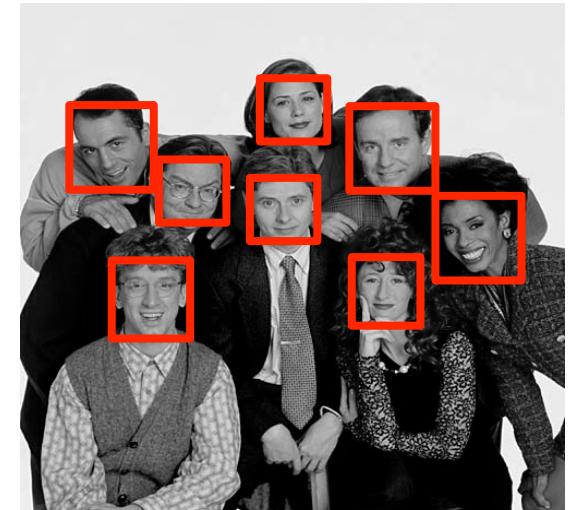
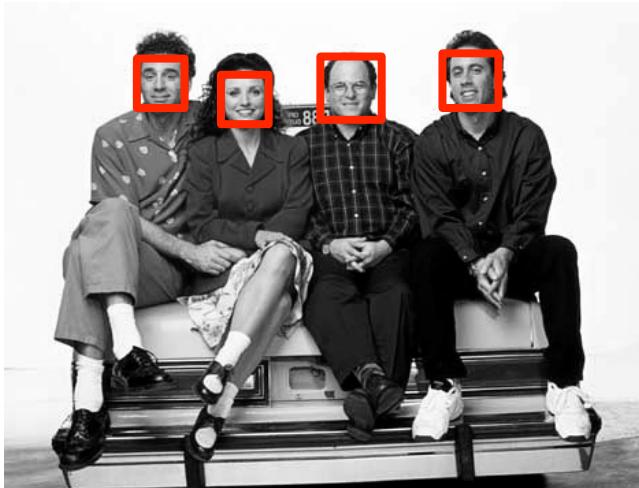
Face detection using sliding window and histograms of oriented gradients

Where are the faces in an image?



Face detection using sliding window and histograms of oriented gradients

Where are the faces in an image?



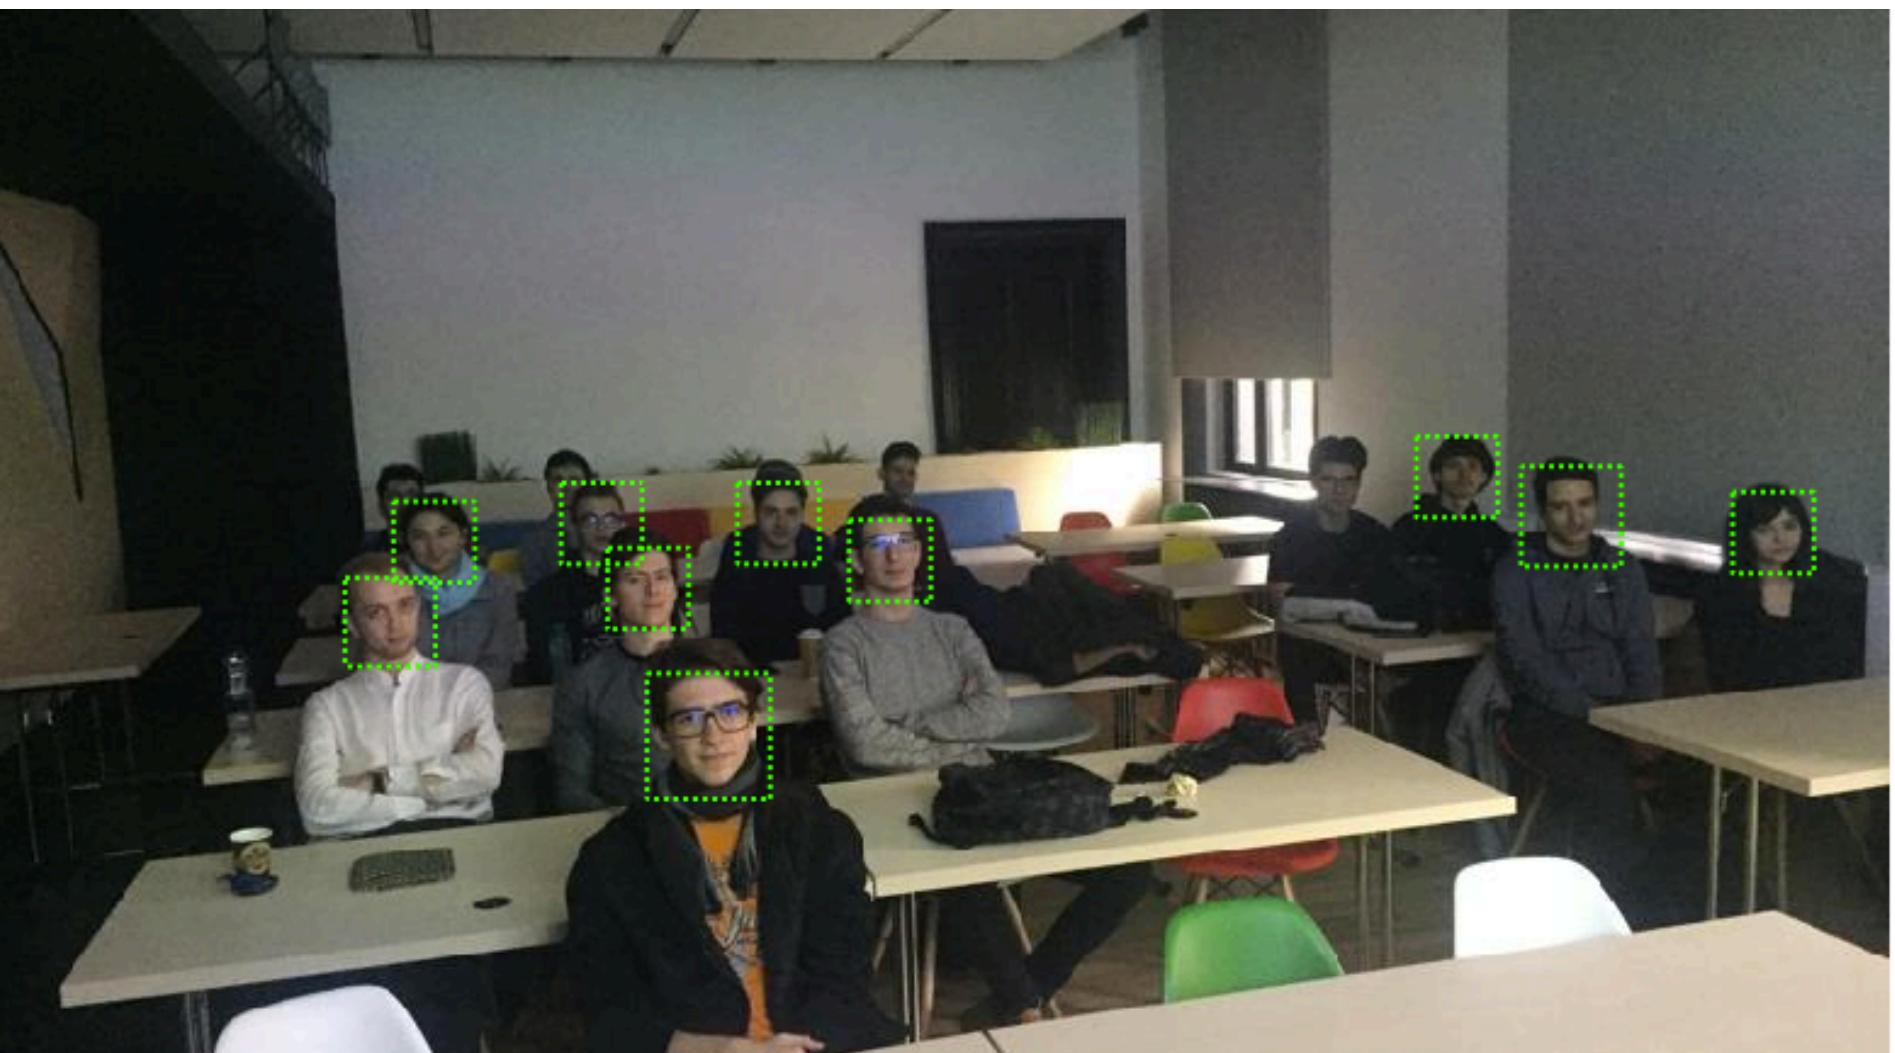
Localization done at the level of a bounding box.

Object detection using sliding window

- detection via classification: **binary classifier** for face → does a window contain a face?
- consider **windows positioned at each pixel**, for **different sizes** (to achieve scale invariance)



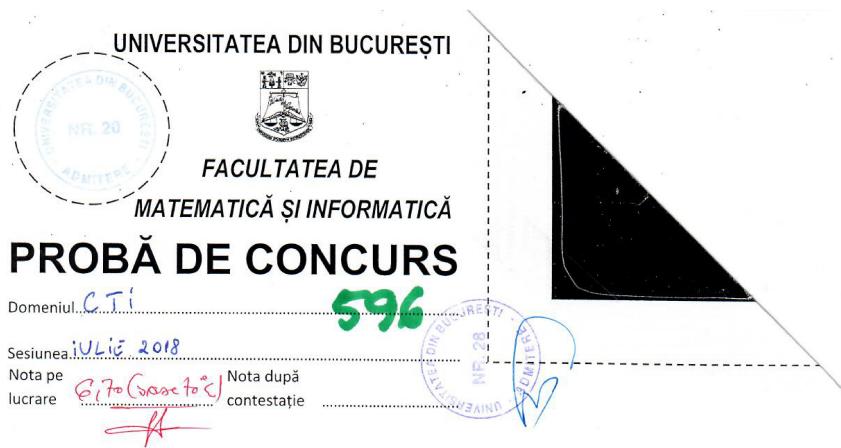






Computer Vision projects
for master students:
Project 1, 2 and 3 proposed last year

Automatic grading of multiple choice tests



MATEMATICĂ

Număr	Răspuns	A	B	C	D
1		X			
2		X			
3		X			
4		X			
5			X		
6		X			
7		X			
8				X	
9			X		
10		X			
11			X		
12			X		
13		X			
14				X	
15				X	

INFORMATICA

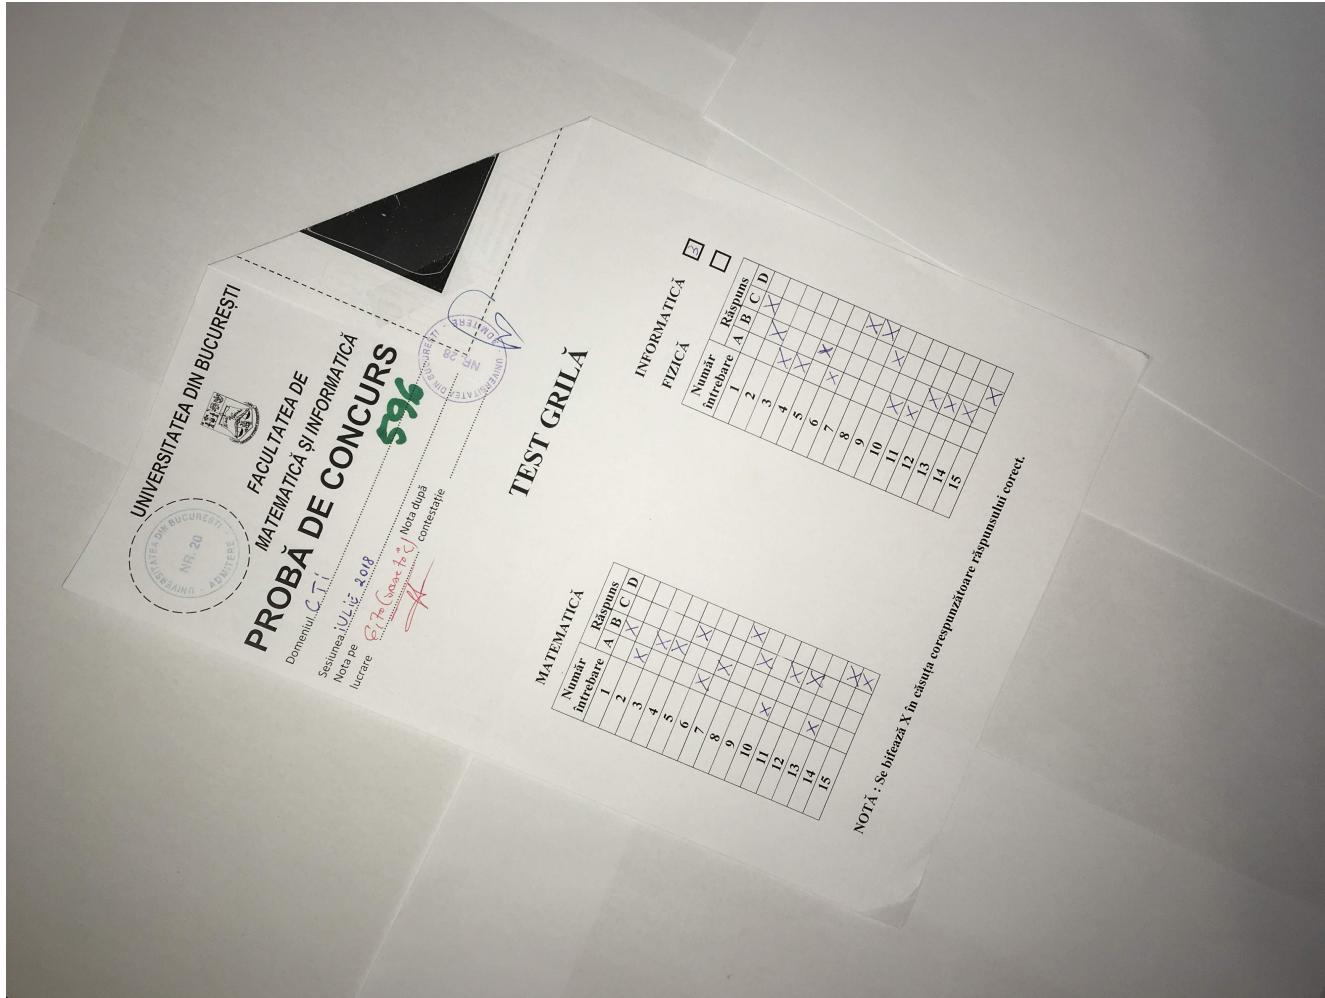
FIZICĂ

Număr	Răspuns	A	B	C	D
1				X	
2			X		
3		X			
4		X			
5			X		
6		X			
7				X	
8				X	
9			X		
10		X			
11		X			
12			X		
13		X			
14			X		
15				X	

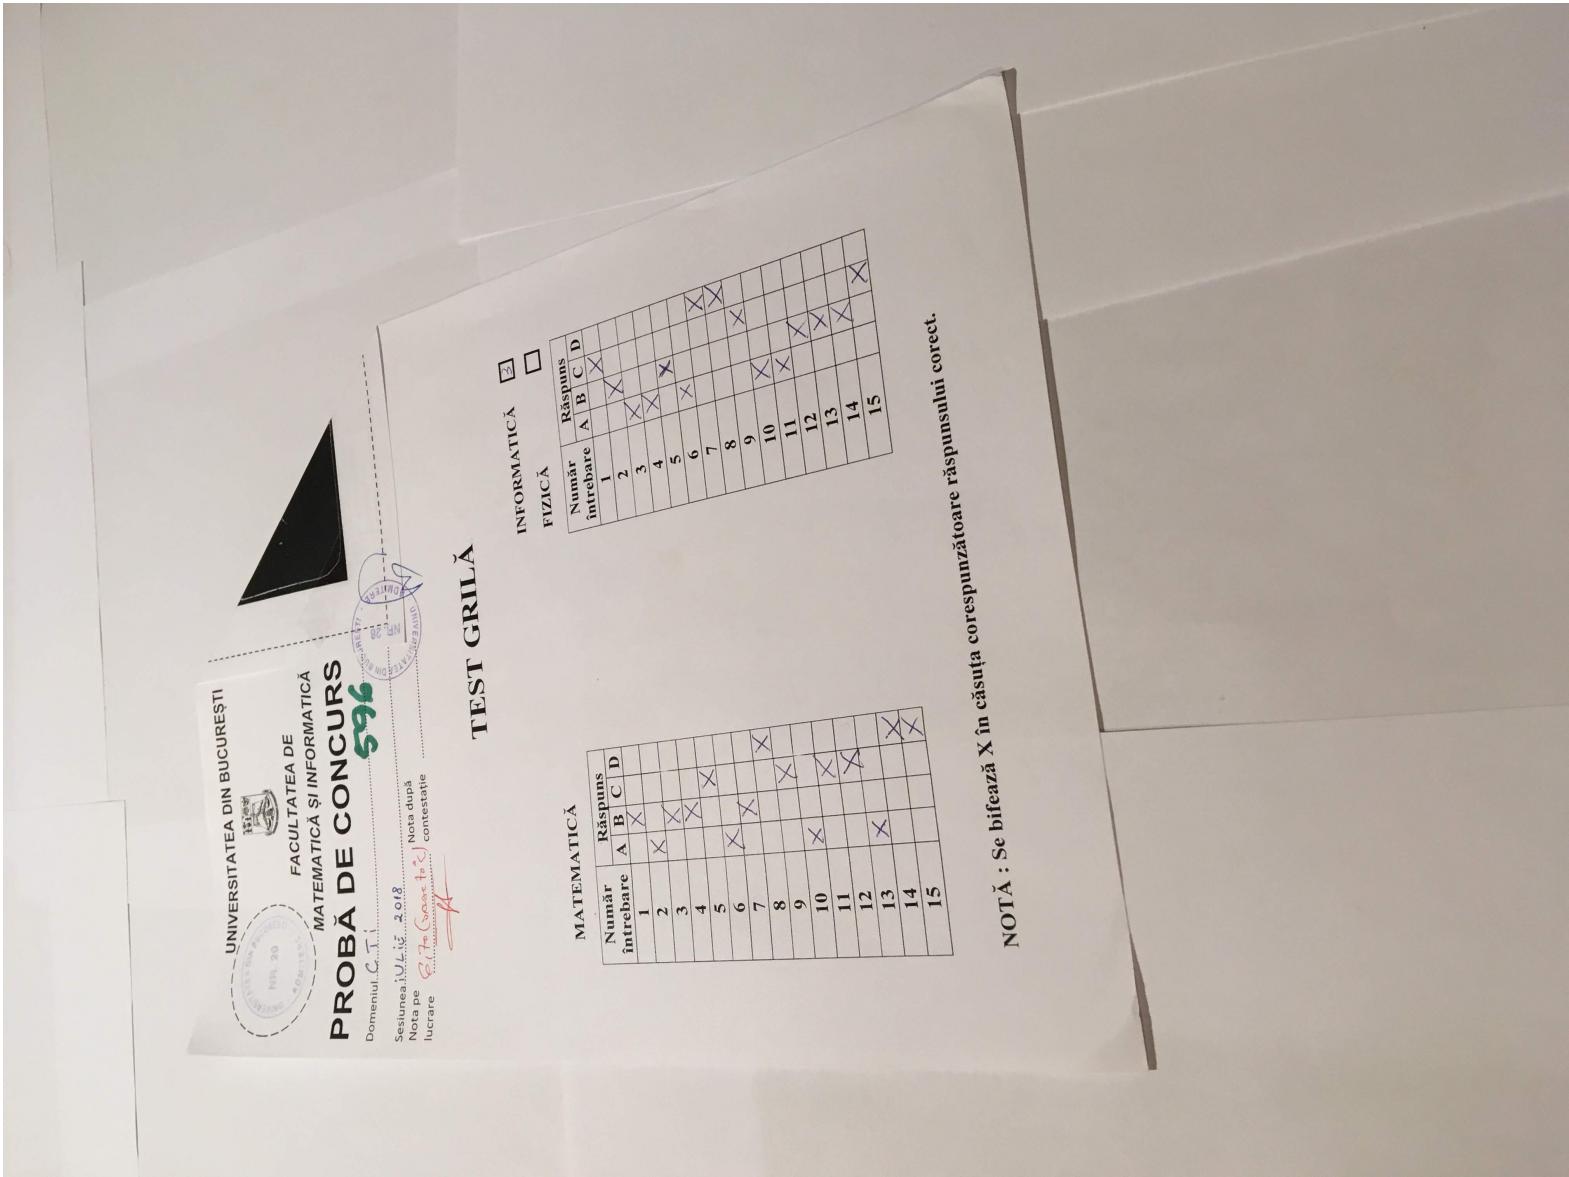
NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

image_2.txt	
1	3
2	B
3	A
4	B
5	C
6	A
7	B
8	D
9	C
10	A
11	C
12	C
13	A
14	D
15	D
16	C
17	B
18	A
19	A
20	B
21	A
22	D
23	D
24	C
25	A
26	A
27	B
28	B
29	B
30	C
	R 19

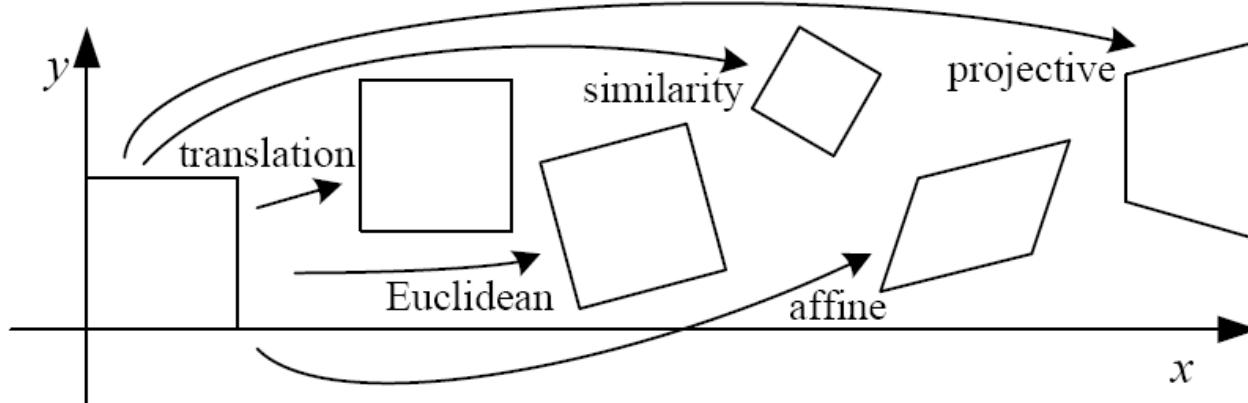
Rotation + scaling



Perspective

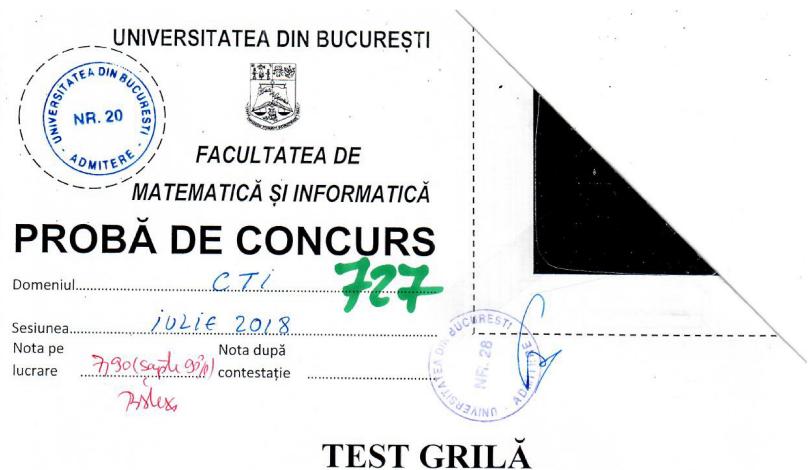


2D Geometric transformations



Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

BONUS: handwritten grade recognition



MATEMATICĂ

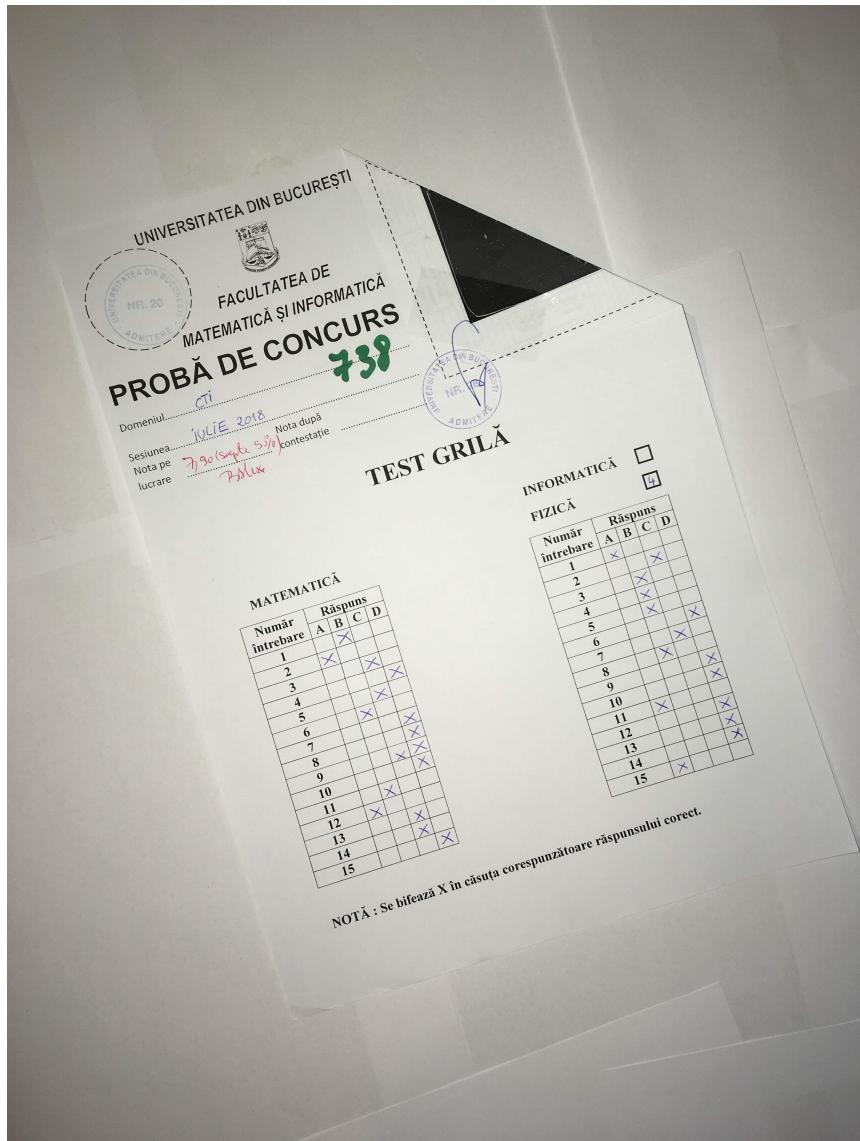
Număr	Răspuns	A	B	C	D
1		X			
2				X	X
3			X		
4		X			
5				X	X
6			X	X	
7		X			
8			X		
9				X	
10		X			
11			X	X	
12			X		
13		X			
14			X	X	
15			X		

INFORMATICĂ

FIZICĂ

Număr	Răspuns	A	B	C	D
1				X	
2			X		
3			X		
4			X		
5				X	
6				X	
7				X	
8				X	
9		X			
10			X		
11			X		
12			X		
13				X	
14		X			
15			X		

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.



Project 1: tinyurl.com/CV-2020-project1

Search ?

Dropbox > CV > 2019-2020 > Project1

Create Share Upload ...

<input type="checkbox"/>	Name	Modified
<input type="checkbox"/>	additional_data	---
<input type="checkbox"/>	ground-truth-correct-answers	---
<input type="checkbox"/>	images	---
<input type="checkbox"/>	output_format	---
<input type="checkbox"/>	test_data	---
<input type="checkbox"/>	project1.pdf	4/7/2020, 12:05 PM
<input type="checkbox"/>	README.txt	4/7/2020, 11:11 AM

Project 2

Video analysis of a snooker footage

This project worths 3.5 points but you can gain a bonus for a total of 4.5 points.

- **Task 1** - you receive an image of the snooker table seen from above. The task is to count the number of balls on the table and specify their color;
- **Task 2** - you receive a video containing a player making a shot. In this video the snooker table is seen from above. The task is to decide whether a ball was potted into a pocket and if so recognize the color of the potted ball and the pocket (one of the six pockets) where the ball was potted;
- **Task 3** - you receive a video containing a player making a shot. In this video the snooker table is seen from above. The task is to track the cue ball and another specified ball.
- (bonus) **Task 4** - you receive a much longer video than the previous ones (from Task 2 and Task 3). In this scenario, it may happen that several frames contain the snooker table seen from other viewpoints than above. The task here is output the score for the current footage based on the balls potted by each player. **1 point**

Task 1 – training data

Image 1

annotation



```
10
1 white
1 black
1 pink
1 blue
1 green
1 brown
1 yellow
3 red
```

Task 1 – training data

Image 8

annotation



8.txt

10
1 white
1 black
1 pink
1 blue
1 green
1 brown
1 yellow
3 red

Task 1 – training data

Image 18

annotation



Task 1 – training data

Image 23

annotation



```
16  
1 white  
0 black  
1 pink  
1 blue  
1 green  
1 brown  
1 yellow  
10 red
```

Task 1 – training data

Image 49

annotation



21
1 white
1 black
1 pink
1 blue
1 green
1 brown
1 yellow
14 red

Project 2 – Task 2

- **Task 2** - you receive a test set of 25 videos containing a player making a shot. In each of the 25 videos the snooker table is seen from above. The task is to decide whether a ball was potted into a pocket and if so recognize the color of the potted ball and the pocket (one of the six pockets) where the ball was potted. By solving correctly this task for a video you will earn 0.04 points as follows: if the correct answer is NO (no ball was potted) you will receive 0.04 points if your output is NO, else you will receive 0.02 points if your output is YES, 0.01 points for specifying the correct pocket (1 - top left corner, 2 - top right corner, 3 - bottom left corner, 4 - bottom right corner, 5 - middle left corner, 6 - middle right corner) and 0.01 points for specifying the correct color of the potted ball. In each of the 25 videos there will be a maximum of one ball being potted. For a test video your output should be similar with the one provided in the directory `training_data/Task2/`. (1 point)

Task 2 – training data

Video 1



annotation

A screenshot of a Mac OS X-style file viewer window titled "1.txt". The file contains the following text:

```
YES
2
red
```

Task 2 – training data

Video 2

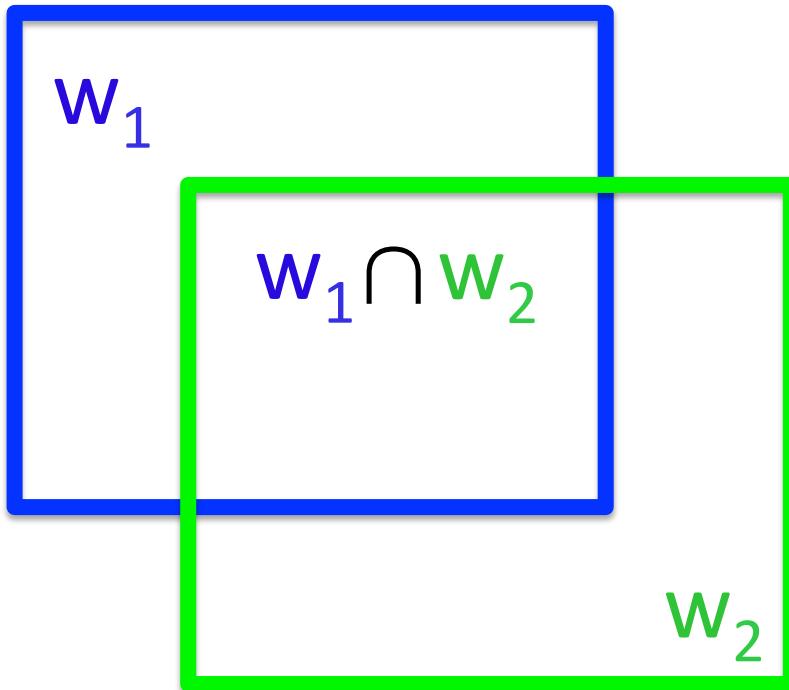
annotation



Project 2 – Task 3

- **Task 3** - you receive a test set of 25 videos containing a player making a shot. In each of the 25 videos the snooker table is seen from above. The task is to track the cue ball (the white ball) and another specified ball. The initial bounding boxes of the two balls to be tracked are provided for the first frame. By correctly solving the tracking problem for a video you will earn 0.04 points (0.02 points for each correctly tracked ball). In each video we will consider that your algorithm *correctly tracks a ball* if in more (greater or equal) than 80% of the frames your algorithm *correctly localizes the ball to be tracked*. We consider that your algorithm correctly localizes the ball to be tracked in a specific frame if the value of the IOU (intersection over union) between the window provided by your algorithm and the ground-truth window is more than 20%. For a test video your output should be similar with the one provided in the folder `training_data/Task3/` (1 point)

Intersection over union



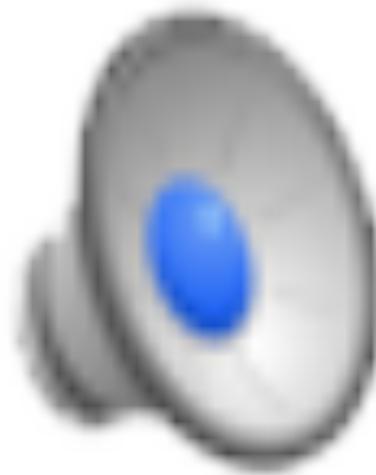
$$IOU(w_1, w_2) = \frac{area(w_1 \cap w_2)}{area(w_1 \cup w_2)}$$

$$IOU(w_1, w_2) = \frac{area(w_1 \cap w_2)}{area(w_1) + area(w_2) - area(w_1 \cap w_2)}$$

Task 3 – training data

Video 1

annotation



	annotation
196	-1 -1 -1 -1 -1
0	859 519 890 550
1	863 522 888 550
2	859 520 889 544
3	861 522 890 549
4	860 521 886 547
5	860 521 888 547
6	858 518 893 549
7	862 520 888 546
8	860 520 887 547
9	861 521 889 548
10	860 521 887 547
178	852 98 873 120
179	853 99 874 121
180	853 99 874 121
181	854 99 875 121
182	854 99 875 121
183	855 99 876 121
184	855 99 876 121
185	855 99 876 121
186	856 99 877 121
187	856 99 877 121
188	855 99 876 121
189	856 99 877 121
190	856 99 877 121

Task 3 – training data

Video 2



Project 2 – Task 4

- (bonus) **Task 4 - -** you receive a test set of 25 videos containing a player making a shot. Different than the previous tasks, in this videos the snooker table can be filmed from different viewpoints (not only from above). The task is to track the cue ball. The task here is much harder as you have different scales of the cue ball, changes in camera viewpoint, the initial bounding box of the white ball is not provided. By correctly solving the tracking problem for a video you will earn 0.04 points. The rules described at Task 3 apply here, meaning that your algorithm should correctly localizes the cue ball in more than 80% of the frames, at each frame correctly localization means that the window provided by your algorithm should have IOU more than 20% with respect to the ground-truth window. For a test video your output should be similar with the one provided in the folder [training_data/Task4/](#) (1 point)

Task 4 – training data



Task 4 – training data



Project 2: tinyurl.com/CV-2020-project2

  Search



Dropbox > CV > 2019-2020 > Project2

Overview

Click here to describe this folder and turn it into a Space

[Show examples](#)

Create 

Share

Upload 

...



Name 

Modified



 format_output



--



 test_data



--



 test_data_ground_truth



--



 training_data



--



 project2.pdf



6/8/2020, 11:14 PM

Project 3 – Image retrieval using bag of visual words model

- build a database with 50 classes
- each class contains a small number (e. g. 10) images with the same scene (landmark)
- the goal: design an image retrieval system that takes as input a query image and outputs most N ($N > 10$) similar in content images from the database
- ingredients:
 - bag-of-visual-word model
 - SIFT descriptors
 - geometric verification with RANSAC
 - inverted file index – use TF-IDF

Class 1



Class 2



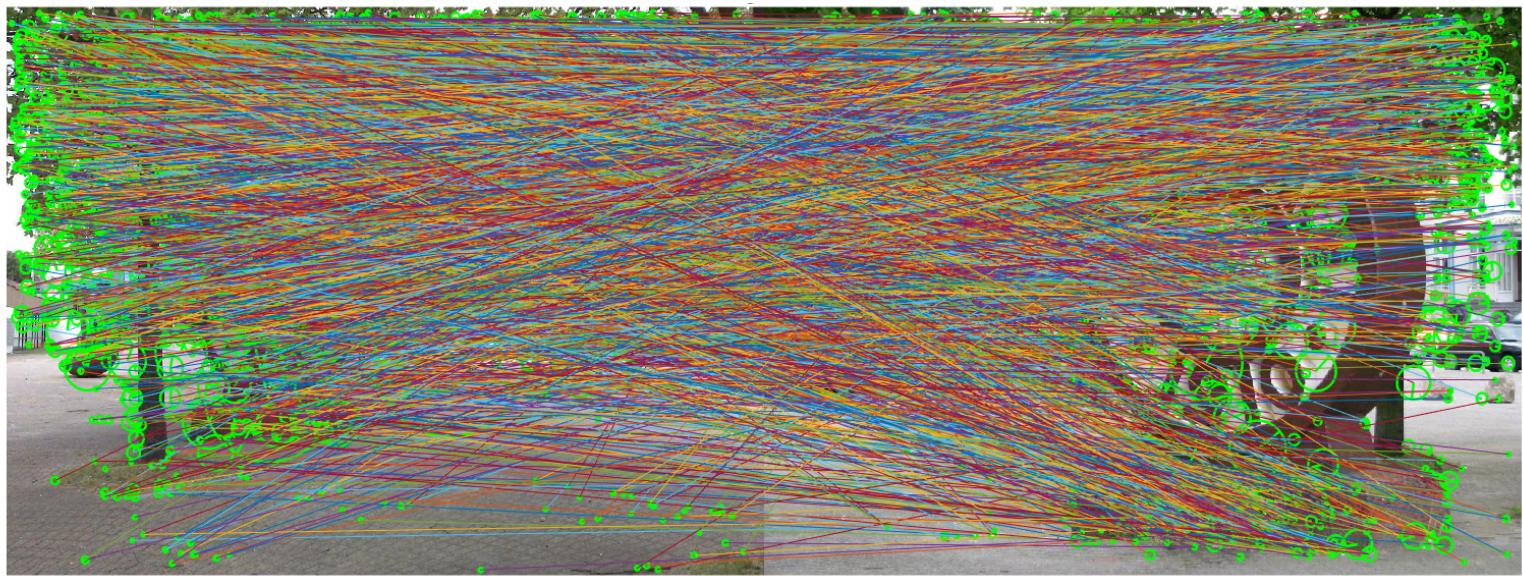
Class 45



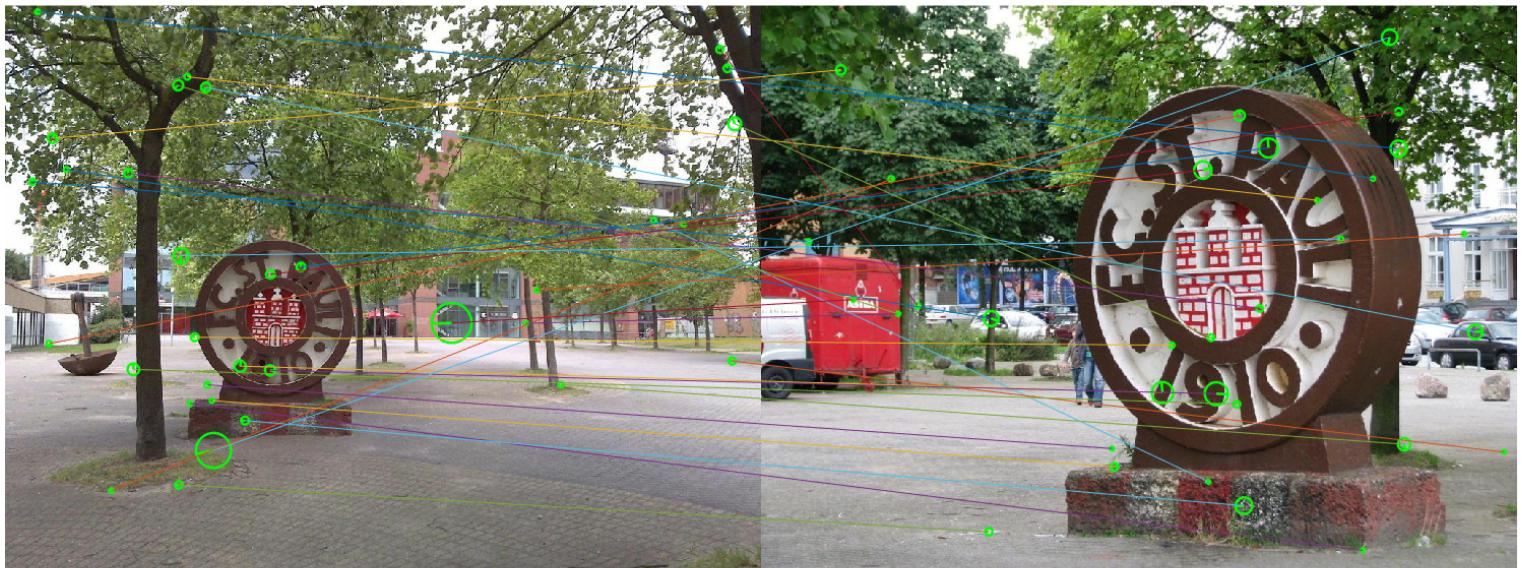
SIFT extraction



SIFT matching using 1stNN



SIFT matching using the ratio $1^{\text{st}}\text{NN}/2^{\text{nd}}\text{NN}$



SIFT matching using the ratio $1^{\text{st}}\text{NN}/2^{\text{nd}}\text{NN}$ + geometric verification with RANSAC

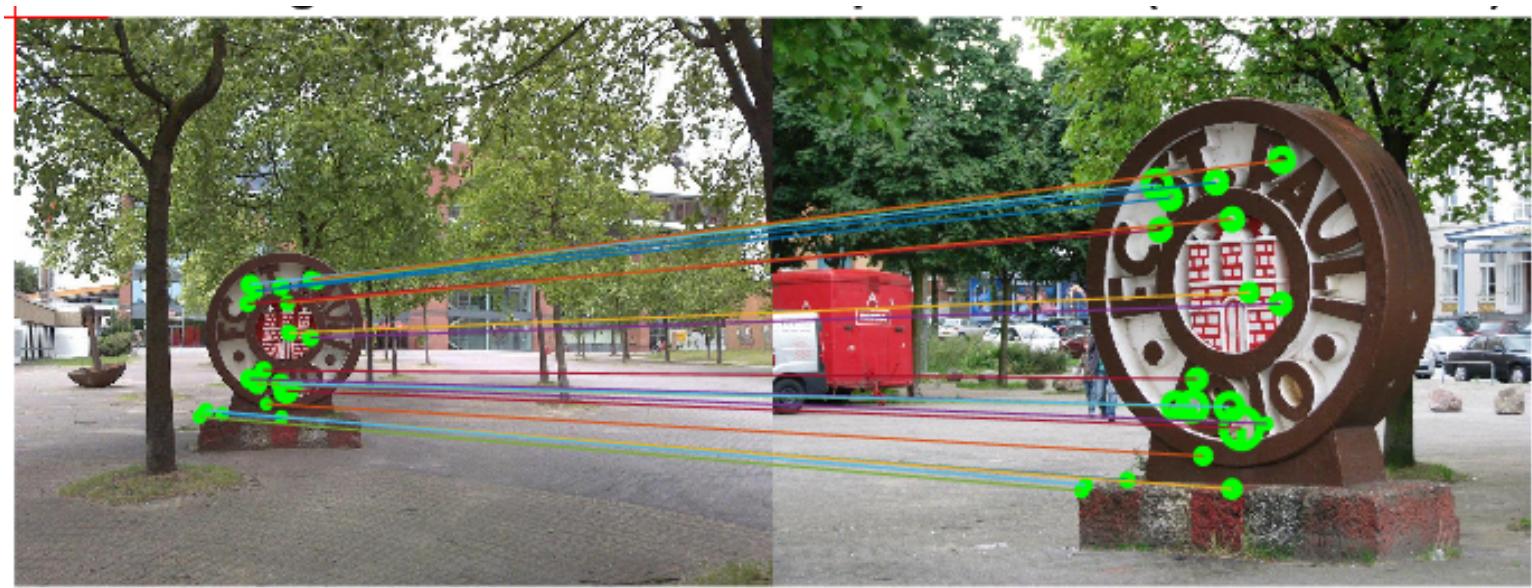


Image retrieval based on inverted file index



Query image



Image retrieval based on inverted file index + geometric verification



Query image

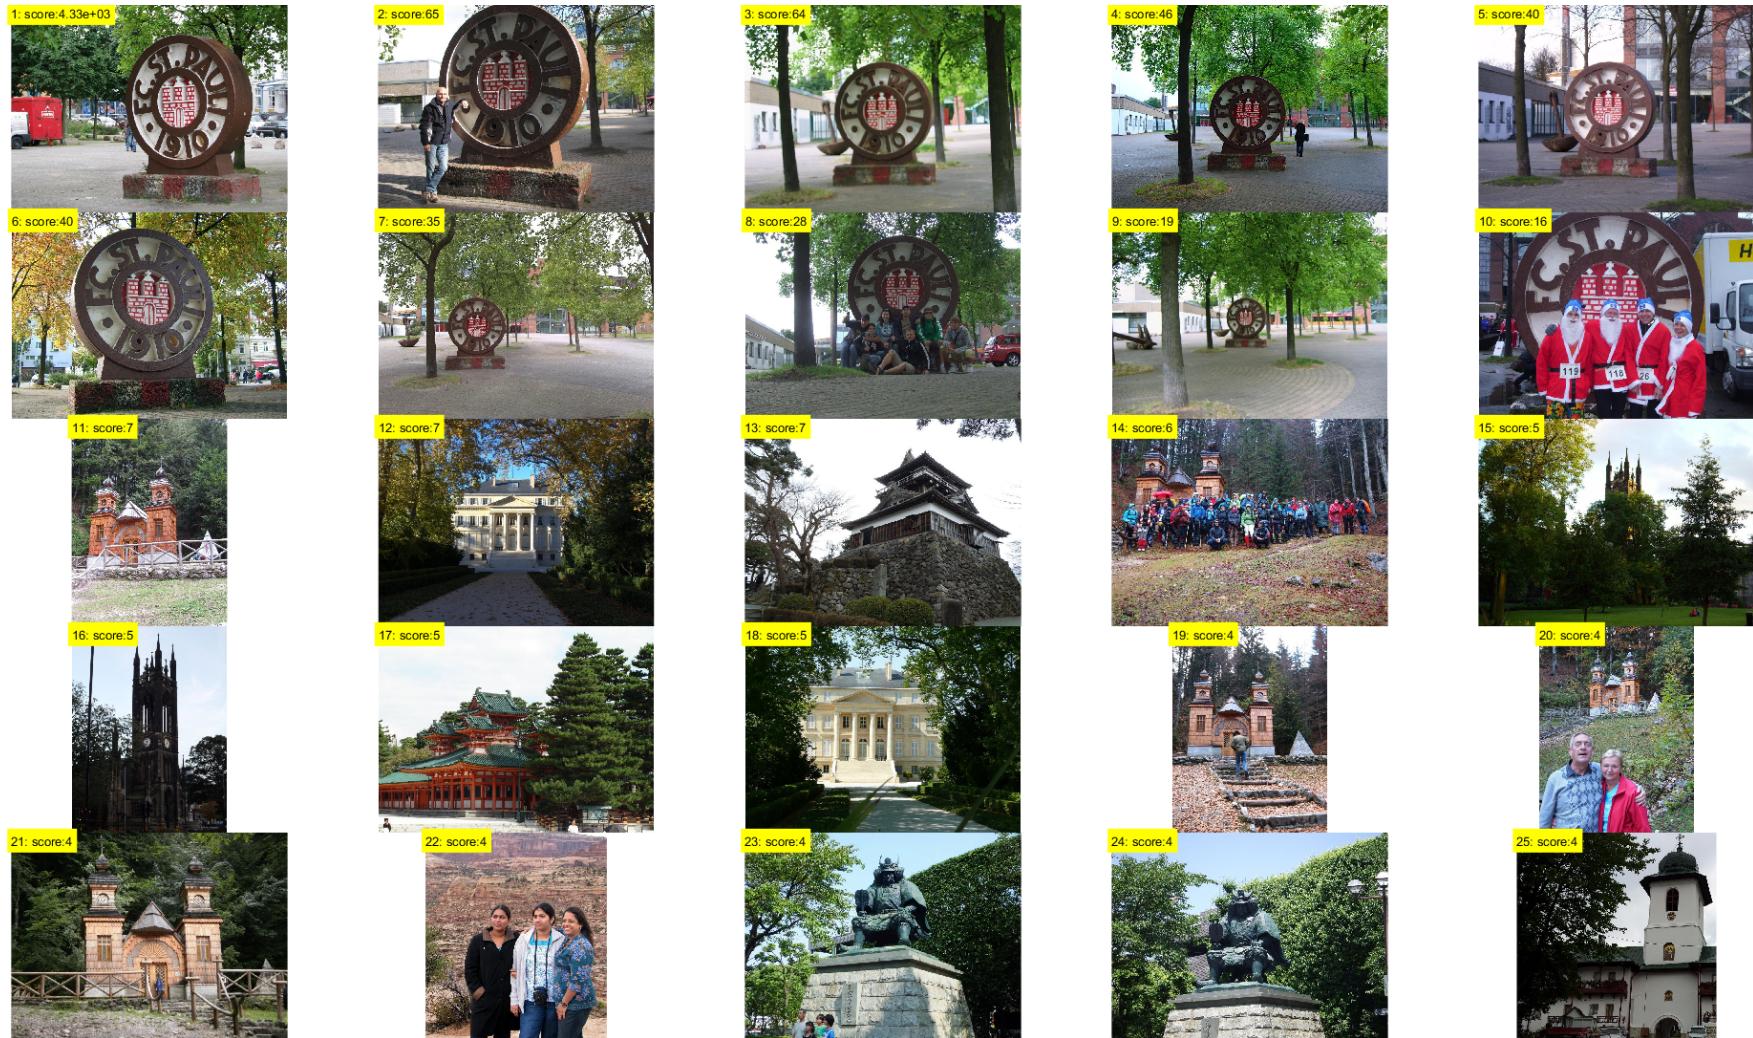
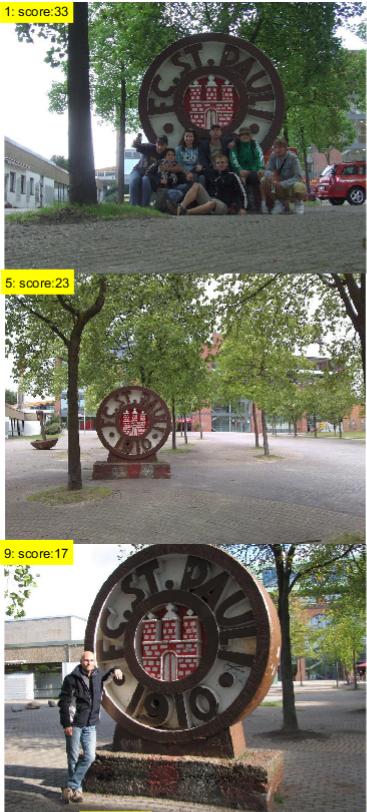


Image retrieval based on inverted file index + geometric verification



Query image



Project 3: tinyurl.com/CV-2020-project3

Search ?

Dropbox > CV > 2019-2020 > Project3

Overview

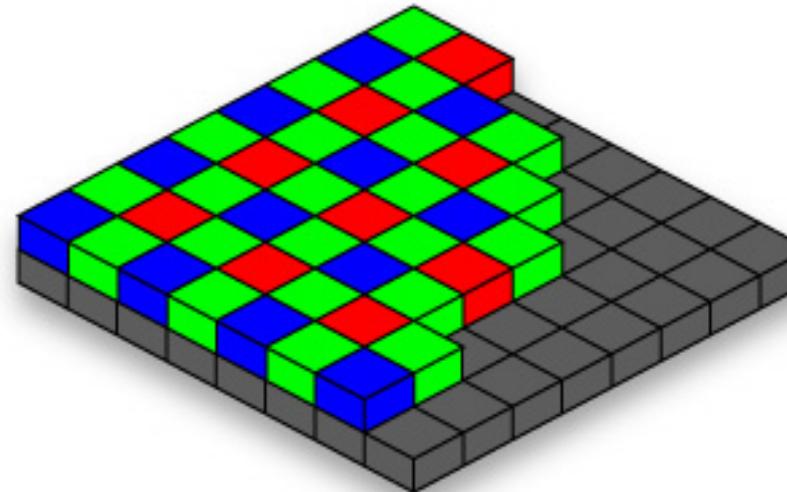
Click here to describe this folder and turn it into a Space Show examples

Create Share Upload ...

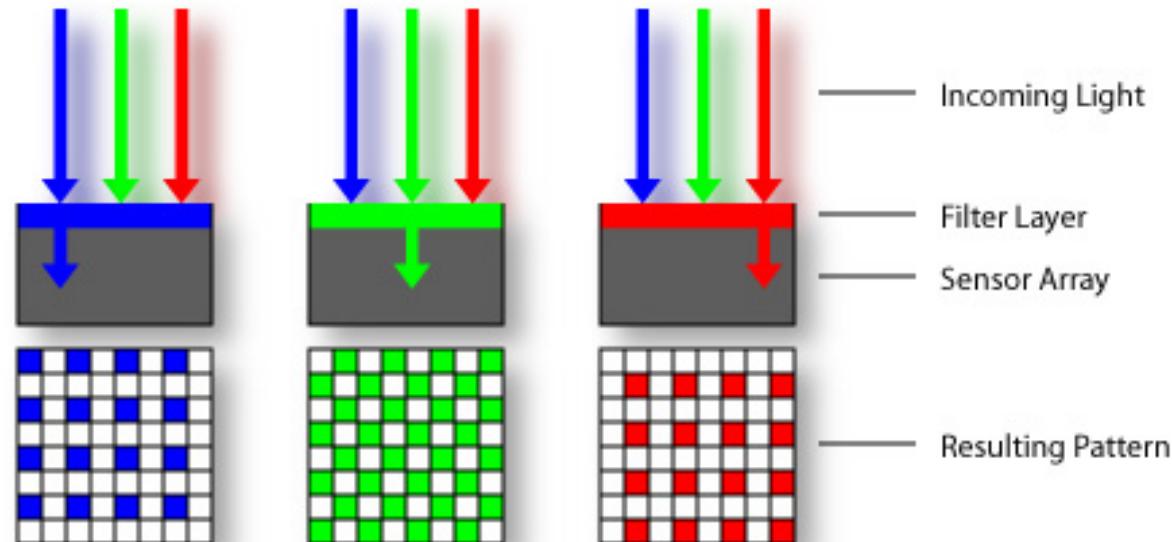
<input type="checkbox"/>	Name ↑	Modified
<input type="checkbox"/>	data	☆ --
<input type="checkbox"/>	BOVW.pdf	☆ 9/8/2020, 11:10 PM

First laboratory class

Color Sensing: Bayer Grid

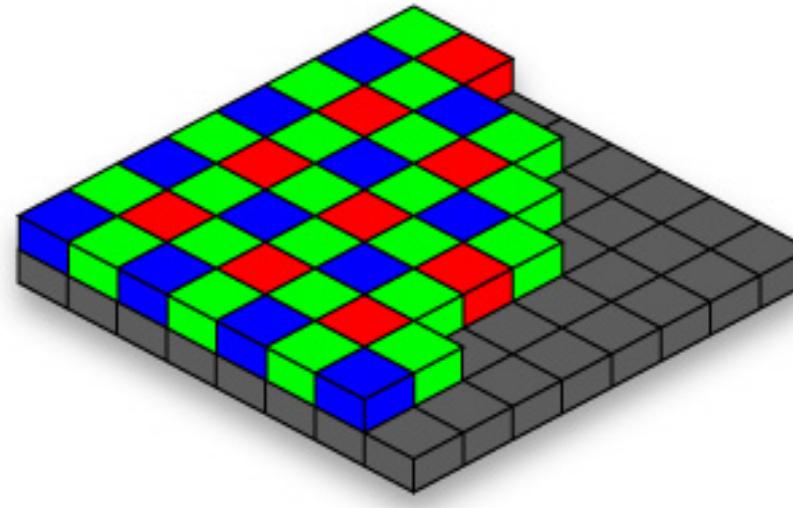
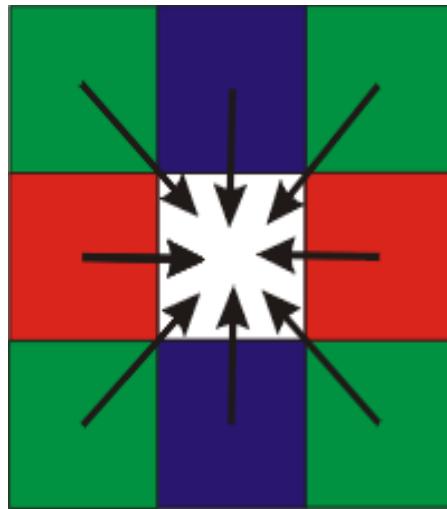


Estimate RGB at each cell from neighboring values



First laboratory class

Color Sensing: Bayer Grid



Estimate RGB at each cell from
neighboring values

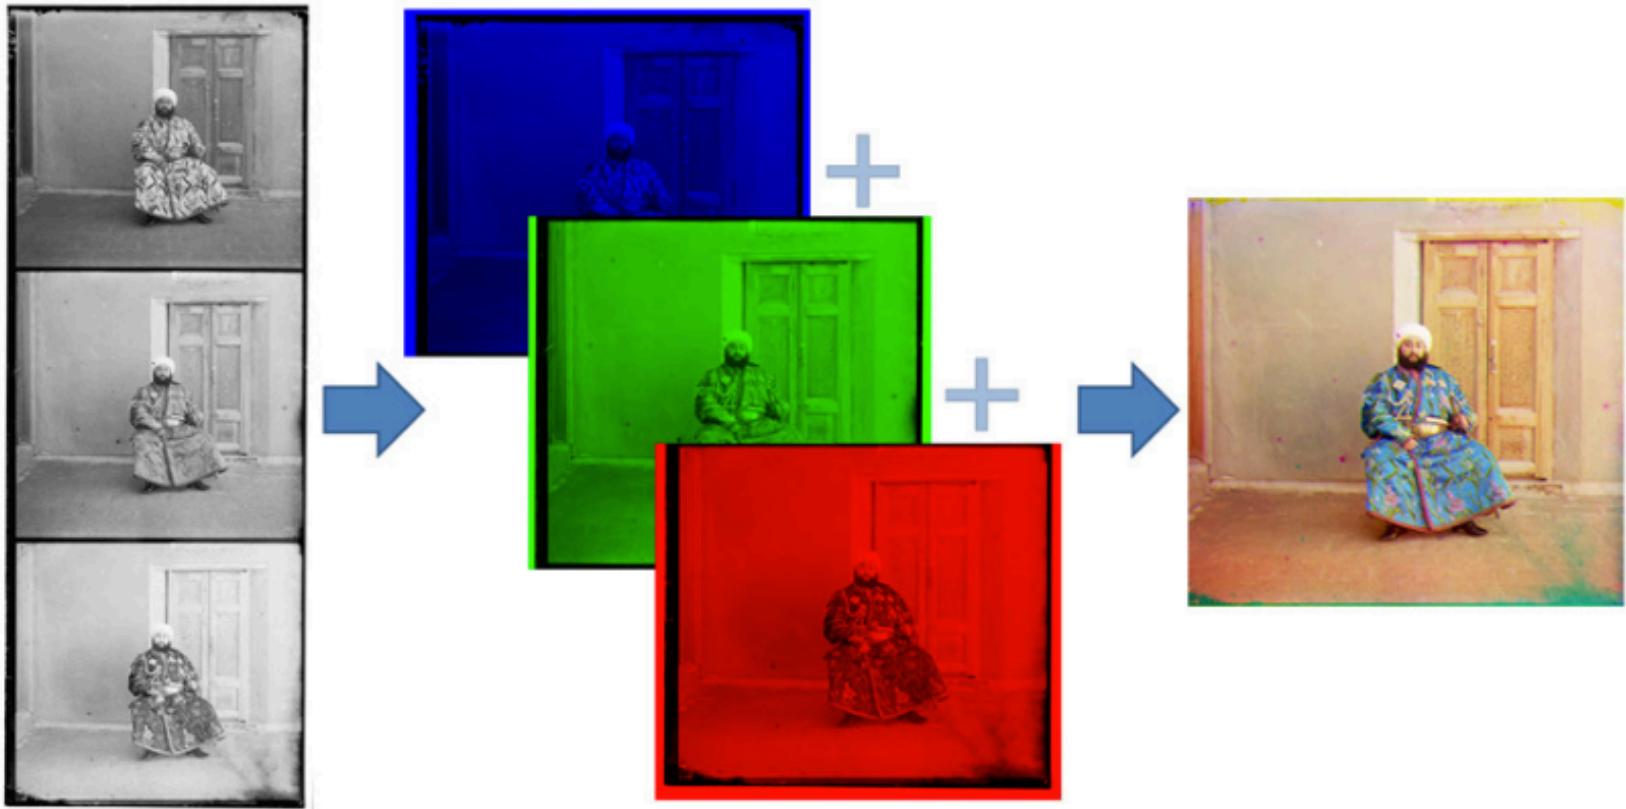
?	Red	?	Red	?	Red
?	?	?	?	?	?
?	Red	?	Red	?	Red
?	?	?	?	?	?
?	Red	?	Red	?	Red
?	?	?	?	?	?

?	Green	?	Green	?	Green
?	?	?	?	?	?
?	Green	?	Green	?	Green
?	?	?	?	?	?
?	Green	?	Green	?	Green
?	?	?	?	?	?

?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?

First laboratory class

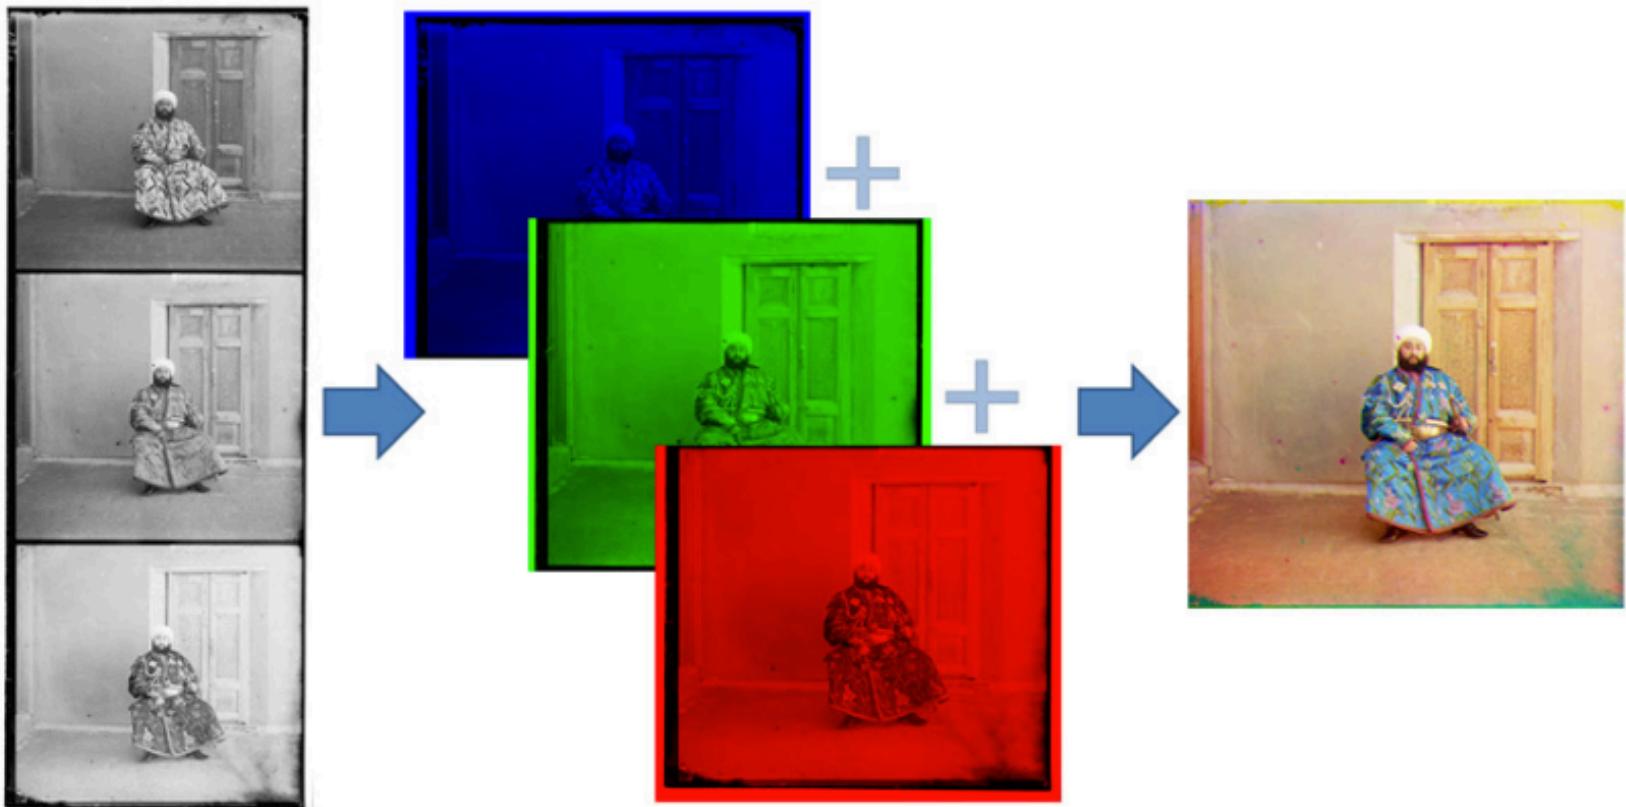
Your first application in Computer Vision



Record three exposures of every scene onto a glass plate using a red, a green, and a blue filter and then project the monochrome pictures with correctly colored light to reproduce the color image

First laboratory class

Your first application in Computer Vision



Your program will take a **glass plate image** as input and produce a **single color image** as output. The program should *divide the image into three equal parts* and **align the second and the third parts (G and R) to the first (B)**.

Not aligned



Aligned

