

# Computer Vision

Bogdan Alexe

[bogdan.alexe@fmi.unibuc.ro](mailto:bogdan.alexe@fmi.unibuc.ro)

University of Bucharest, 2<sup>nd</sup> semester, 2020-2021

# Course structure

## 1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

## 2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

## 3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

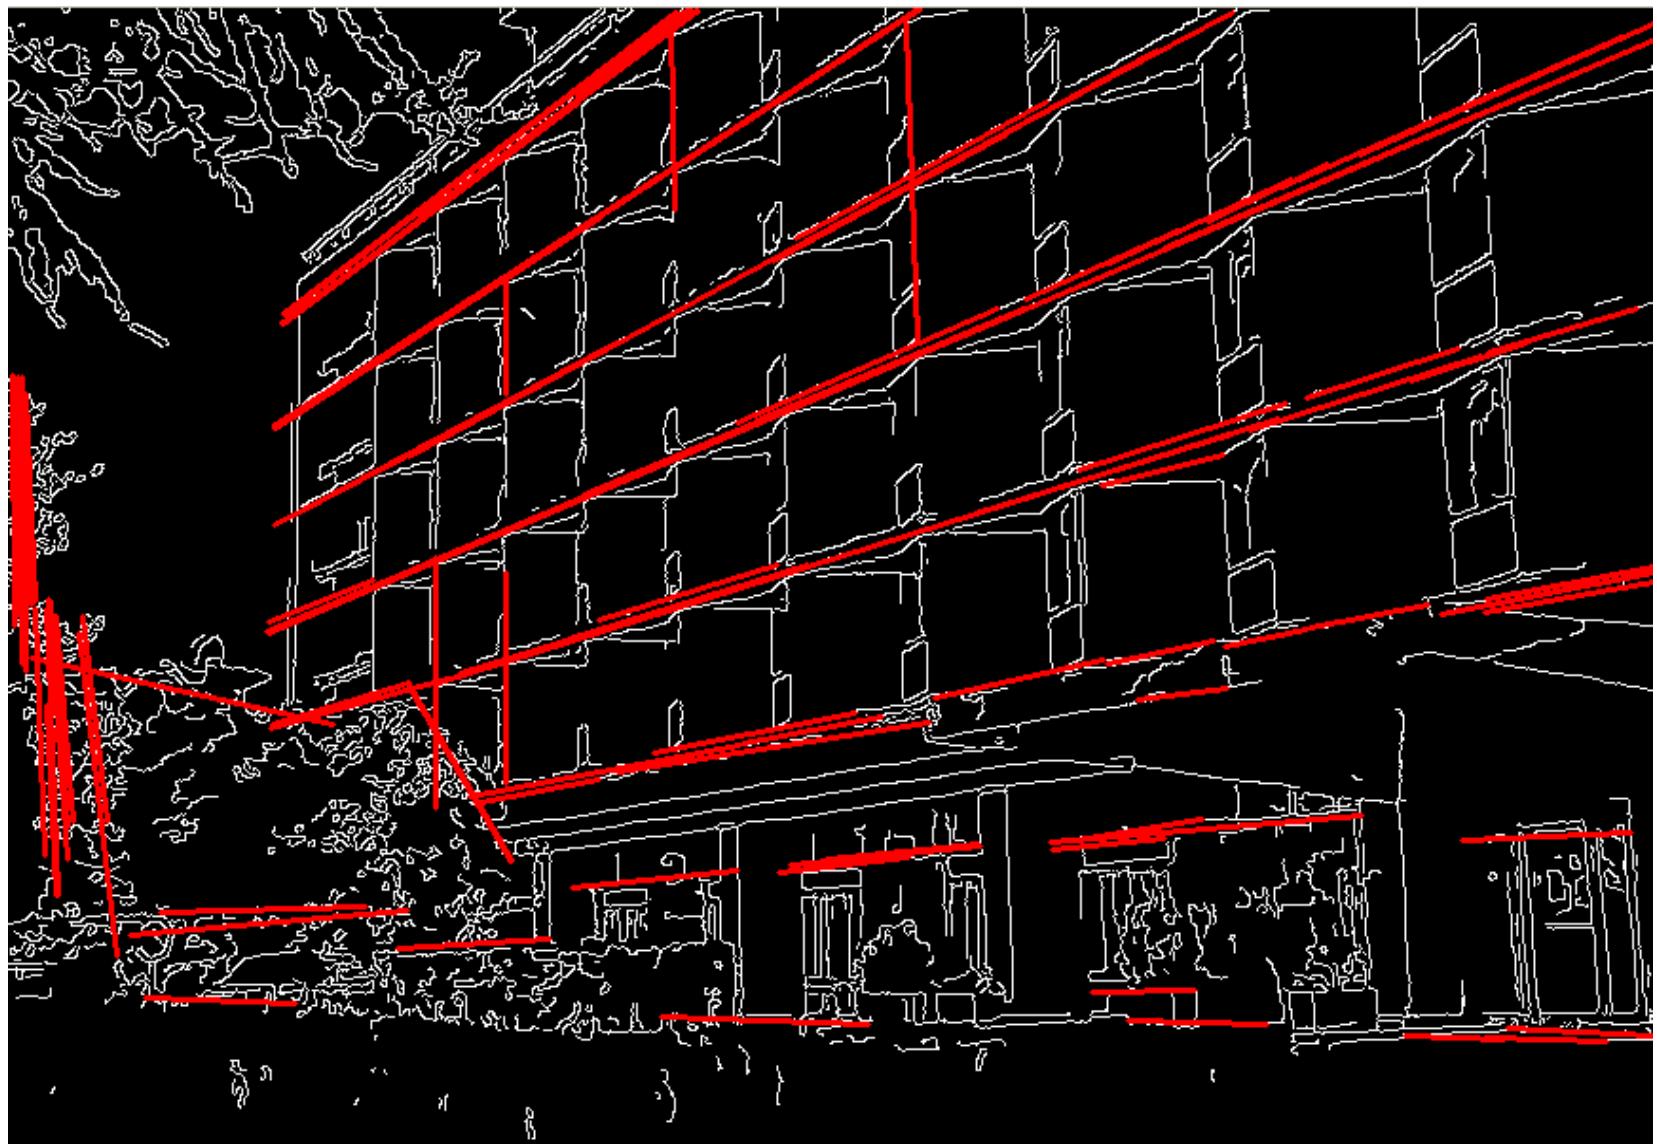
## 4. Object Recognition: high – level vision

Object classification, object detection, part based models, bovw models

## 5. Video understanding

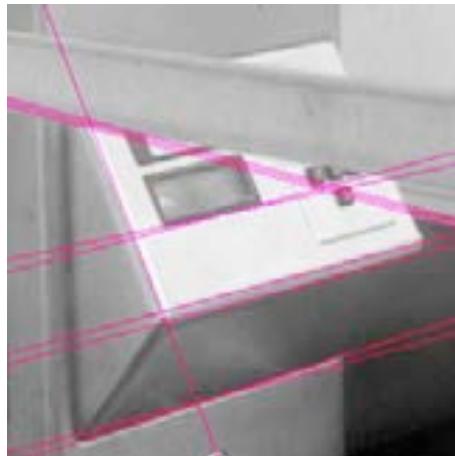
Object tracking, background subtraction, motion descriptors, optical flow

# Fitting



# Fitting

- Choose a *parametric model* to represent a set of features



simple model: lines



simple model: circles



complicated model: car



# Lab 2: Line detection - application

**PROBĂ DE CONCURS**

Domeniu... *CTi* **730**

Sesiunea... *IULIE 2018*

Nota pe lucrare *4,90 (patru zeci și nouă)* Nota după contestație *PDLEX*

UNIVERSITATEA DIN BUCURESTI  
NF. 78 ADH

**TEST GRILĂ**

**MATEMATICĂ**

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2	X			
3		X		
4	X			
5		X		
6		X		
7		X		
8				X
9			X	
10	X			
11		X		
12		X		
13			X	
14			X	
15			X	

**INFORMATICĂ**

**FIZICĂ**  1

Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2			X	
3			X	
4	X			
5			X	
6	X			
7	X			
8			X	
9			X	
10			X	
11		X		
12				X
13			X	
14	X			
15			X	

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

# Lab 2: Line detection - application

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				
9				X
10	X			X
11		X		
12		X		
13			X	
14				X
15			X	

	Număr întrebare	Răspuns			
		A	B	C	D
1			X		
2			X		
3				X	
4		X			
5			X		
6				X	
7		X			
8				X	
9				X	
10		X			
11			X		
12			X		
13				X	
14				X	
15				X	

# Lab 2: Line detection - application

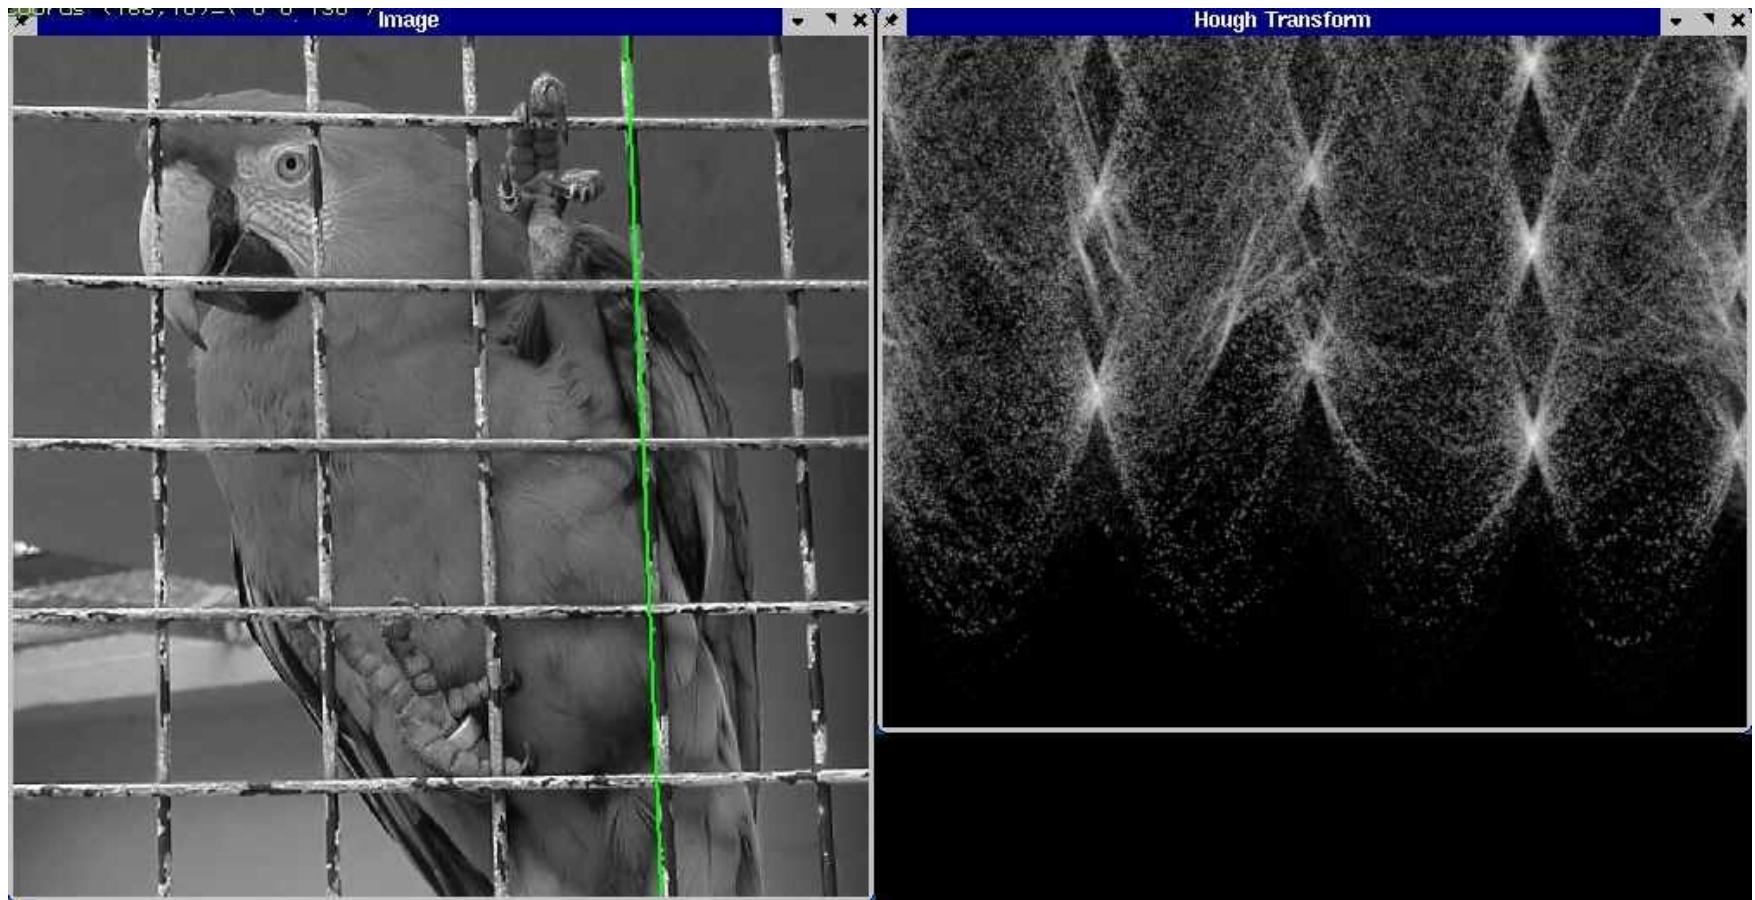
Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	X
7		X		
8				X
9				X
10	X			X
11		X		
12		X		
13			X	
14				X
15				X

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4		X		
5			X	
6			X	
7		X		
8				X
9				X
10		X		
11			X	
12		X		
13			X	
14				X
15				X

# Fitting: challenges

- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform

# Fitting: The Hough transform

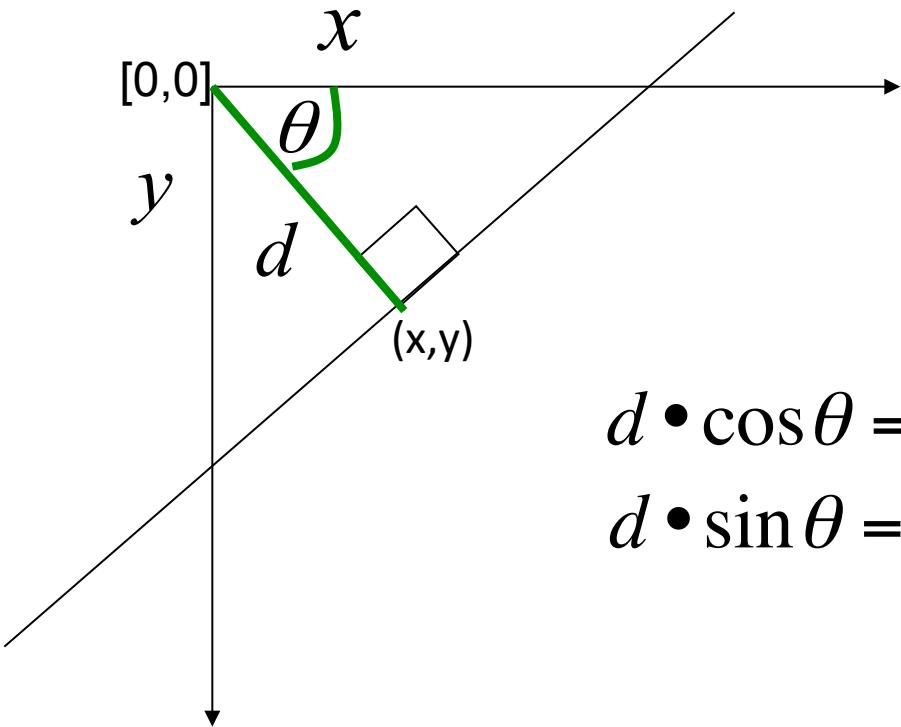


# Polar representation

Problems with the (m,b) space:

- Unbounded parameter domains
- Vertical lines require infinite m

Alternative: **polar representation**



$d$ : distance from origin to line  
(length of the perpendicular)

$\theta$ : angle between the  
perpendicular and Ox axis

$$d \cdot \cos \theta = x$$

$$d \cdot \sin \theta = y$$

$$d \cdot (\cos \theta)^2 = x \cdot \cos \theta$$

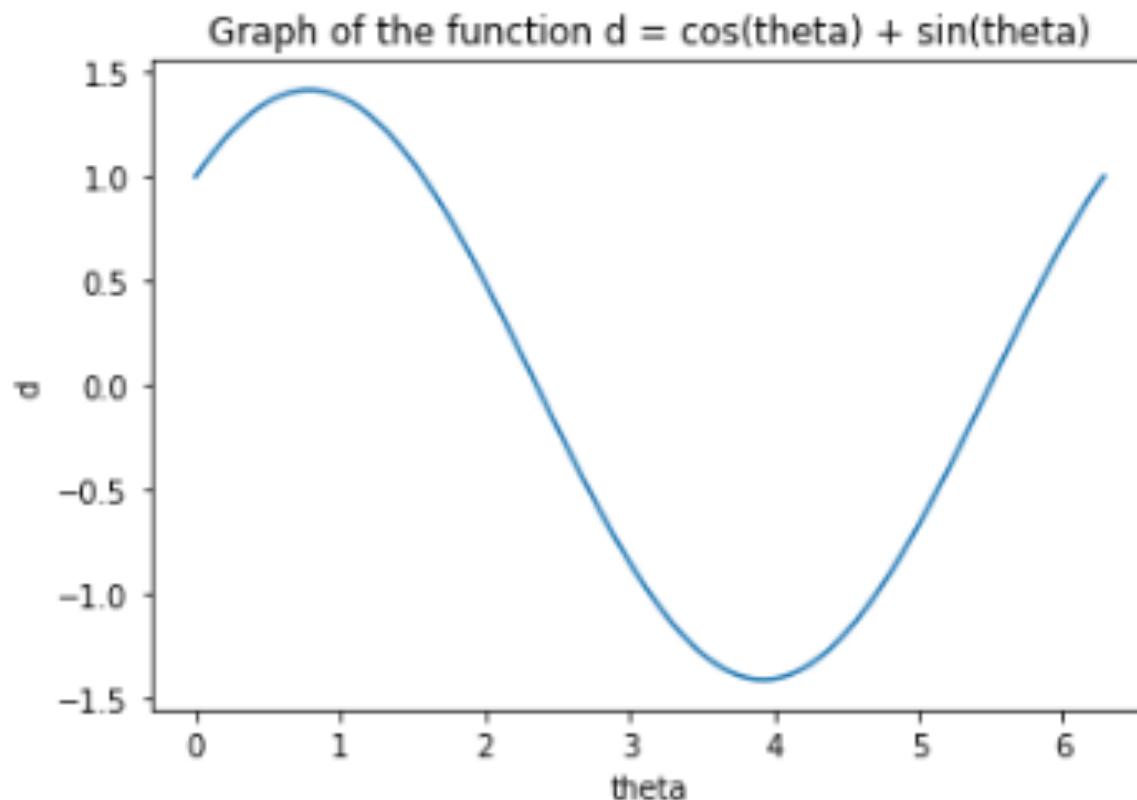
$$d \cdot (\sin \theta)^2 = y \cdot \sin \theta$$

---

$$d = x \cdot \cos \theta + y \cdot \sin \theta$$

Point in the image space  $\rightarrow$  sinusoid curve in the Hough space

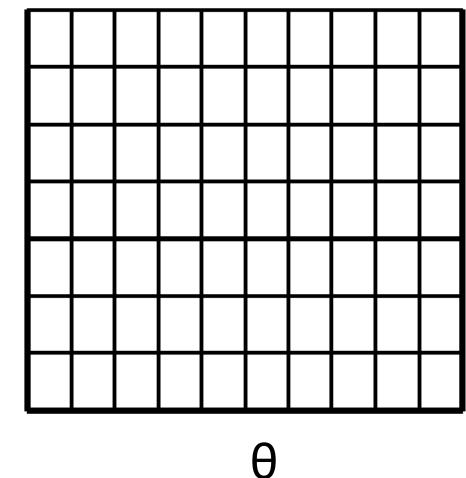
```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 x, y = 1,1
4 theta = np.linspace(0,2*np.pi,100)
5 d = x * np.cos(theta) + y * np.sin(theta)
6 plt.figure,
7 plt.plot(theta, d)
8 plt.xlabel('theta')
9 plt.ylabel('d')
10 plt.title('Graph of the function d = cos(theta) + sin(theta)')
11 plt.show()
```



# Hough transform-algorithm outline

- Initialize accumulator  $H$  to all zeros
- For each feature point  $(x, y)$  in the image
  - For  $\theta = 0$  to  $180$ 
    - $d = x \cos \theta + y \sin \theta$
    - $H(\theta, d) = H(\theta, d) + 1$
  - end
- Find the value(s) of  $(\theta, d)$  where  $H(\theta, d)$  is a local maximum
  - The detected line in the image is given by
$$d = x \cos \theta + y \sin \theta$$

$H$ : accumulator array (votes)



$\theta$

end

end

- Find the value(s) of  $(\theta, d)$  where  $H(\theta, d)$  is a local maximum

- The detected line in the image is given by

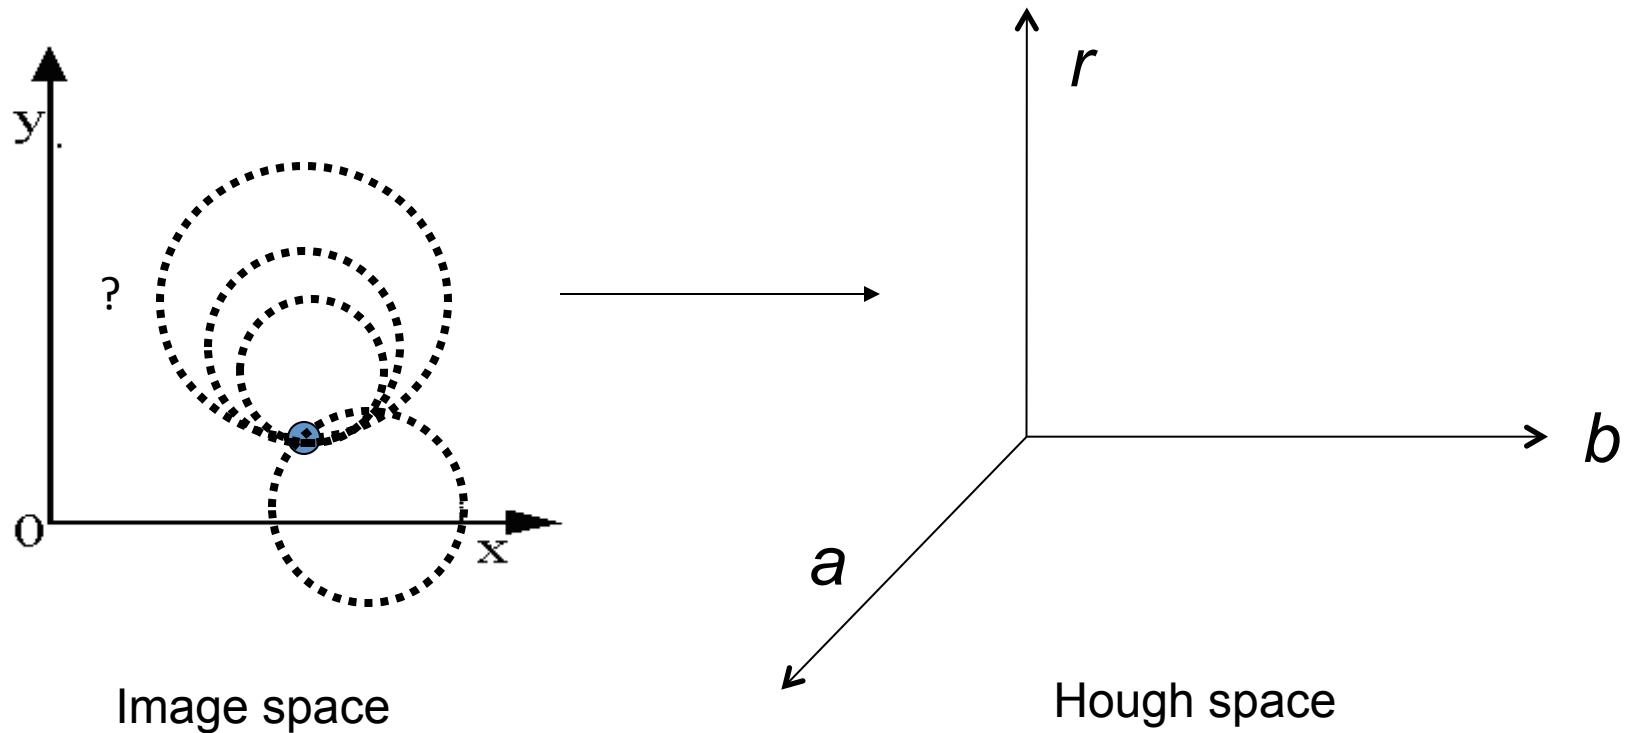
$$d = x \cos \theta + y \sin \theta$$

# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$ . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For unknown radius  $r$



# Hough transform for circles

- Circle: center  $(a, b)$  and radius  $r$ . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For unknown radius  $r$

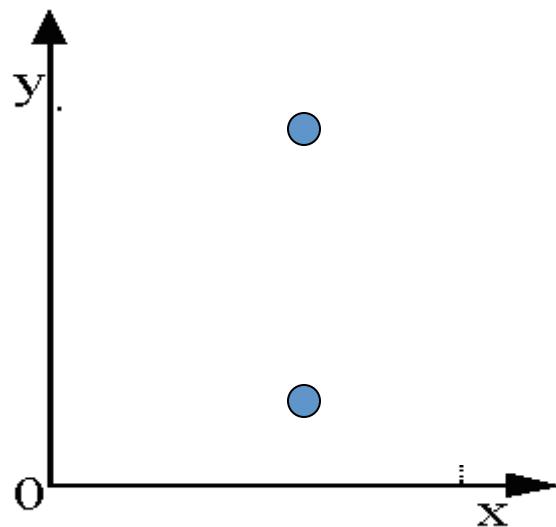
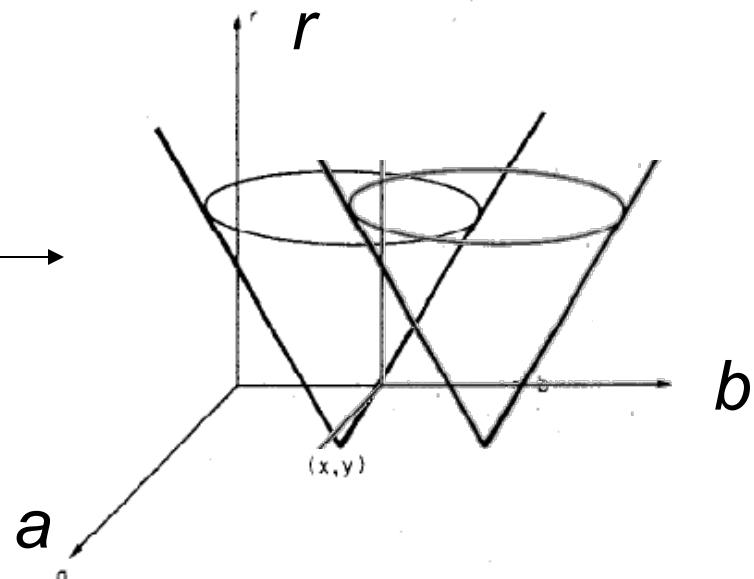


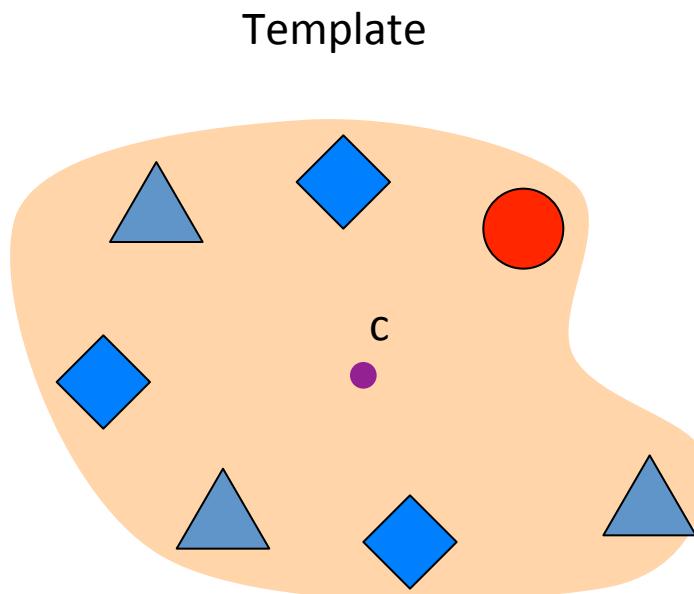
Image space



Hough space

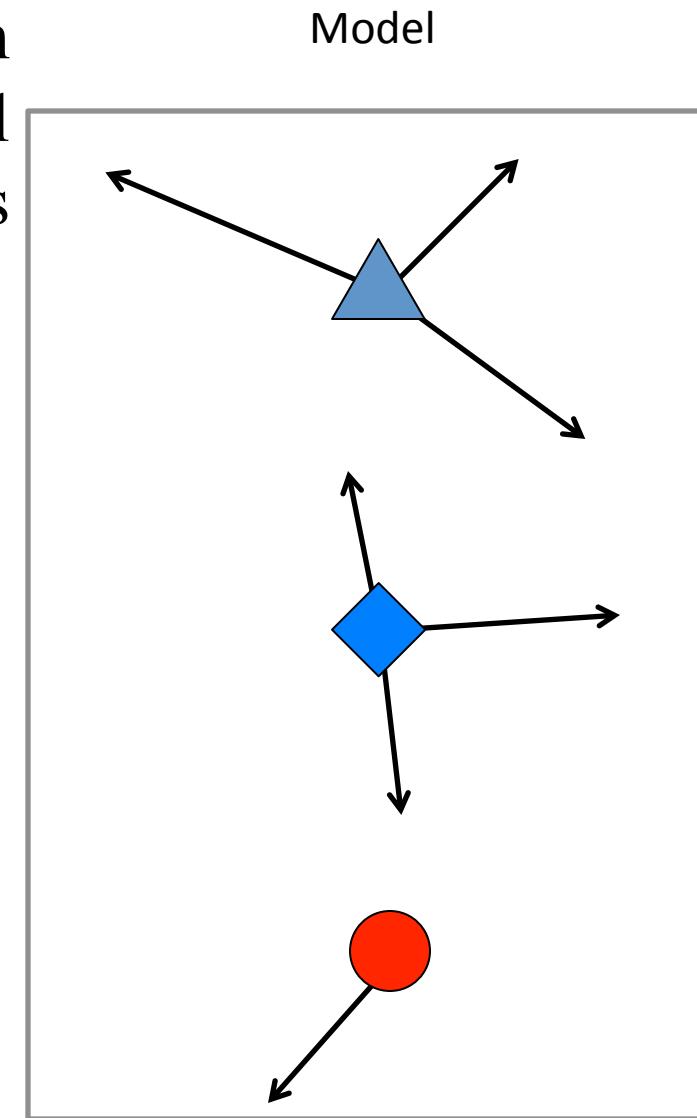
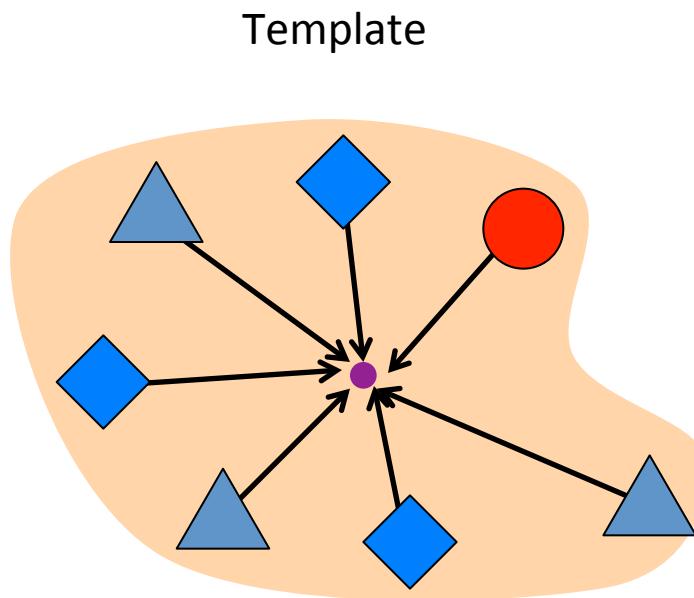
# Generalized Hough transform

- We want to find a template defined by its reference point (center) and several distinct types of landmark points in stable spatial configuration



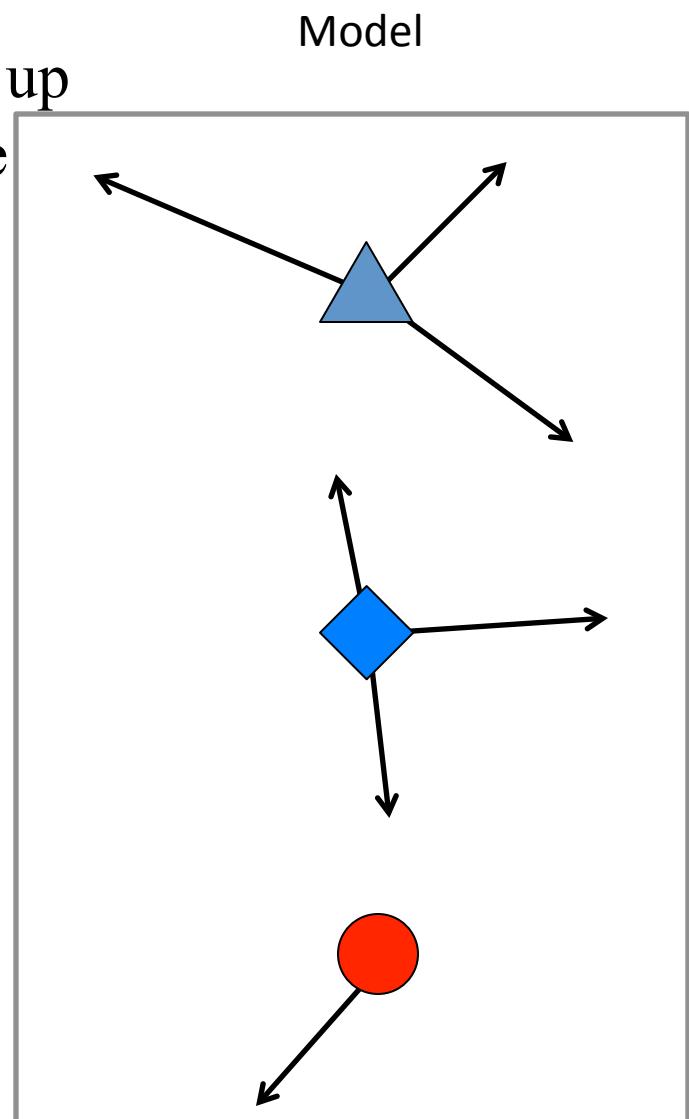
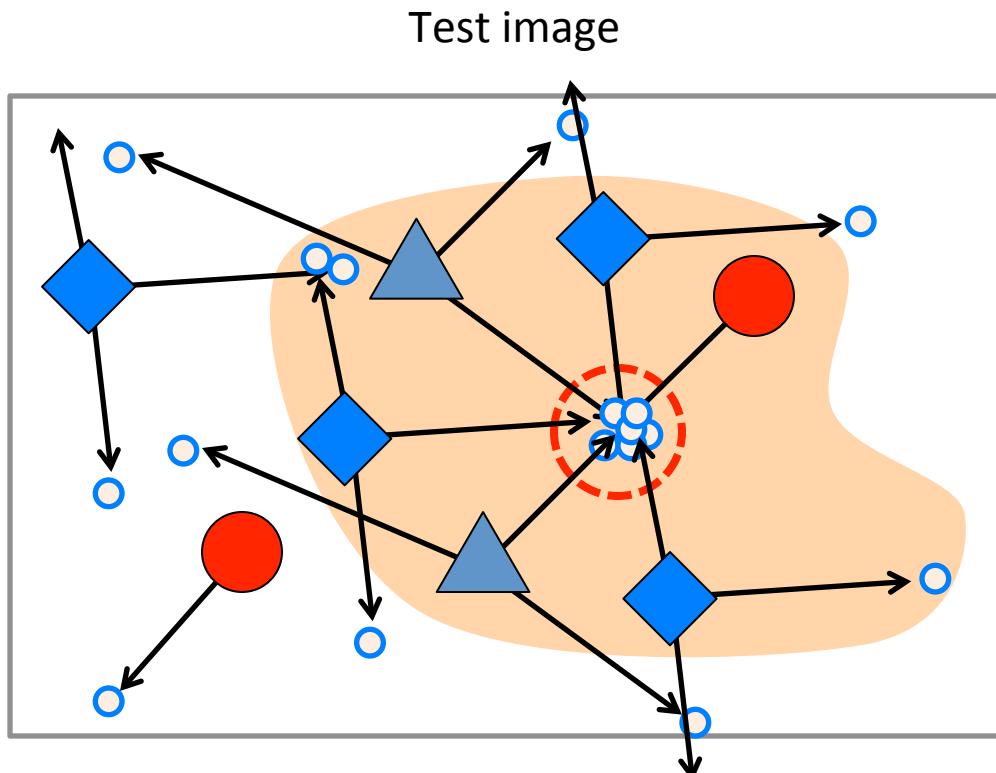
# Generalized Hough transform

- Template representation: for each type of landmark point, store all possible displacement vectors towards the center



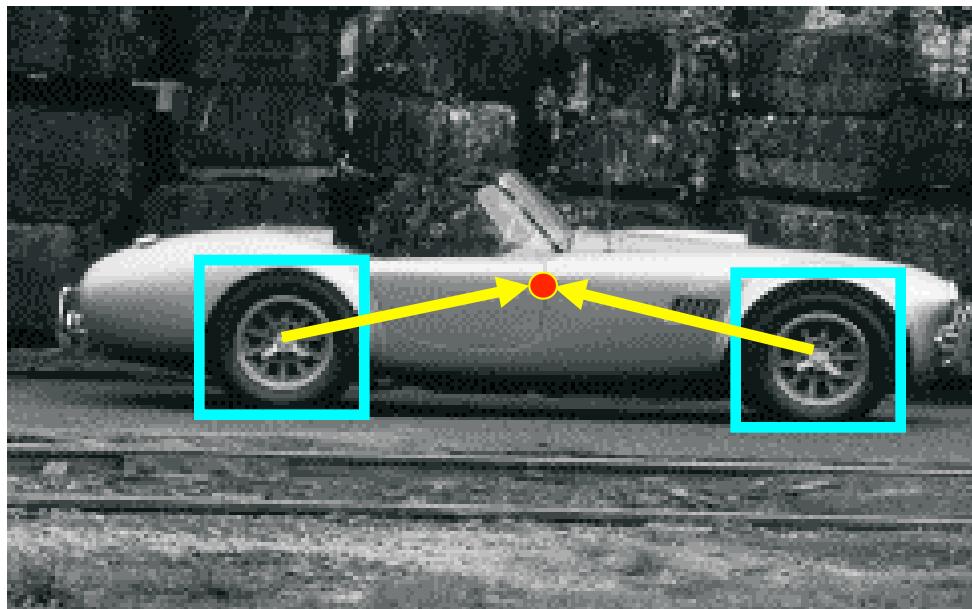
# Generalized Hough transform

- Detecting the template:
  - For each feature in a new image, look up that feature type in the model and vote for the possible center locations associated with that type in the model

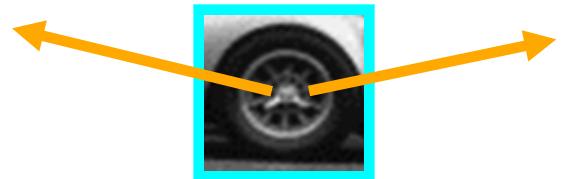


# Application in recognition

- Index displacements by “visual codeword”



training image



visual codeword with  
displacement vectors

# Application in recognition

- Index displacements by “visual codeword”



test image

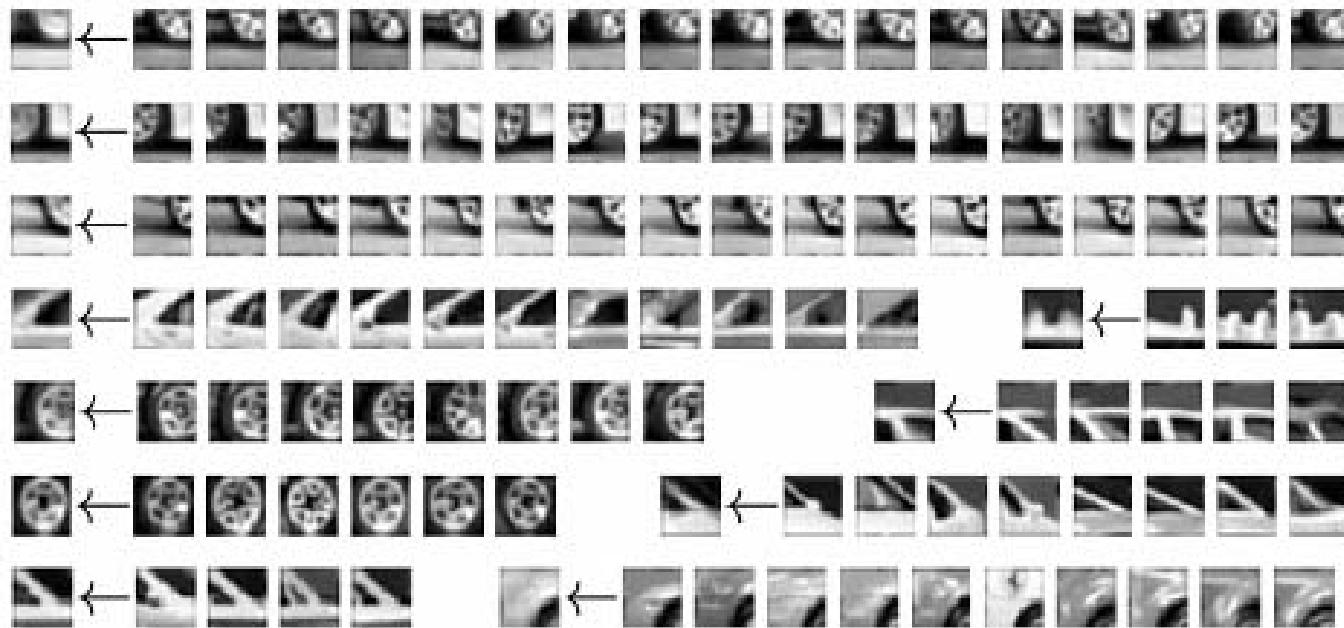
B. Leibe, A. Leonardis, and B. Schiele,

Combined Object Categorization and Segmentation with an Implicit Shape Model, ECCV

Workshop on Statistical Learning in Computer Vision 2004

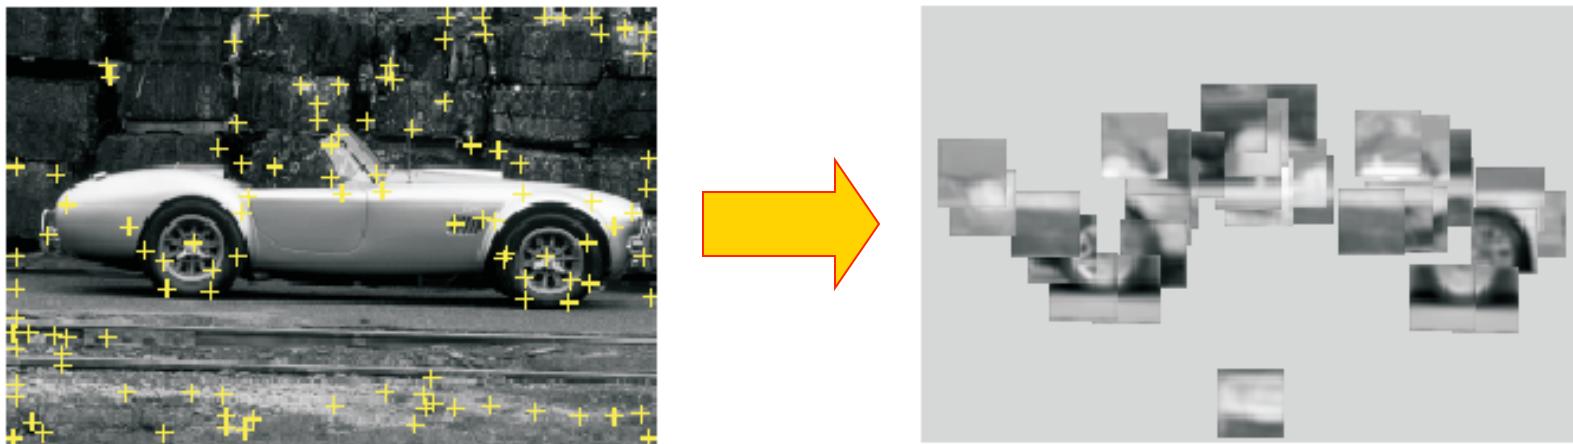
# Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering



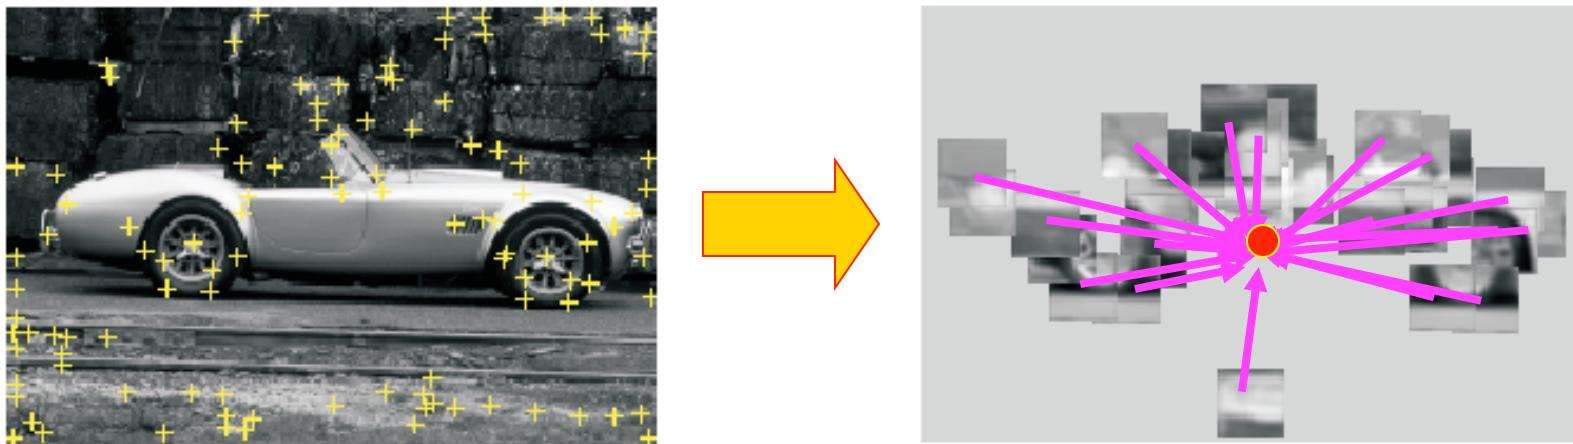
# Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry



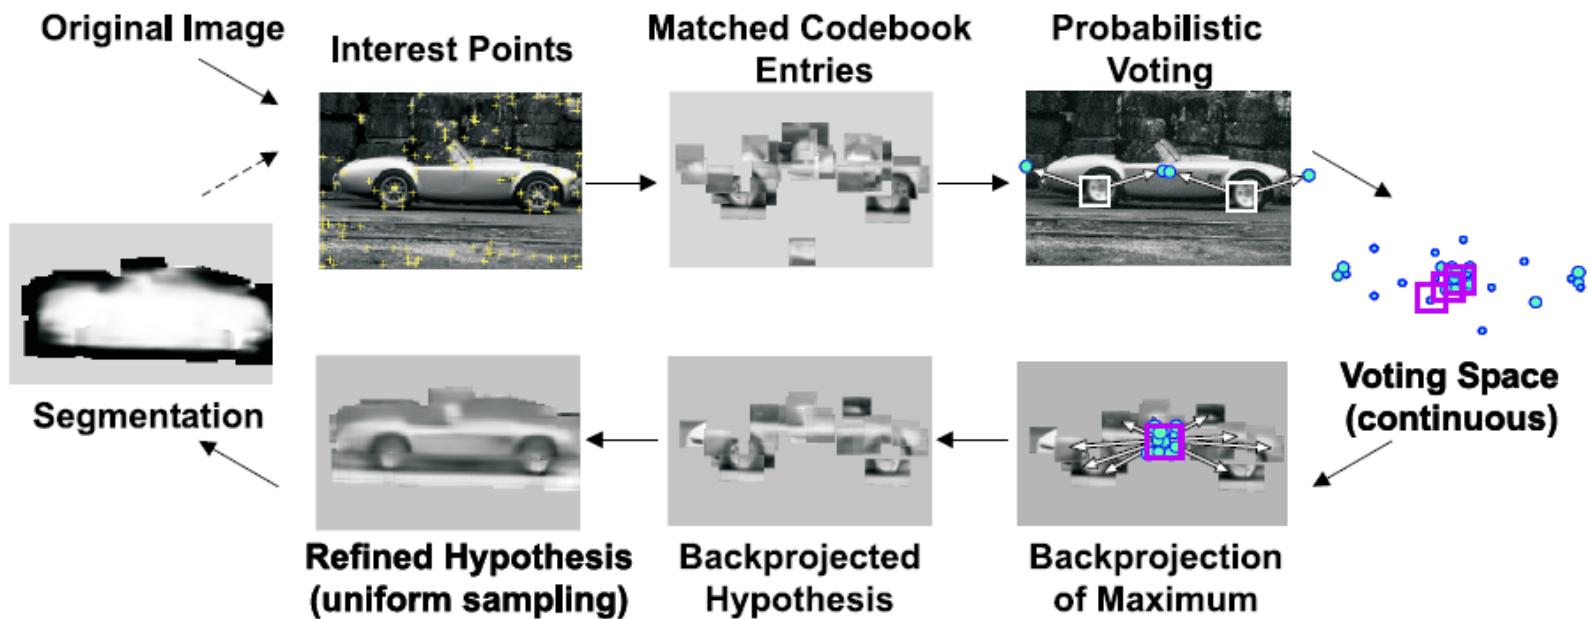
# Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry
3. For each codebook entry, store all positions it was found, relative to object center

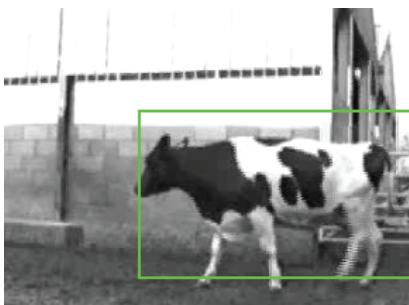
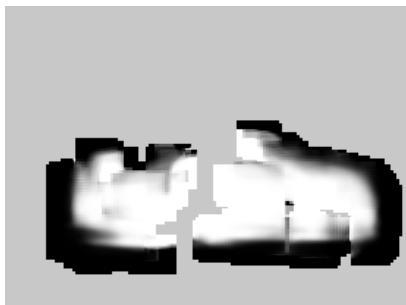
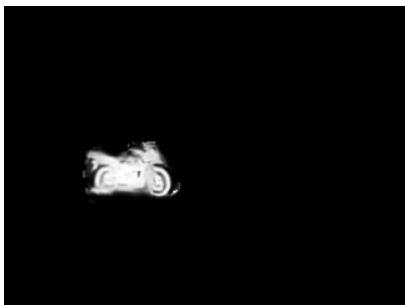
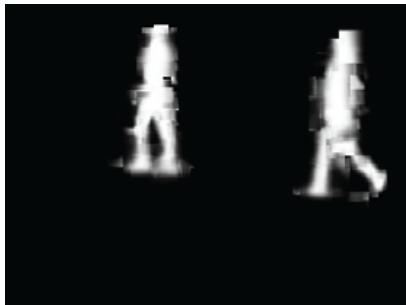
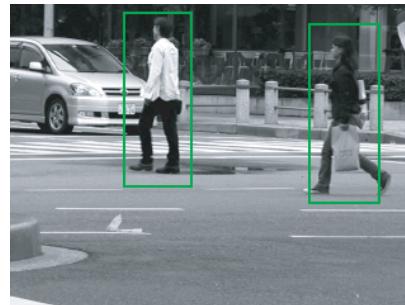
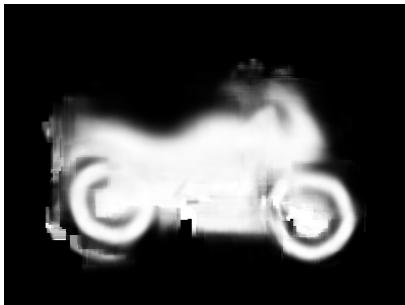


# Implicit shape models: Testing

1. Given test image, extract patches, match to codebook entry
2. Cast votes for possible positions of object center
3. Search for maxima in voting space
4. Extract weighted segmentation mask based on stored masks for the codebook occurrences

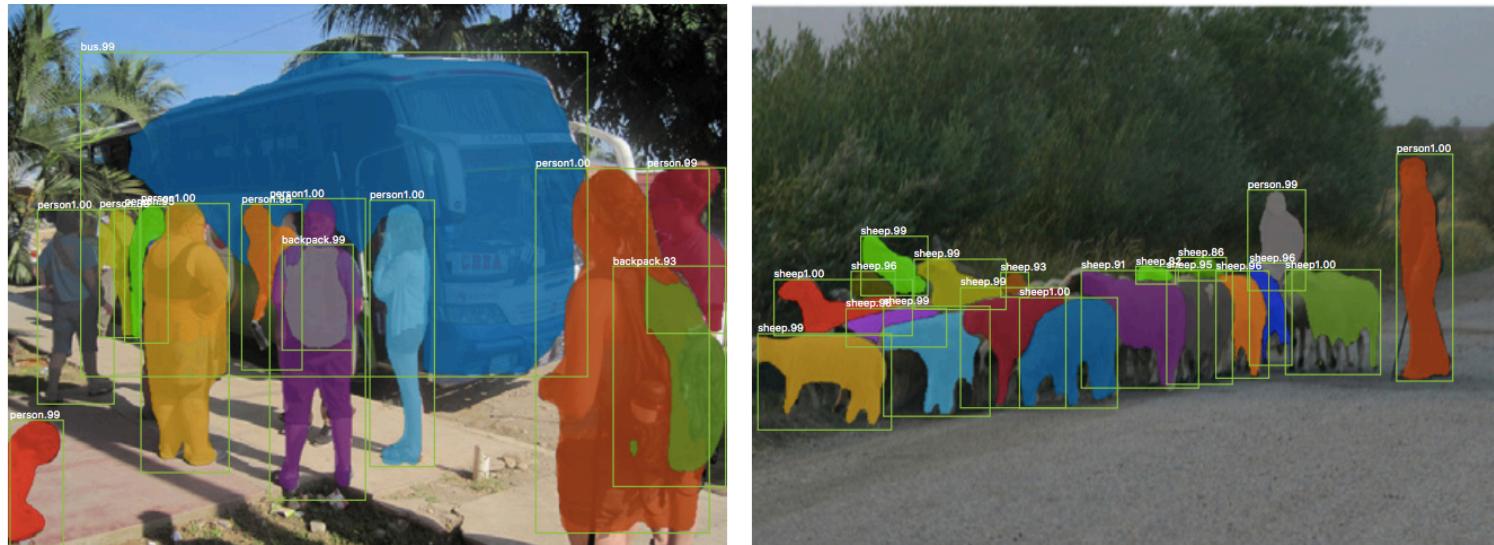
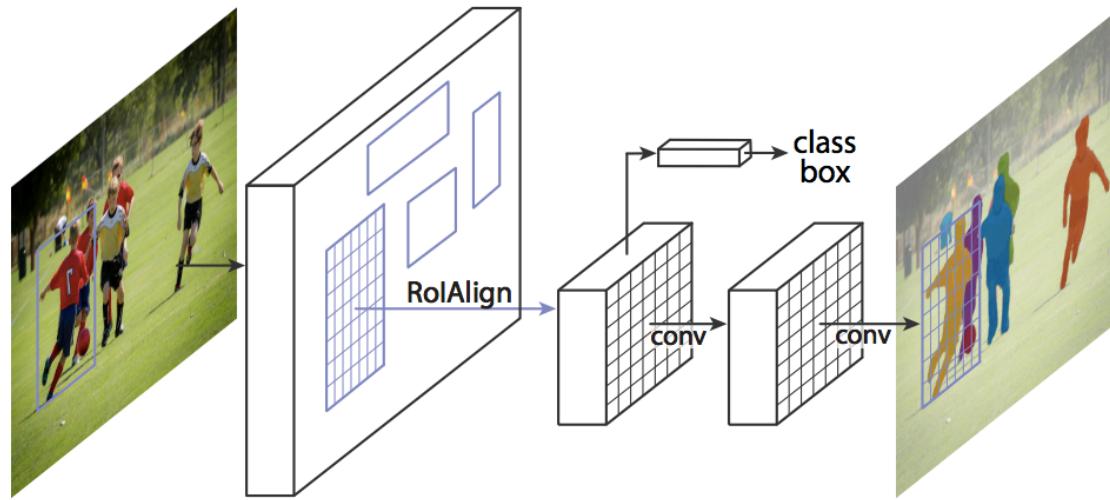


# Additional examples



B. Leibe, A. Leonardis, and B. Schiele,  
[Robust Object Detection with Interleaved Categorization and Segmentation](#),  
IJCV 77 (1-3), pp. 259-289, 2008.

# These days: Mask R-CNN



# Fitting: challenges

- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform

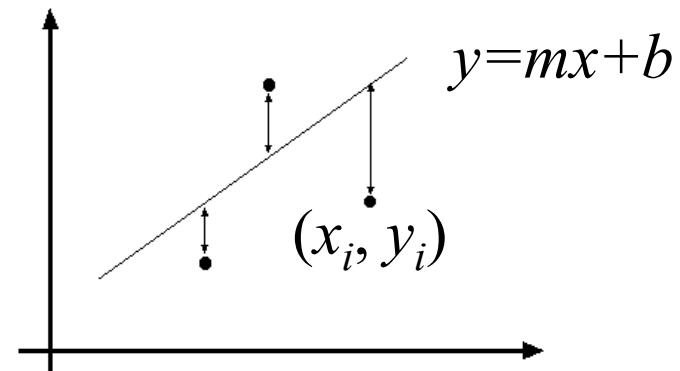
# Least squares line fitting

Data:  $(x_1, y_1), \dots, (x_n, y_n)$

Line equation:  $y_i = mx_i + b$

Find  $(m, b)$  to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \|Y - XB\|^2 \quad \text{where} \quad Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \quad X = \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad B = \begin{bmatrix} m \\ b \end{bmatrix}$$

$$E = \|Y - XB\|^2 = (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB)$$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$X^T XB = X^T Y$  *Normal equations:* least squares solution to  $XB=Y$

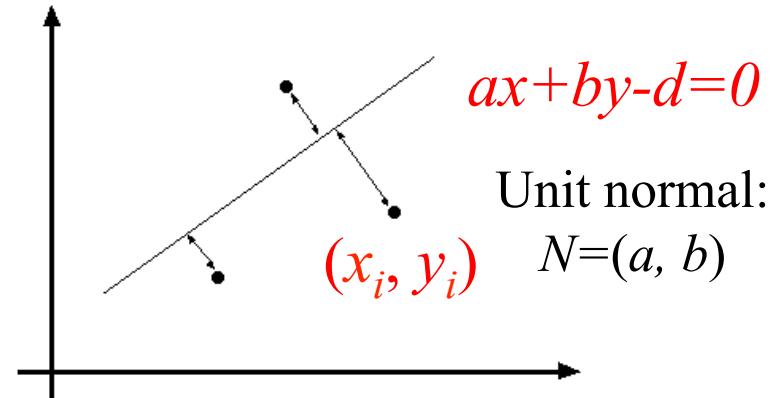
$$B = (X^T X)^{-1} X^T Y \quad \text{Solution, if } X^T X \text{ is invertible}$$

# Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines

# Total least squares

Distance between point  $(x_i, y_i)$  and line  
 $ax+by-d=0$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$

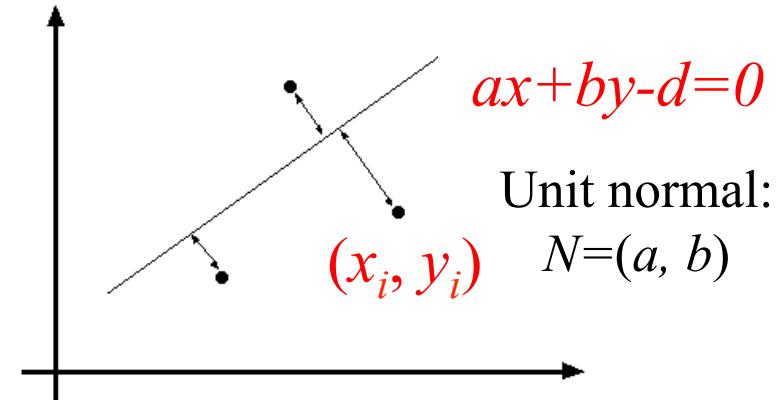


# Total least squares

Distance between point  $(x_i, y_i)$  and line  
 $ax+by-d=0$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$

Find  $(a, b, d)$  to minimize the sum of squared *perpendicular* distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$



# Total least squares

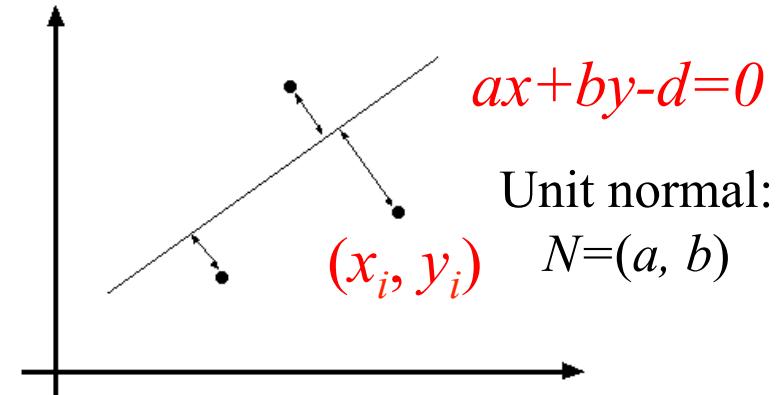
Distance between point  $(x_i, y_i)$  and line  
 $ax+by-d=0$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$

Find  $(a, b, d)$  to minimize the sum of squared *perpendicular* distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$



$$d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$

Solution to  $(U^T U)N = 0$ , subject to  $\|N\|^2 = 1$ : eigenvector of  $U^T U$  associated with the smallest eigenvalue (least squares solution to *homogeneous linear system*  $UN=0$ )

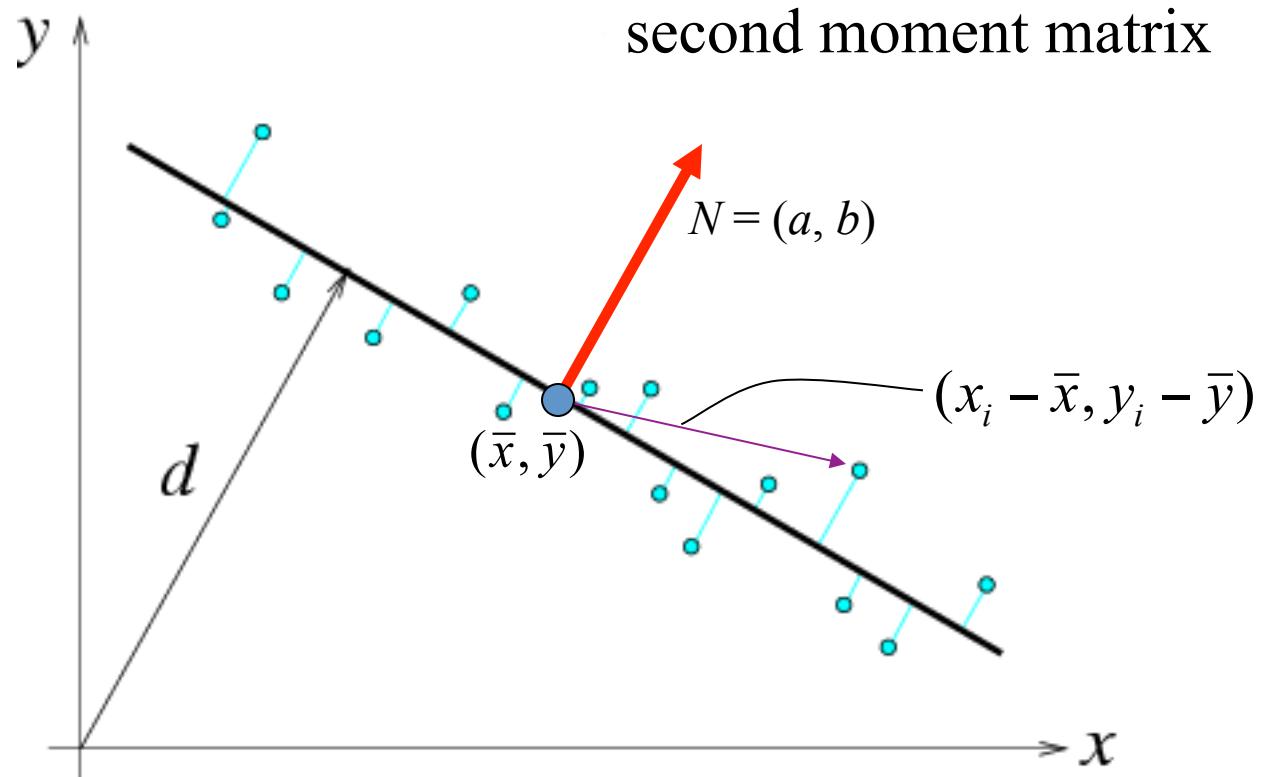
# Total least squares

$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \quad U^T U = \begin{bmatrix} \sum_{i=1}^n (x_i - \bar{x})^2 & \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \\ \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) & \sum_{i=1}^n (y_i - \bar{y})^2 \end{bmatrix}$$

second moment matrix

# Total least squares

$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \quad U^T U = \begin{bmatrix} \sum_{i=1}^n (x_i - \bar{x})^2 & \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \\ \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) & \sum_{i=1}^n (y_i - \bar{y})^2 \end{bmatrix}$$

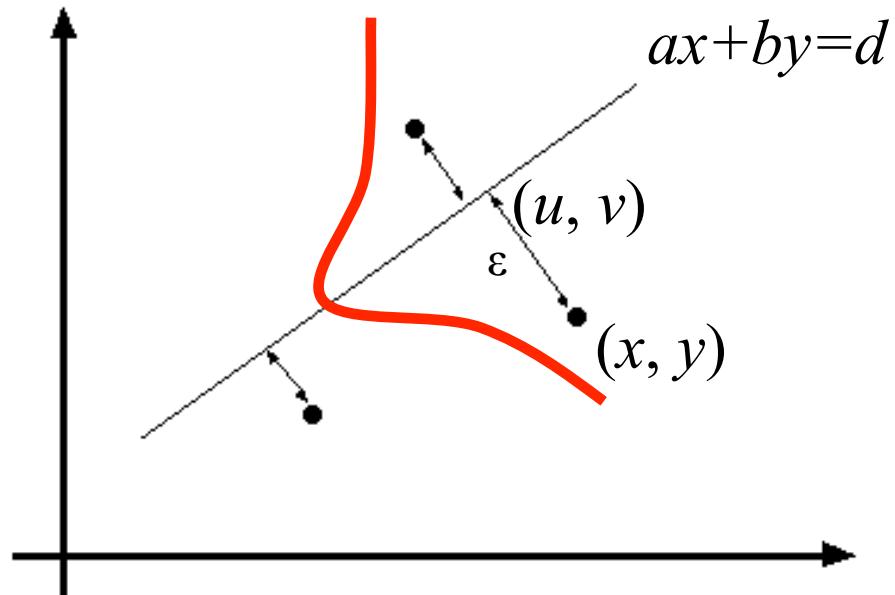


# Least squares as likelihood maximization

- **Generative model:** line points are sampled independently and corrupted by Gaussian noise in the direction perpendicular to the line

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} u \\ v \end{pmatrix} + \varepsilon \begin{pmatrix} a \\ b \end{pmatrix}$$

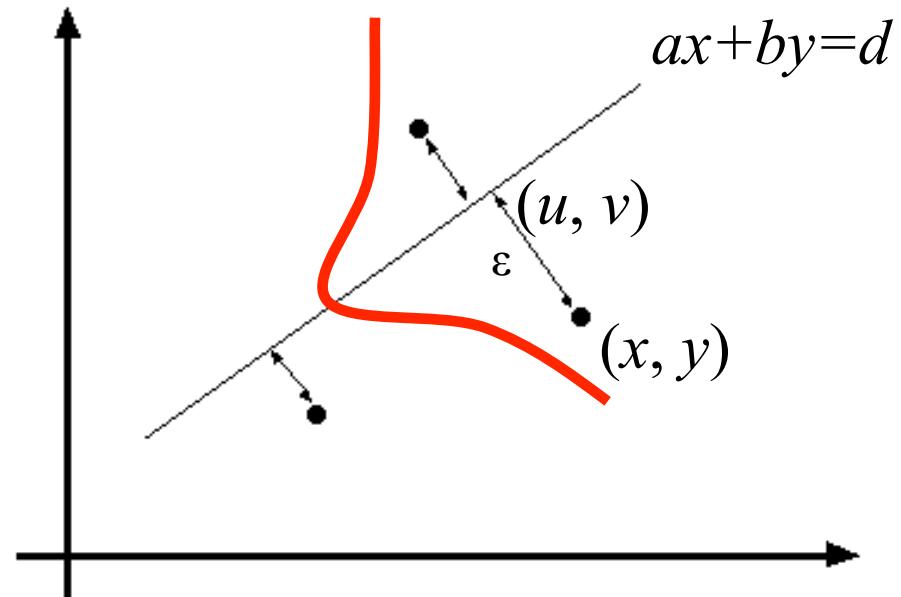
point on the line      noise: sampled from zero-mean Gaussian with std. dev.  $\sigma$       normal direction



# Least squares as likelihood maximization

- **Generative model:** line points are sampled independently and corrupted by Gaussian noise in the direction perpendicular to the line

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} u \\ v \end{pmatrix} + \varepsilon \begin{pmatrix} a \\ b \end{pmatrix}$$



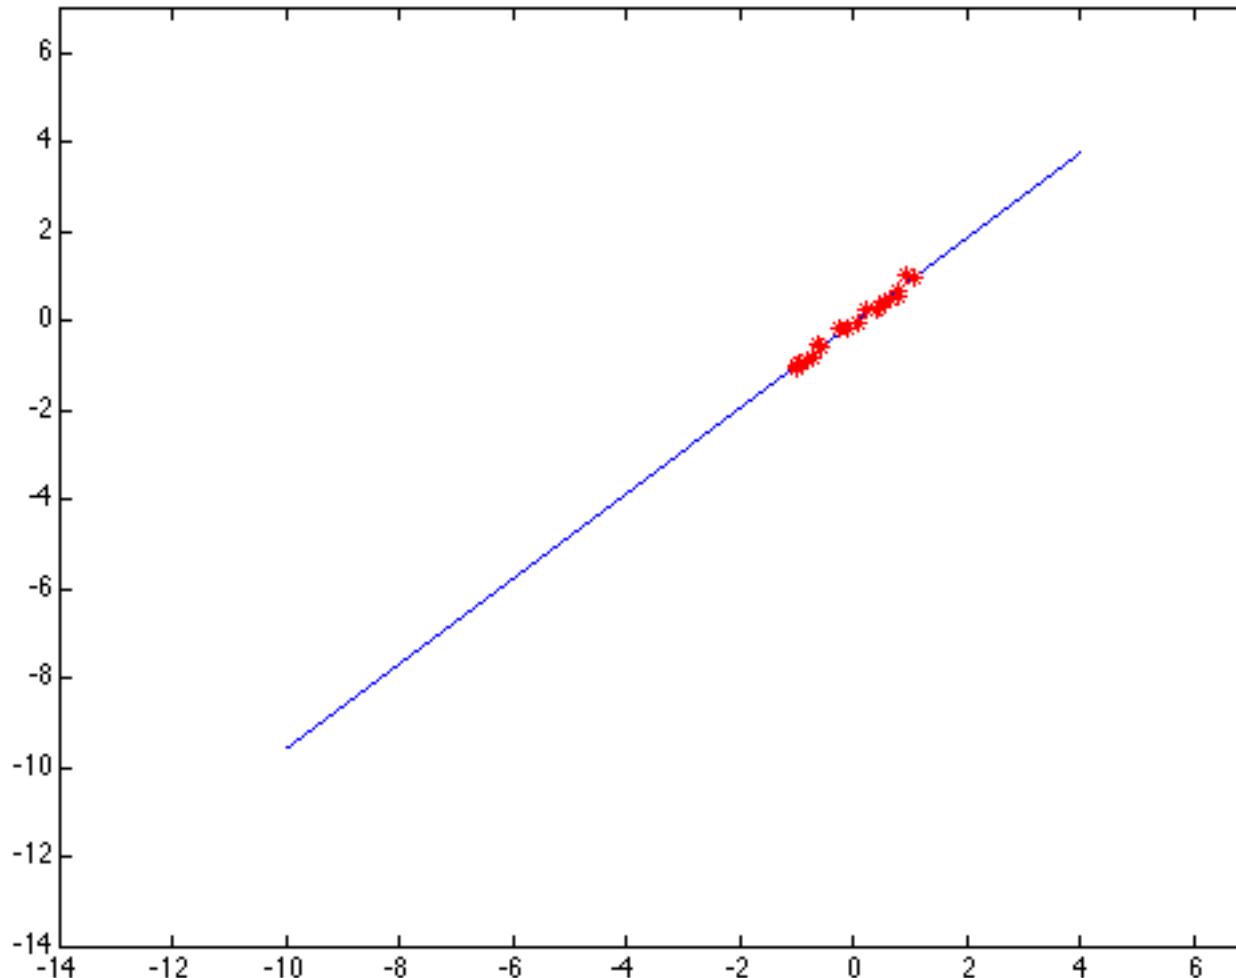
*Likelihood* of points given line parameters  $(a, b, d)$ :

$$P(x_1, y_1, \dots, x_n, y_n | a, b, d) = \prod_{i=1}^n P(x_i, y_i | a, b, d) \propto \prod_{i=1}^n \exp\left(-\frac{(ax_i + by_i - d)^2}{2\sigma^2}\right)$$

Log-likelihood:  $L(x_1, y_1, \dots, x_n, y_n | a, b, d) = -\frac{1}{2\sigma^2} \sum_{i=1}^n (ax_i + by_i - d)^2$

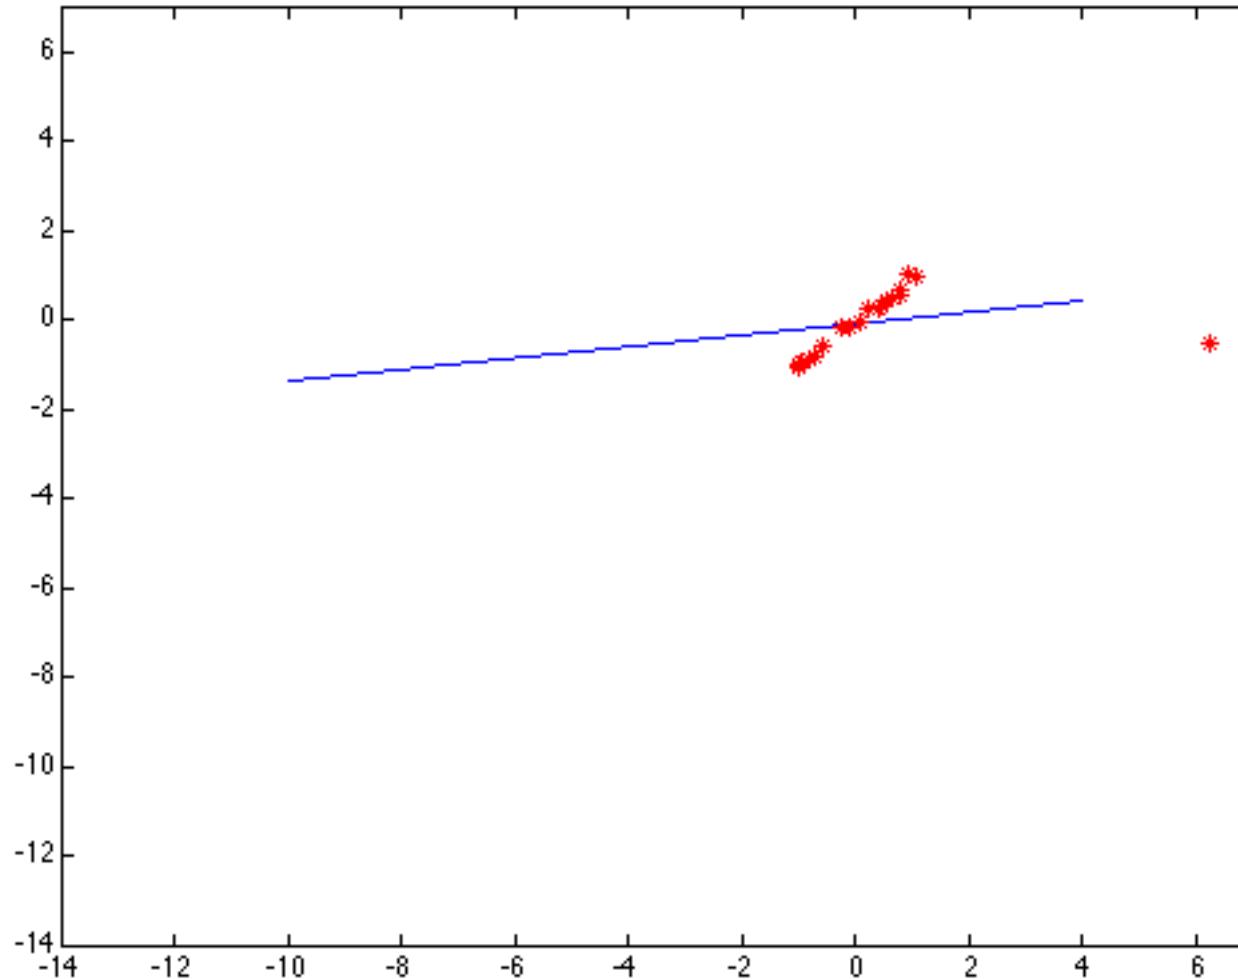
# Least squares: Robustness to noise

Least squares fit to the red points:



# Least squares: Robustness to noise

Least squares fit with an outlier:



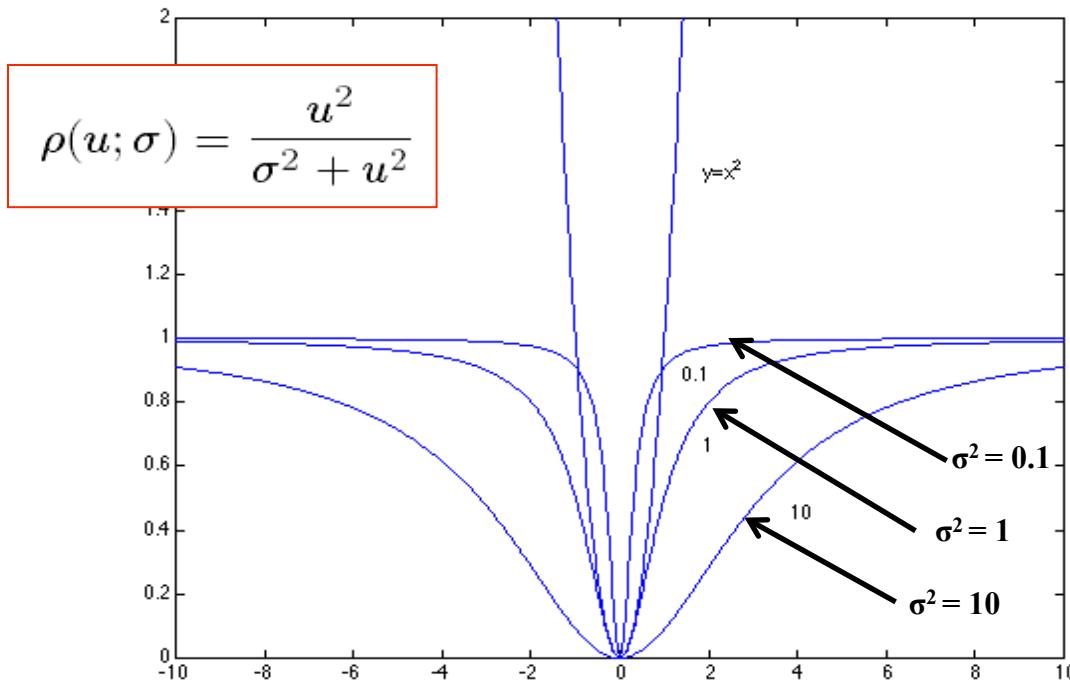
Problem: squared error heavily penalizes outliers

# Robust estimators

- General approach: find model parameters  $\theta$  that minimize

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$

- $r_i(x_i, \theta)$  – residual of  $i^{\text{th}}$  point w.r.t. model parameters  $\theta$
- $\rho$  – robust function with scale parameter  $\sigma$



The robust function  $\rho$  behaves like squared distance for small values of the residual  $u$  but saturates for larger values of  $u$

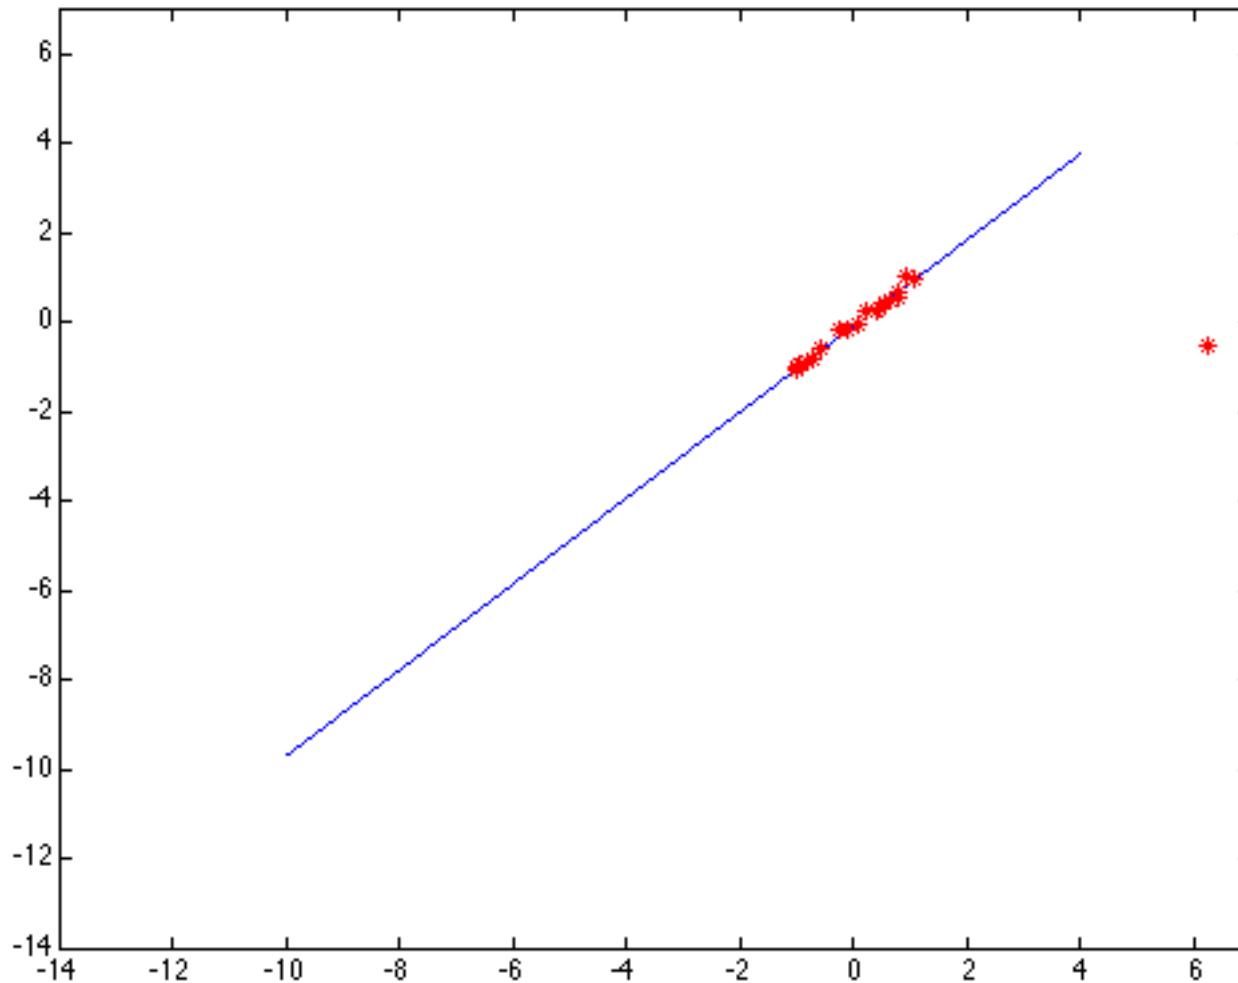
# Robust estimators

- General approach: find model parameters  $\theta$  that minimize

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$

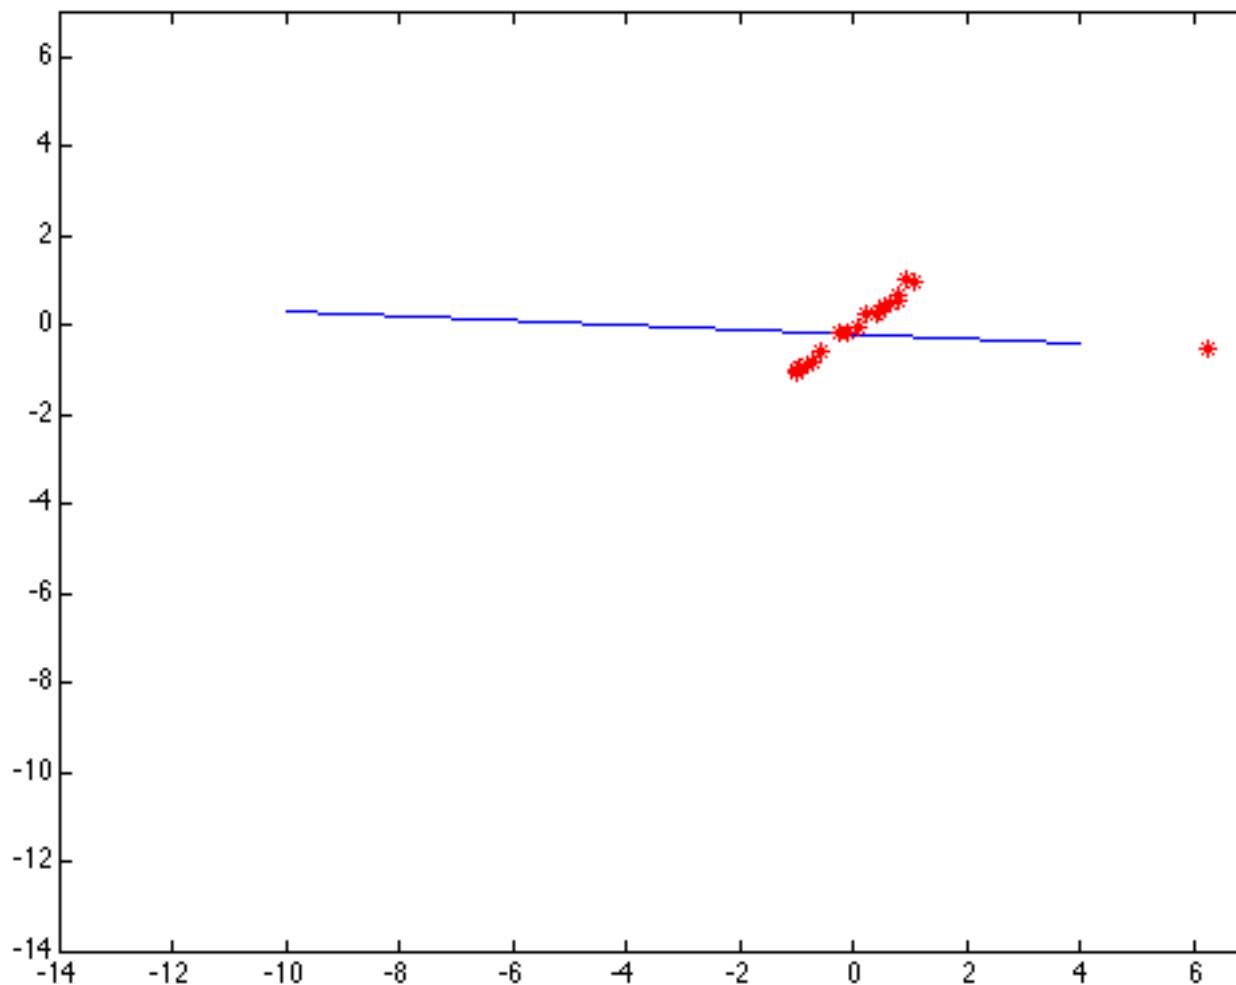
- $r_i(x_i, \theta)$  – residual of  $i^{\text{th}}$  point w.r.t. model parameters  $\theta$
- $\rho$  – robust function with scale parameter  $\sigma$
- Robust fitting is a nonlinear optimization problem that must be solved iteratively
- Least squares solution can be used for initialization
- Scale of robust function should be chosen carefully

# Choosing the scale: just right



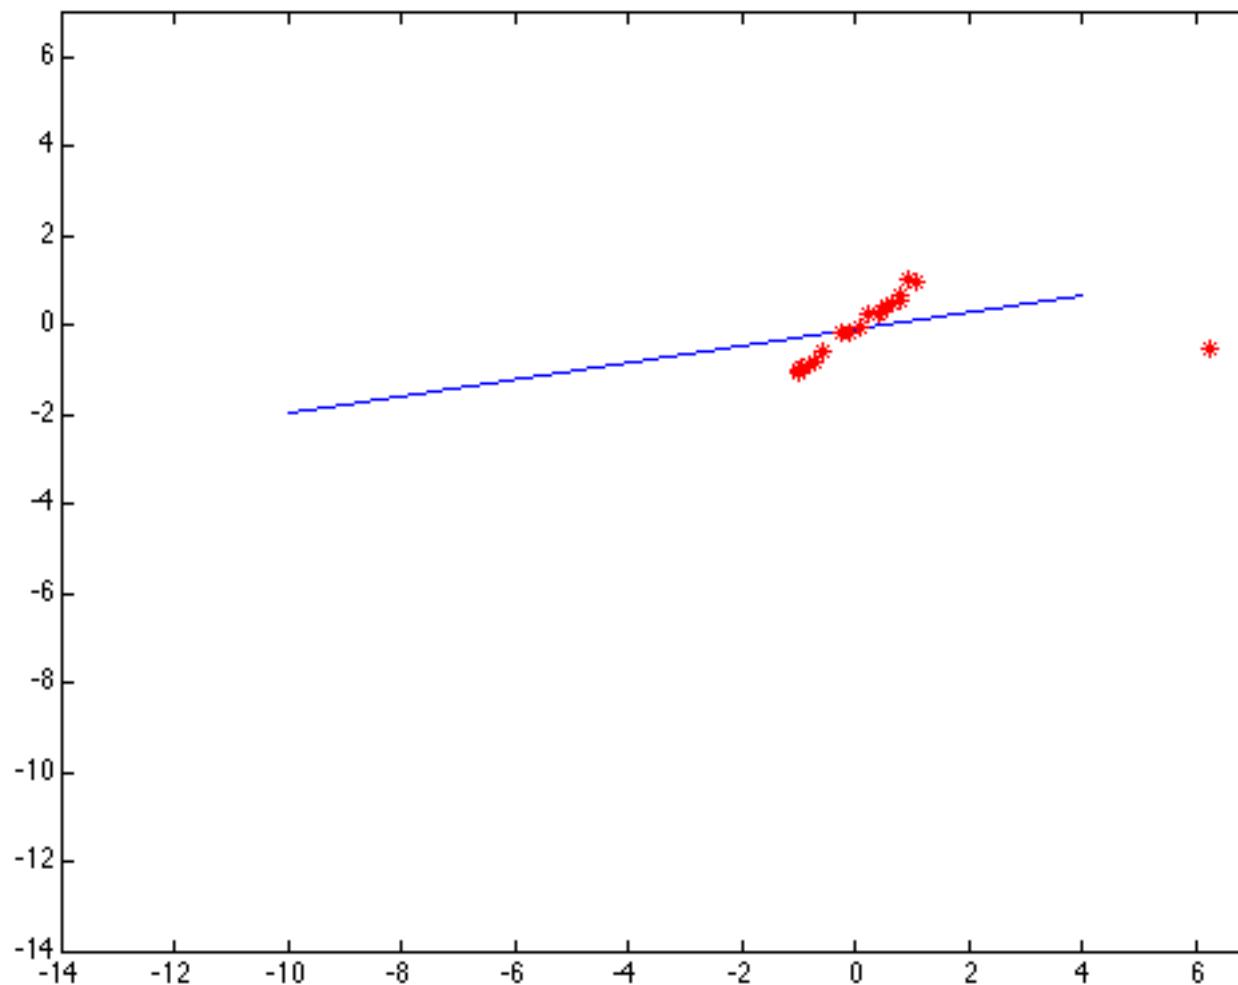
The effect of the outlier is minimized

# Choosing the scale: too small



The error value is almost the same for every point and the fit is very poor

# Choosing the scale: too large

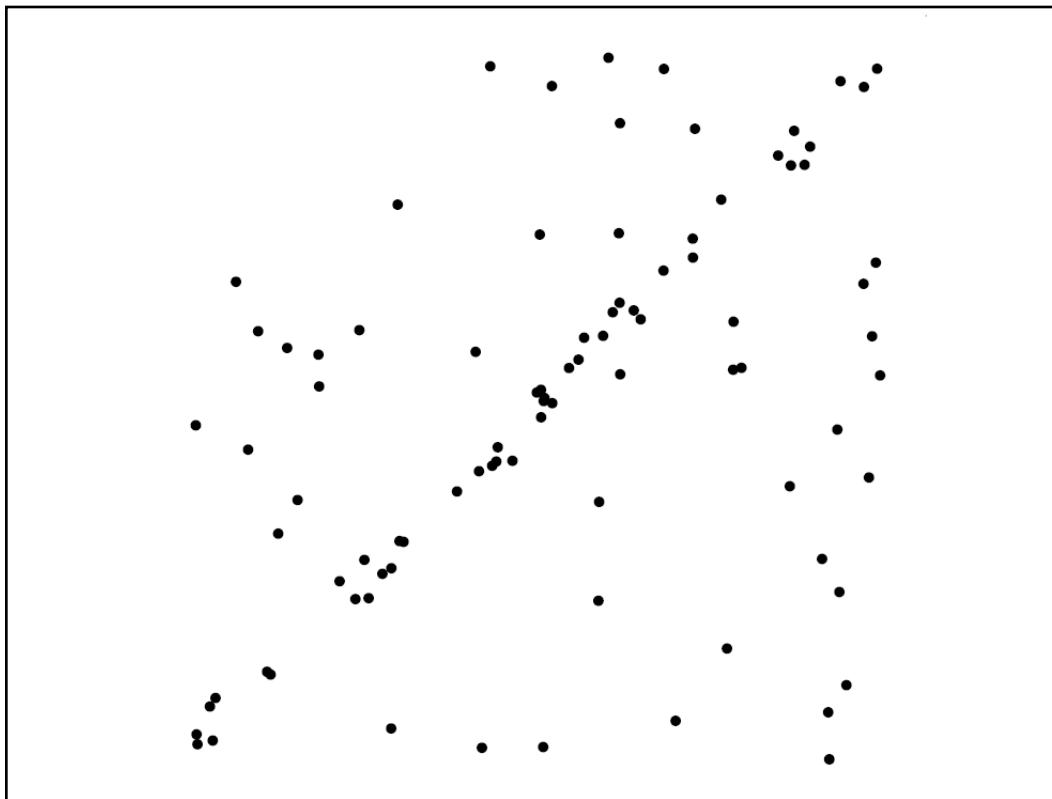


Behaves much the same as least squares

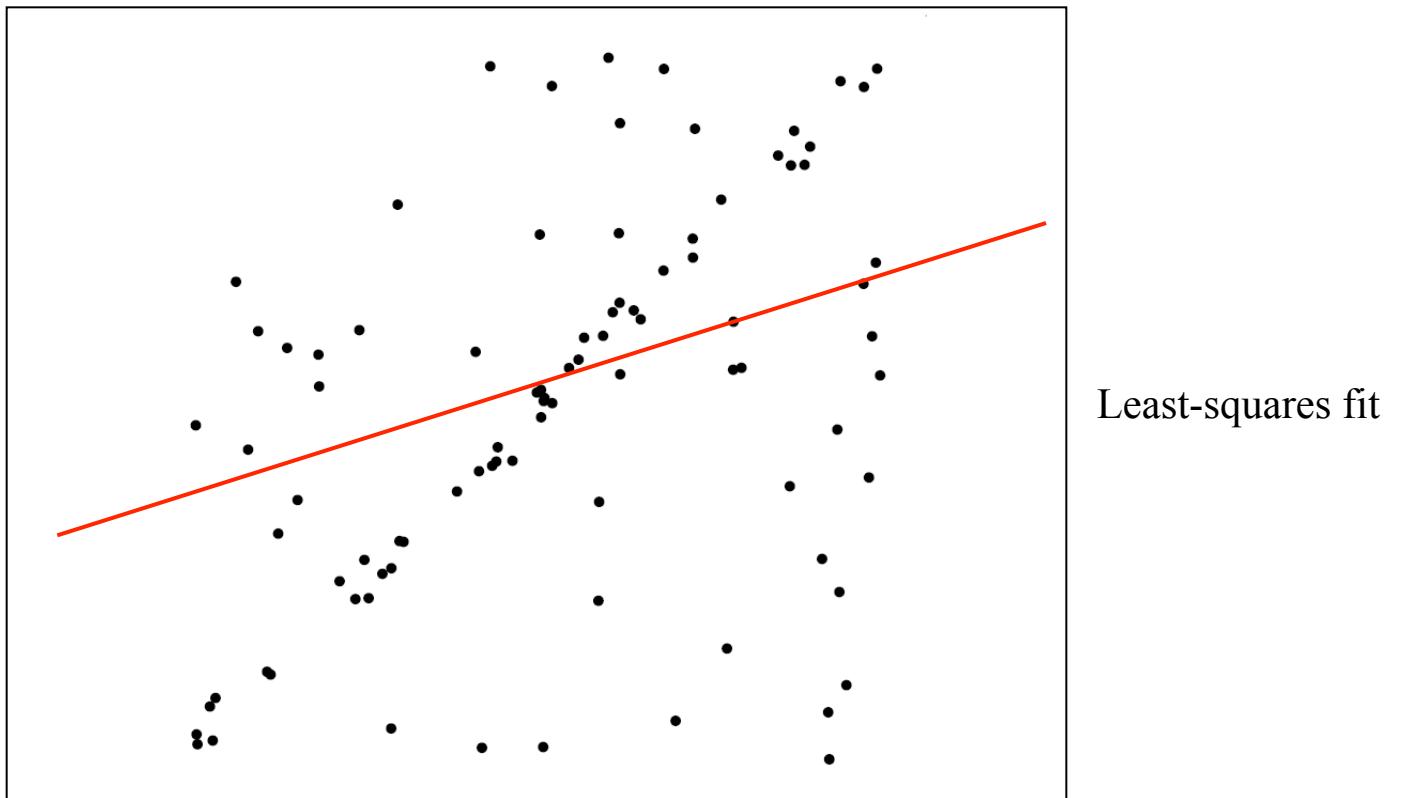
# RANSAC

- Robust fitting can deal with a few outliers – what if we have many?
- Random sample consensus (RANSAC):  
Very general framework for model fitting in the presence of outliers
- Outline
  - Choose a small subset of points uniformly at random
  - Fit a model to that subset
  - Find all remaining points that are “close” to the model and reject the rest as outliers
  - Do this many times and choose the best model

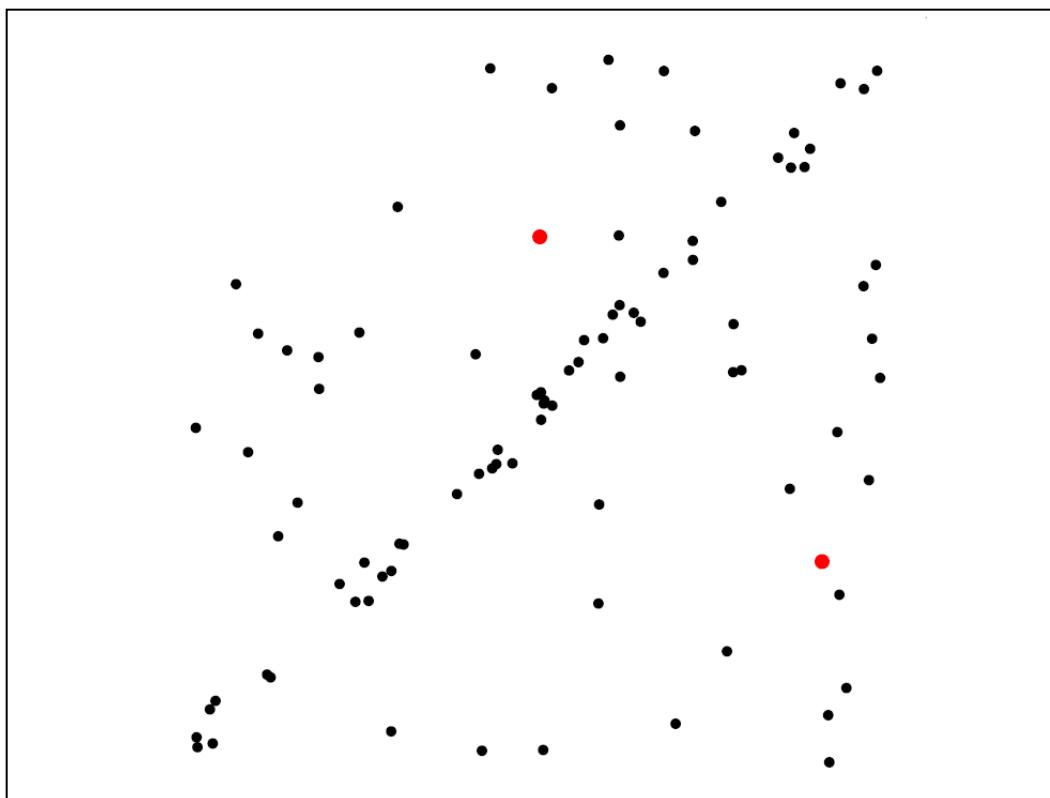
# RANSAC for line fitting example



# RANSAC for line fitting example

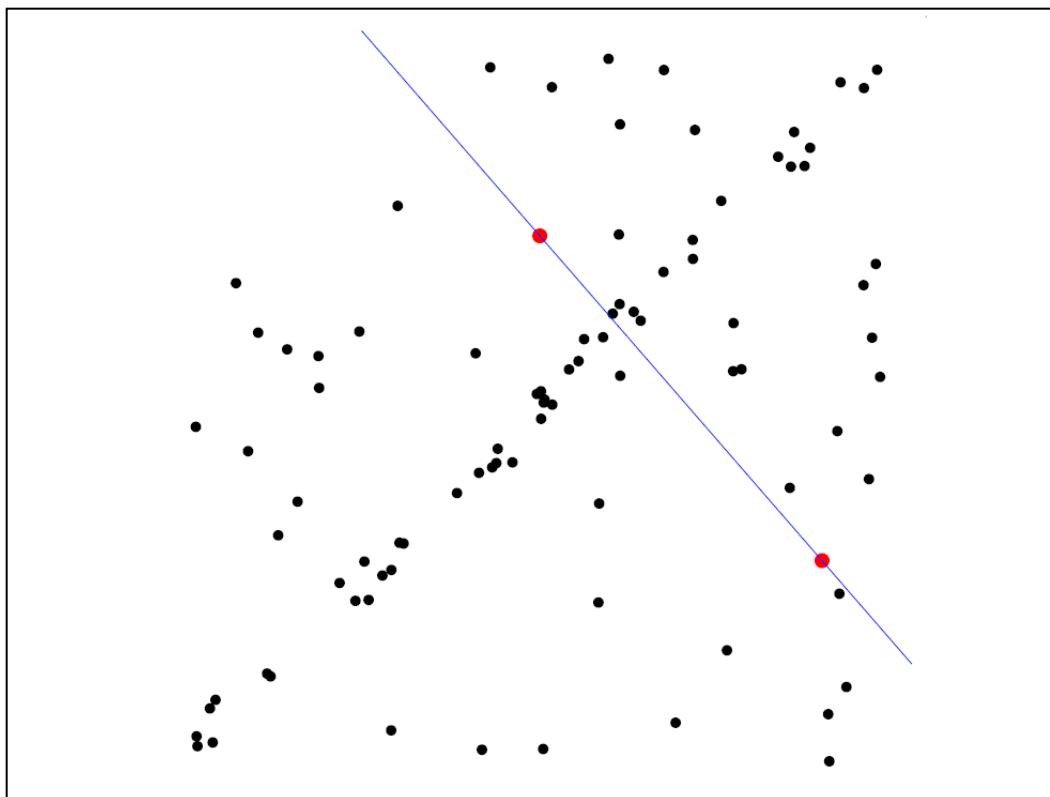


# RANSAC for line fitting example



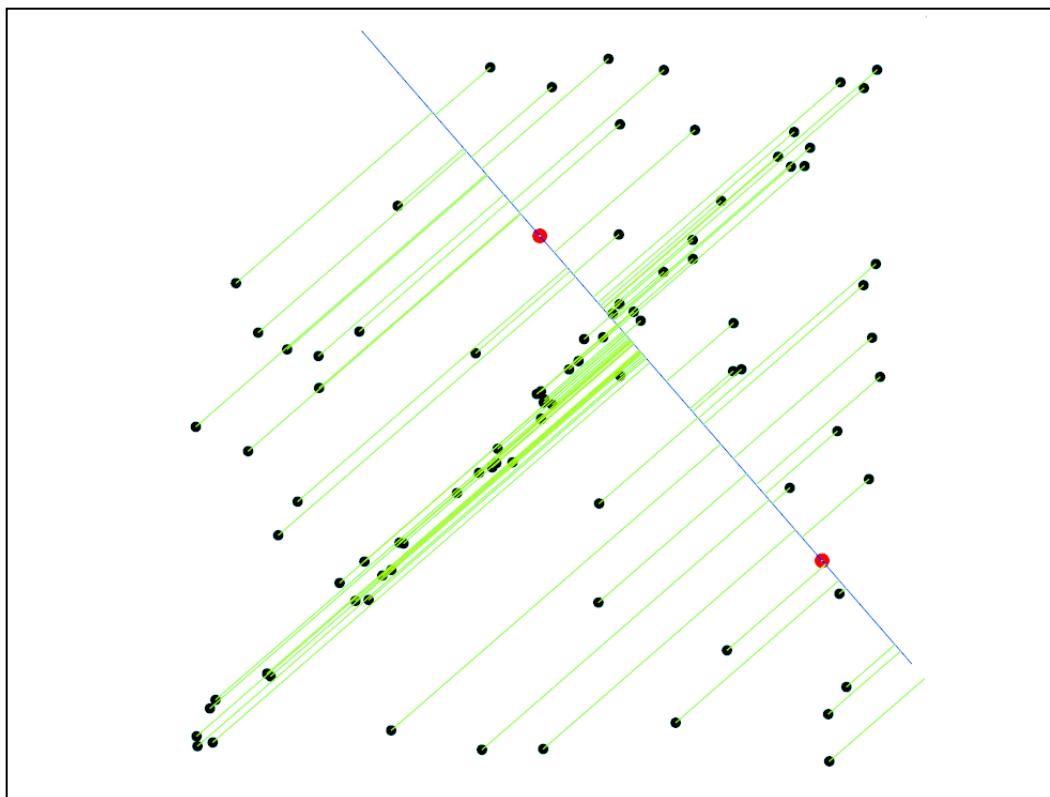
1. Randomly select minimal subset of points

# RANSAC for line fitting example



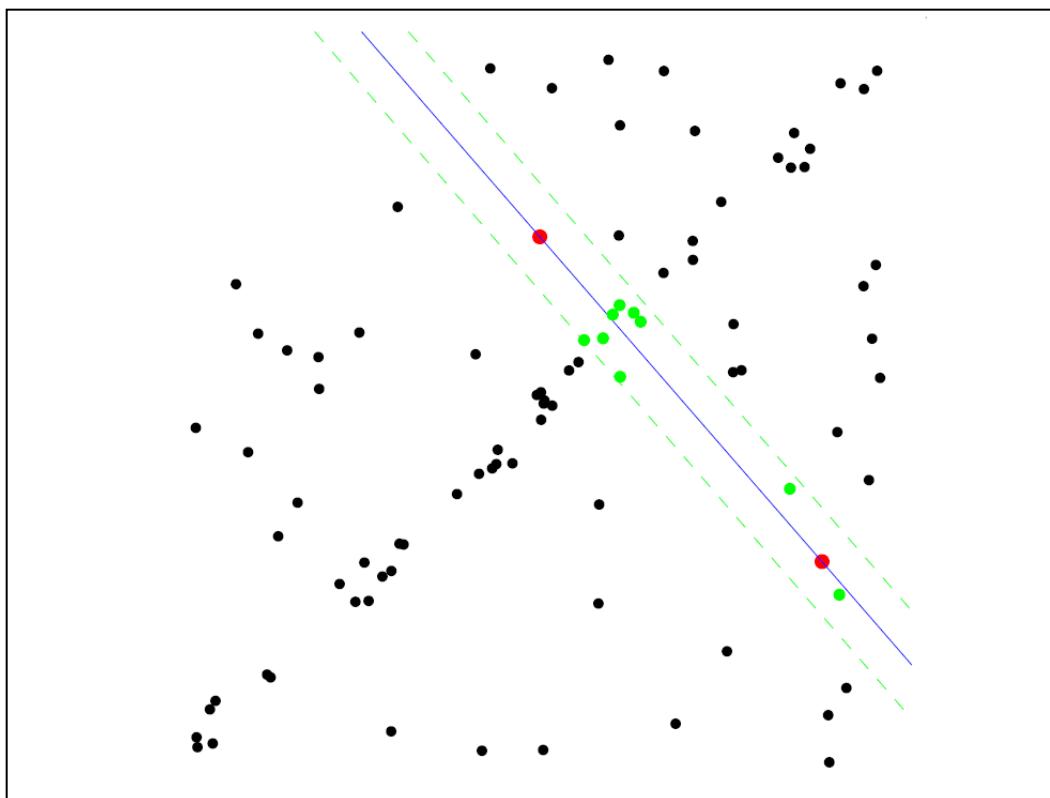
1. Randomly select minimal subset of points
2. Hypothesize a model

# RANSAC for line fitting example



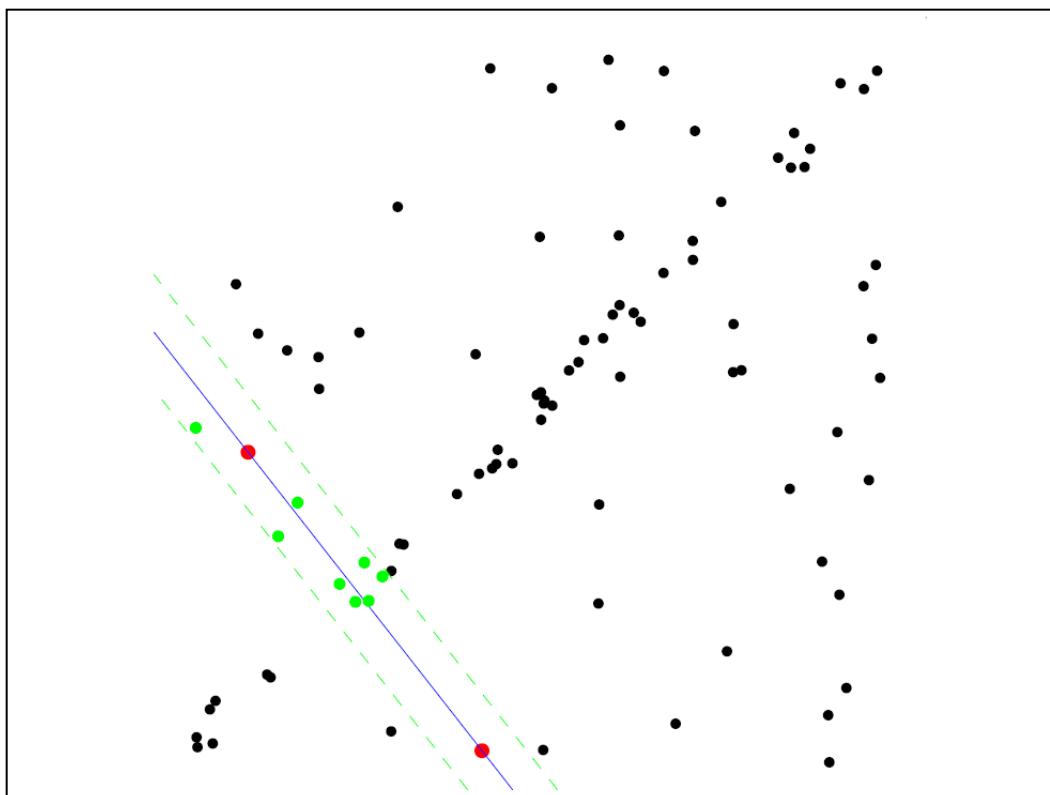
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

# RANSAC for line fitting example



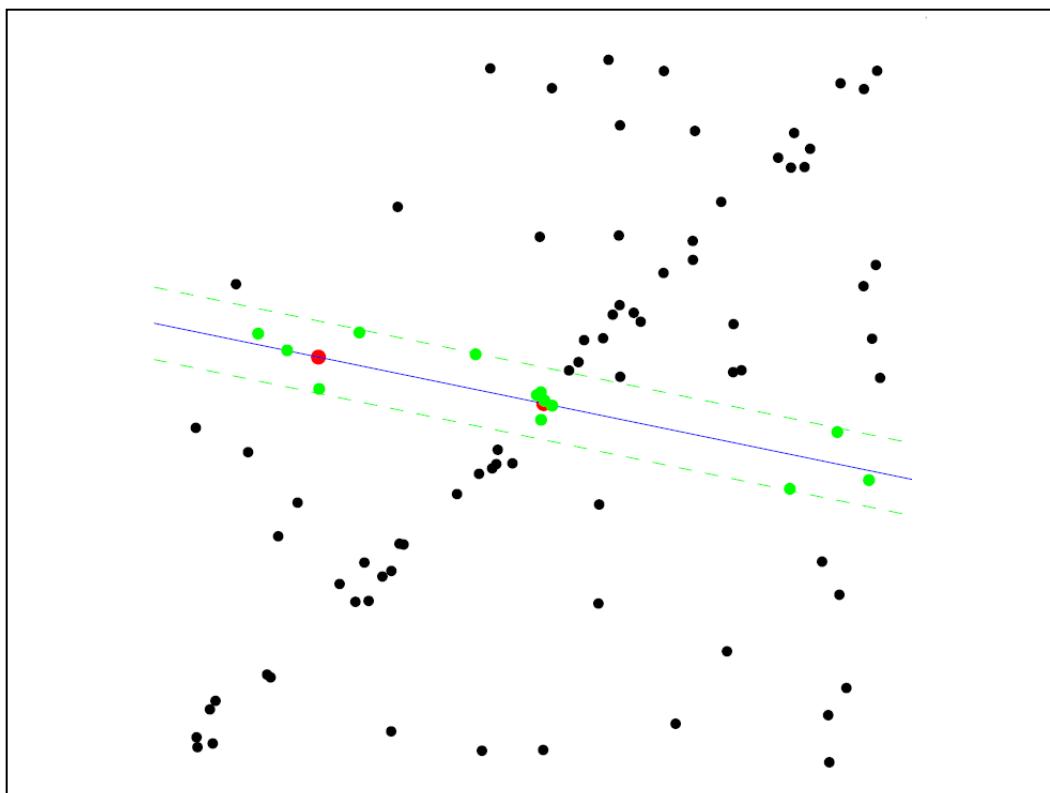
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. **Select points consistent with model**

# RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

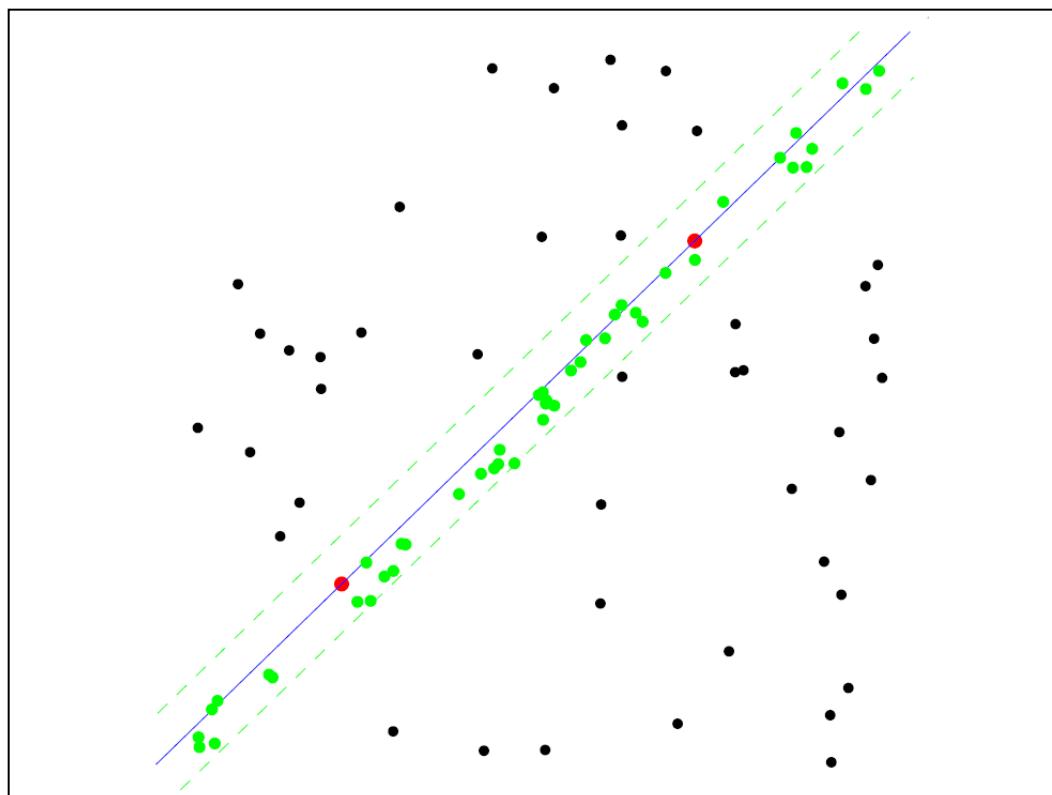
# RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

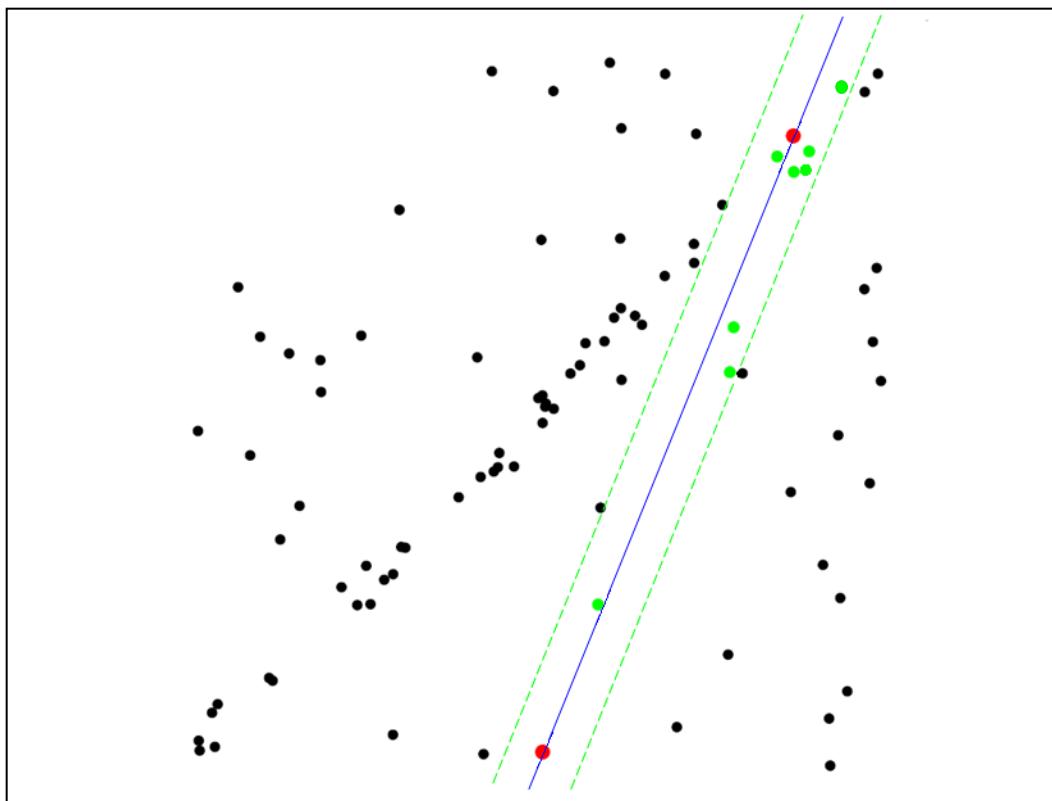
# RANSAC for line fitting example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting

- Repeat  $N$  times:
  - Draw  $s$  points uniformly at random
  - Fit line to these  $s$  points
  - Find *inliers* to this line among the remaining points (i.e., points whose distance from the line is less than  $t$ )
  - If there are  $d$  or more inliers, accept the line and refit using all inliers

# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$ ,  $t=1.9\sigma$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )

# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

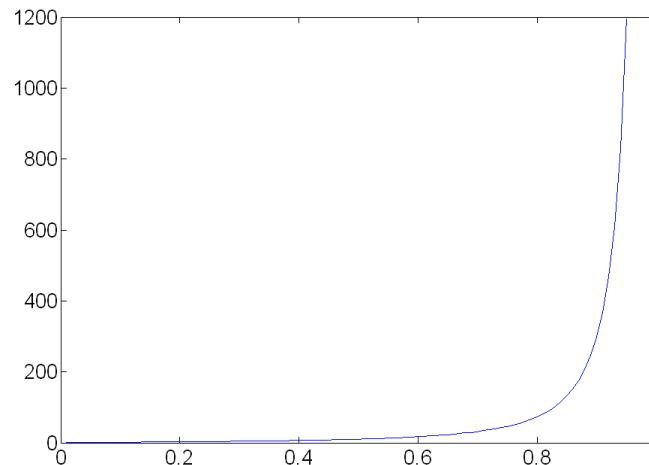
s	proportion of outliers $e$							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$



# Choosing the parameters

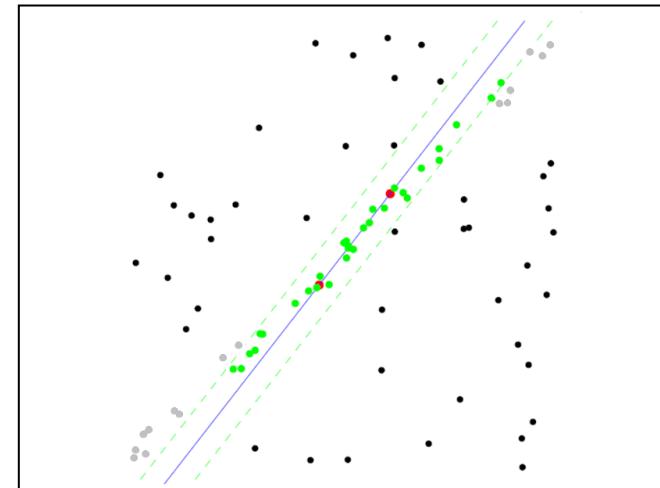
- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )
- Consensus set size  $d$ 
  - Should match expected inlier ratio

# Adaptively determining the number of samples

- Outlier ratio  $e$  is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield  $e=0.2$
- Adaptive procedure:
  - $N = \infty$ ,  $sample\_count = 0$
  - While  $N > sample\_count$ 
    - Choose a sample and count the number of inliers
    - If inlier ratio is highest of any found so far, set  $e = 1 - (\text{number of inliers})/(\text{total number of points})$
    - Recompute  $N$  from  $e$ :
$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$
    - Increment the  $sample\_count$  by 1

# RANSAC pros and cons

- Pros
  - Simple and general
  - Applicable to many different problems
  - Often works well in practice
- Cons
  - Lots of parameters to tune
  - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
  - Can't always get a good initialization of the model based on the minimum number of samples



# Course structure

## 1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

## 2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

## 3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

## 4. Object Recognition: high – level vision

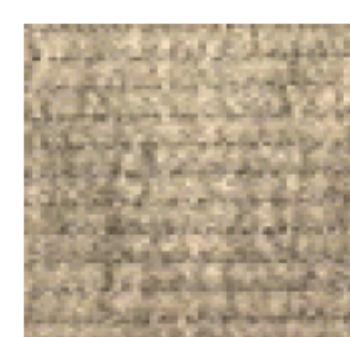
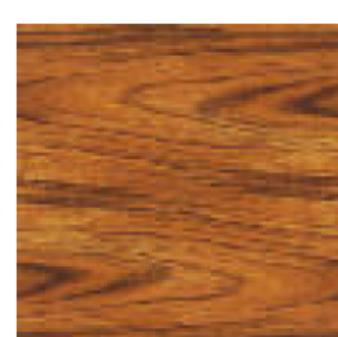
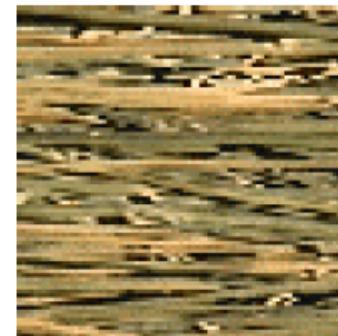
Object classification, object detection, part based models, bovw models

## 5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

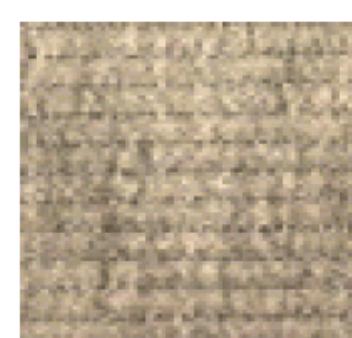
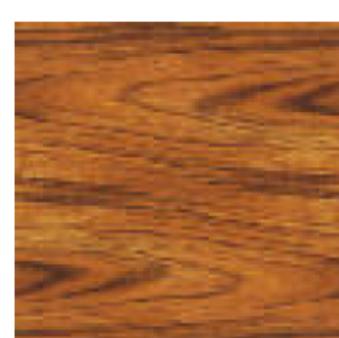
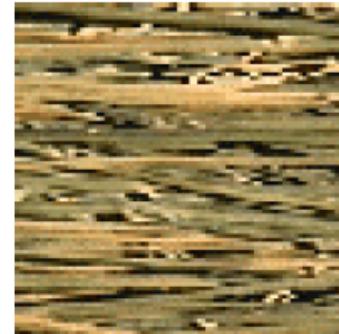
# Texture

# Texture



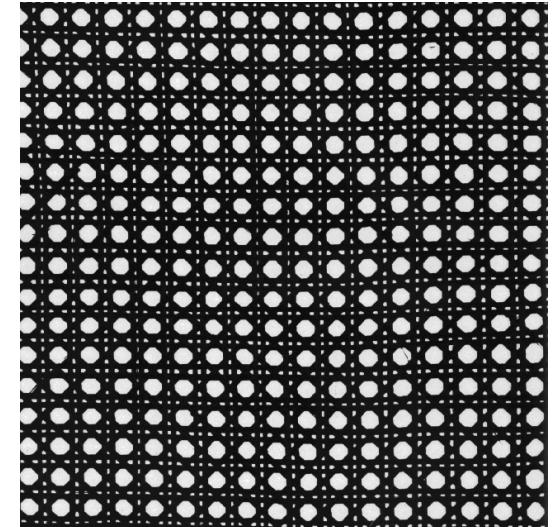
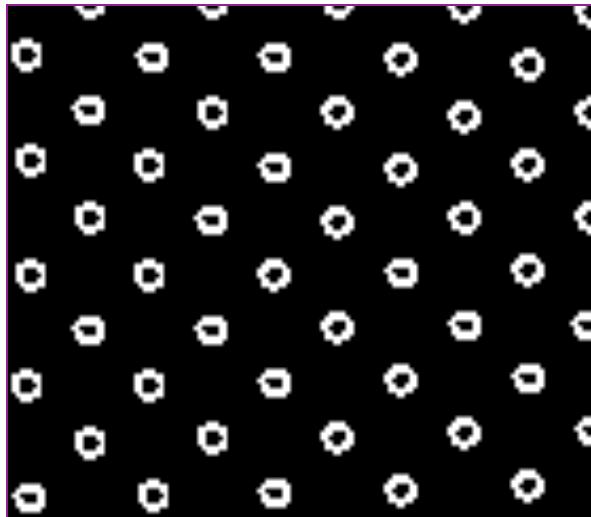
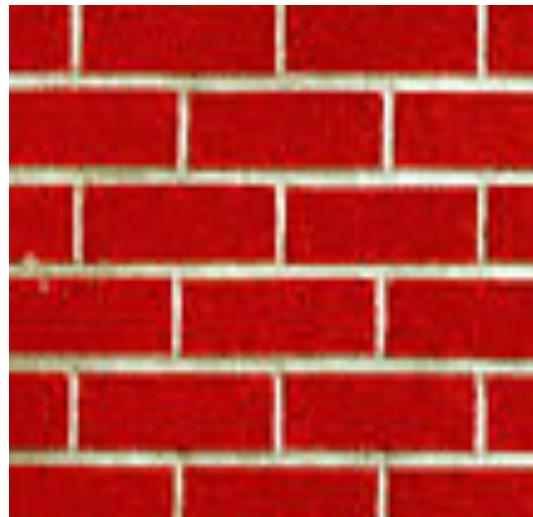
What defines a texture?

# Texture



Similar structures of pixels (patterns) that repeat themselves  
Similar aspect at a fixe scale  
Follow some statistics properties

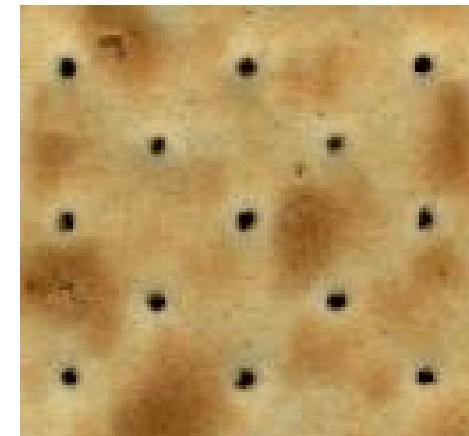
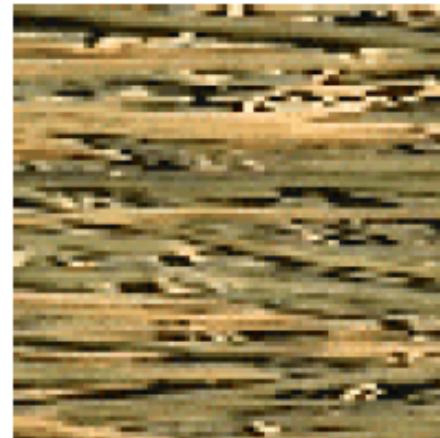
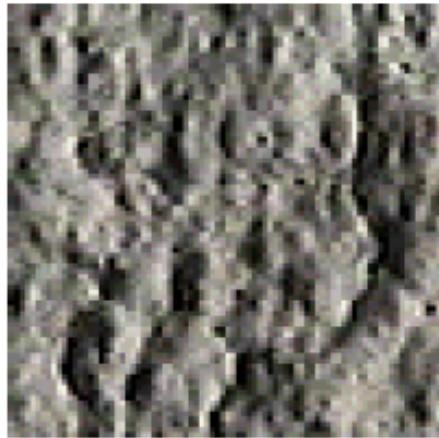
# Includes: more regular patterns



Texture = set of *texels* in some regular pattern

*Texels* = *texture elements*

# Includes: more random patterns



Texture = set of *texels* in some regular or repeated pattern

*Texels* = *texture elements*

# Scale and texture



Leaf - object



Tree - texture

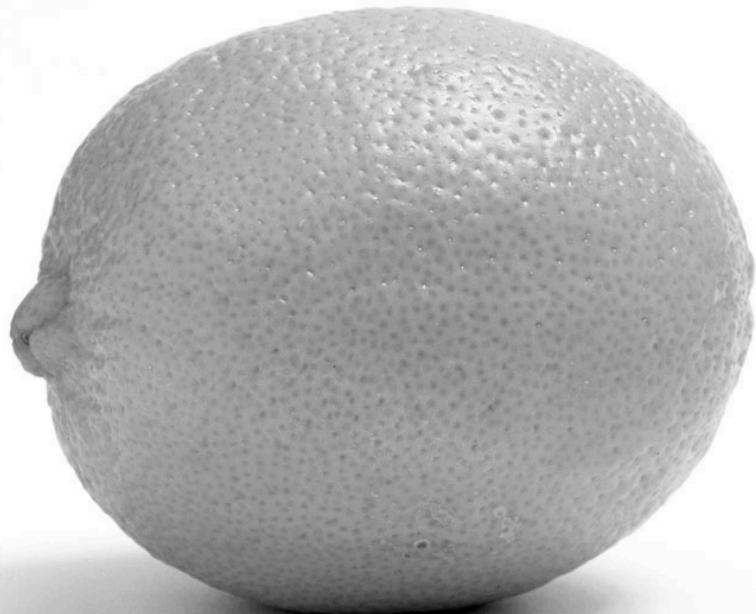
# Why analyze texture?

Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture



Slide credit: Kristen Grauman



Slide credit: Kristen Grauman



# Why analyze texture?

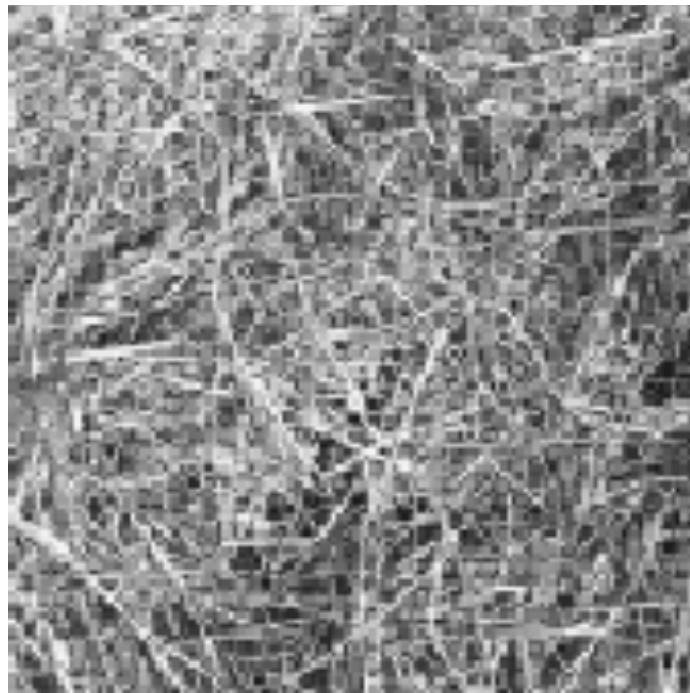
Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

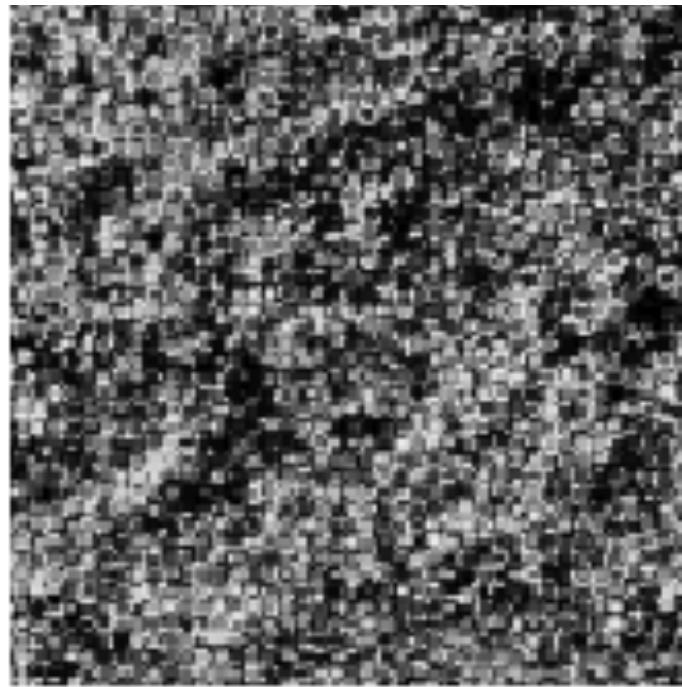
Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.

# Natural Textures



grass



leaves

What/Where are the texels?

# The Case for Statistical Texture

- Segmenting out texels is difficult or impossible in real images.
- Numeric quantities or statistics that describe a texture can be computed from the grayscale (or colors) images alone.
- This approach is less intuitive, but is computationally efficient.
- It can be used for both classification and segmentation.

# Some Simple Statistical Texture Measures

## Edge Density and Direction

- Use an edge detector as the first step in texture analysis.
- The number of edge pixels in a fixed-size region tells us how busy that region is.
- The directions of the edges also help characterize the texture

# Edge-based Texture Measures

1. edgeness per unit area

$$\text{Edgeness} = |\{ p \mid \text{gradient\_magnitude}(p) \geq \text{threshold}\}| / N$$

where  $N$  is the size of the unit area

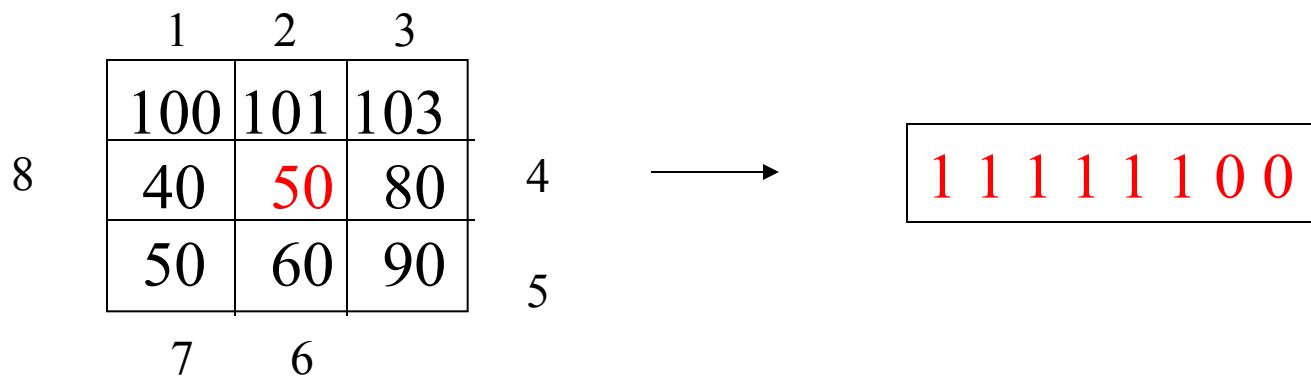
2. edge magnitude and direction histograms

$$F_{magdir} = ( H_{magnitude}, H_{direction} )$$

where these are the normalized histograms of gradient magnitudes and gradient directions, respectively.

# Local Binary Pattern Measure

- For each pixel  $p$ , create an 8-bit number  $b_1 b_2 b_3 b_4 b_5 b_6 b_7 b_8$ , where  $b_i = 0$  if neighbor  $i$  has value less than or equal to  $p$ 's value and 1 otherwise.
- Represent the texture in the image (or a region) by the histogram of these numbers.



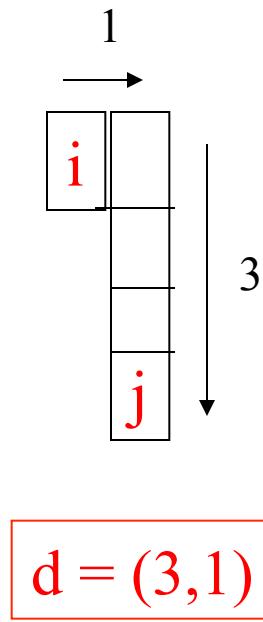
# Co-occurrence Matrix Features

A co-occurrence matrix is a 2D array  $C$  in which

- Both the rows and columns represent a set of possible image values.
- $C_d(i,j)$  indicates how many times value  $i$  co-occurs with value  $j$  in a particular spatial relationship  $d$ .
- The spatial relationship is specified by a vector  $d = (dr, dc)$ .

# Co-occurrence Example

1	1	0	0
1	1	0	0
0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2



0	1	2
0	1	3
1	2	2
2	0	1

$C_d$

co-occurrence  
matrix

grayscale  
image

From  $C_d$  we can compute  $N_d$ , the normalized co-occurrence matrix, where each value is divided by the sum of all the values.

# Co-occurrence Features

- extract numerical features from the co-occurrence matrix which can be used for representing and comparing textures

$$\begin{aligned} Energy &= \sum_i \sum_j N_d^2(i, j) \\ Entropy &= -\sum_i \sum_j N_d(i, j) \log_2 N_d(i, j) \\ Contrast &= \sum_i \sum_j (i - j)^2 N_d(i, j) \\ Homogeneity &= \sum_i \sum_j \frac{N_d(i, j)}{1 + |i - j|} \\ Correlation &= \frac{\sum_i \sum_j (i - \mu_i)(j - \mu_j) N_d(i, j)}{\sigma_i \sigma_j} \end{aligned} \tag{7.11}$$

where  $\mu_i, \mu_j$  are the means and  $\sigma_i, \sigma_j$  are the standard deviations of the row and column

# But how do you choose d?

- This is actually a critical question with **all** the statistical texture methods.
- Are the “texels” tiny, medium, large, all three ...?
- Not really a solved problem.

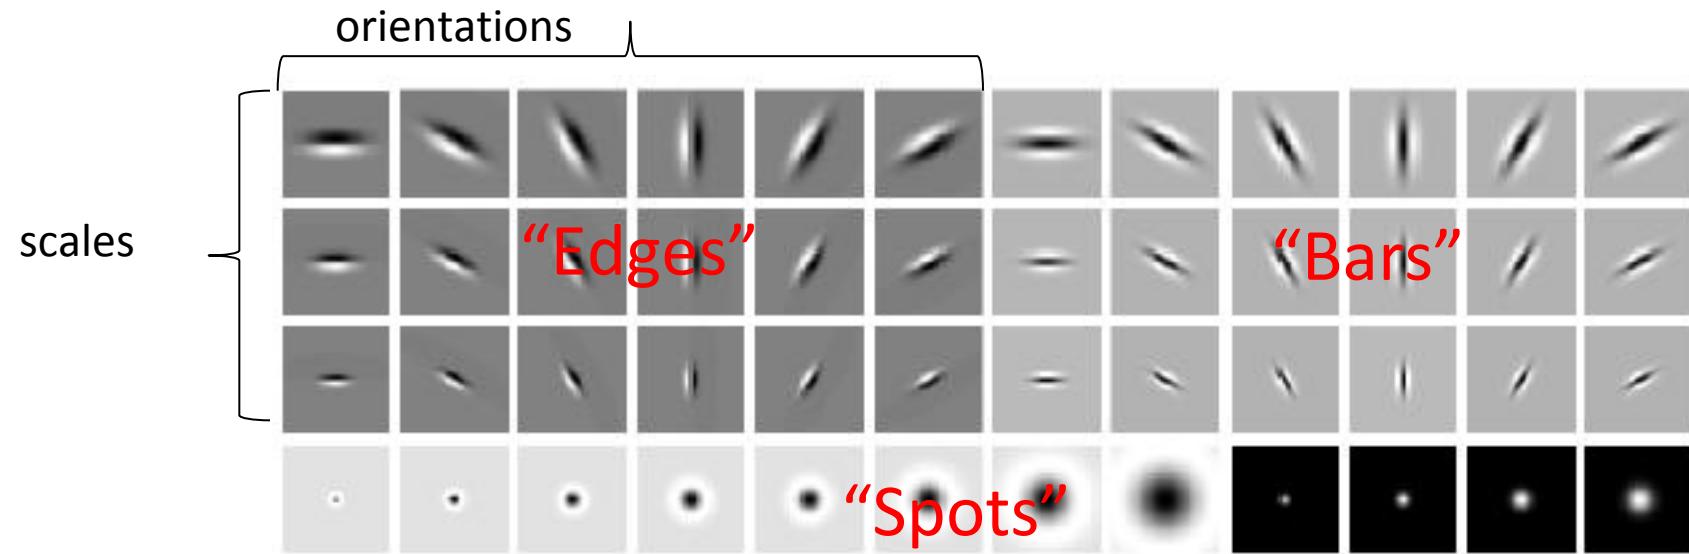
# Alternative approach

Signal-processing-based algorithms use texture filters applied to the image to create filtered images from which texture features are computed.

Textures are made up of repeated local patterns, so:

- Find the patterns
  - Use filters that look like patterns (spots, bars, raw patches...)
  - Consider magnitude of response
- Describe their statistics within each local window, e.g.,
  - Mean, standard deviation
  - Histogram
  - Histogram of “prototypical” feature occurrences

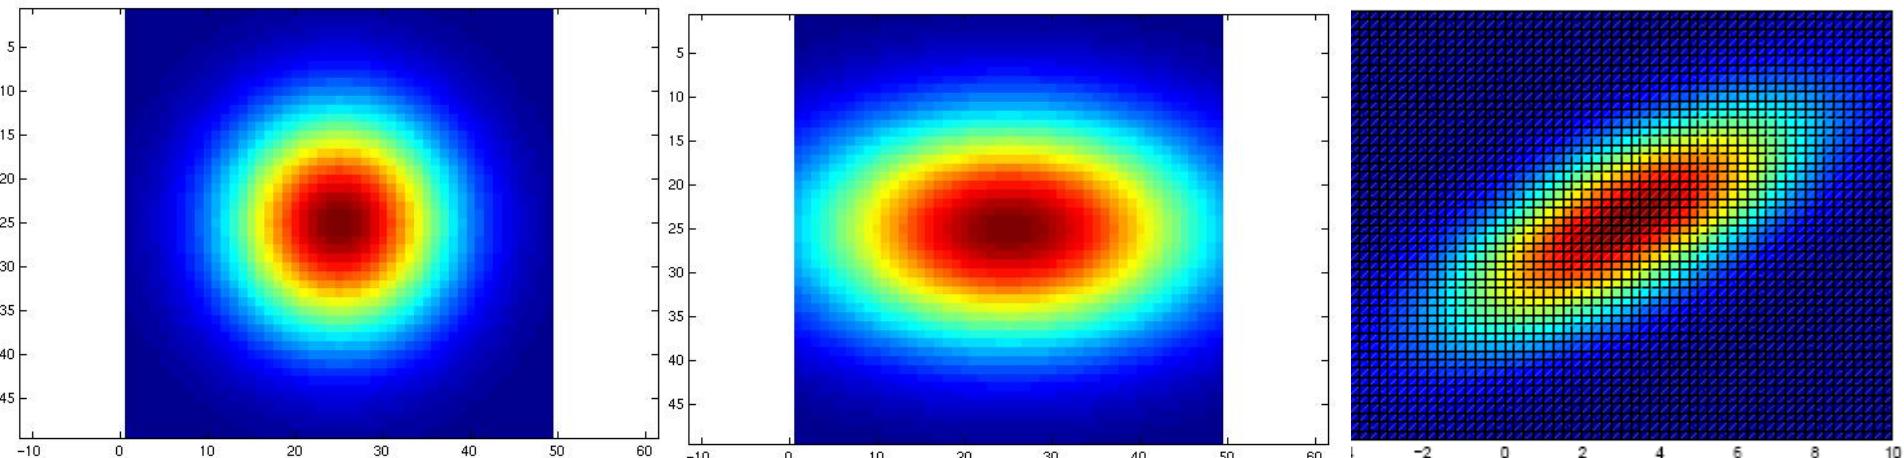
# Filter banks



- What filters to put in the bank?
  - Typically we want a combination of scales and orientations, different types of patterns.

# Multivariate Gaussian

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right).$$



$$\Sigma = \begin{bmatrix} 9 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 16 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 10 & 5 \\ 5 & 5 \end{bmatrix}$$

Slide credit: Kristen  
Grauman

# Filter bank

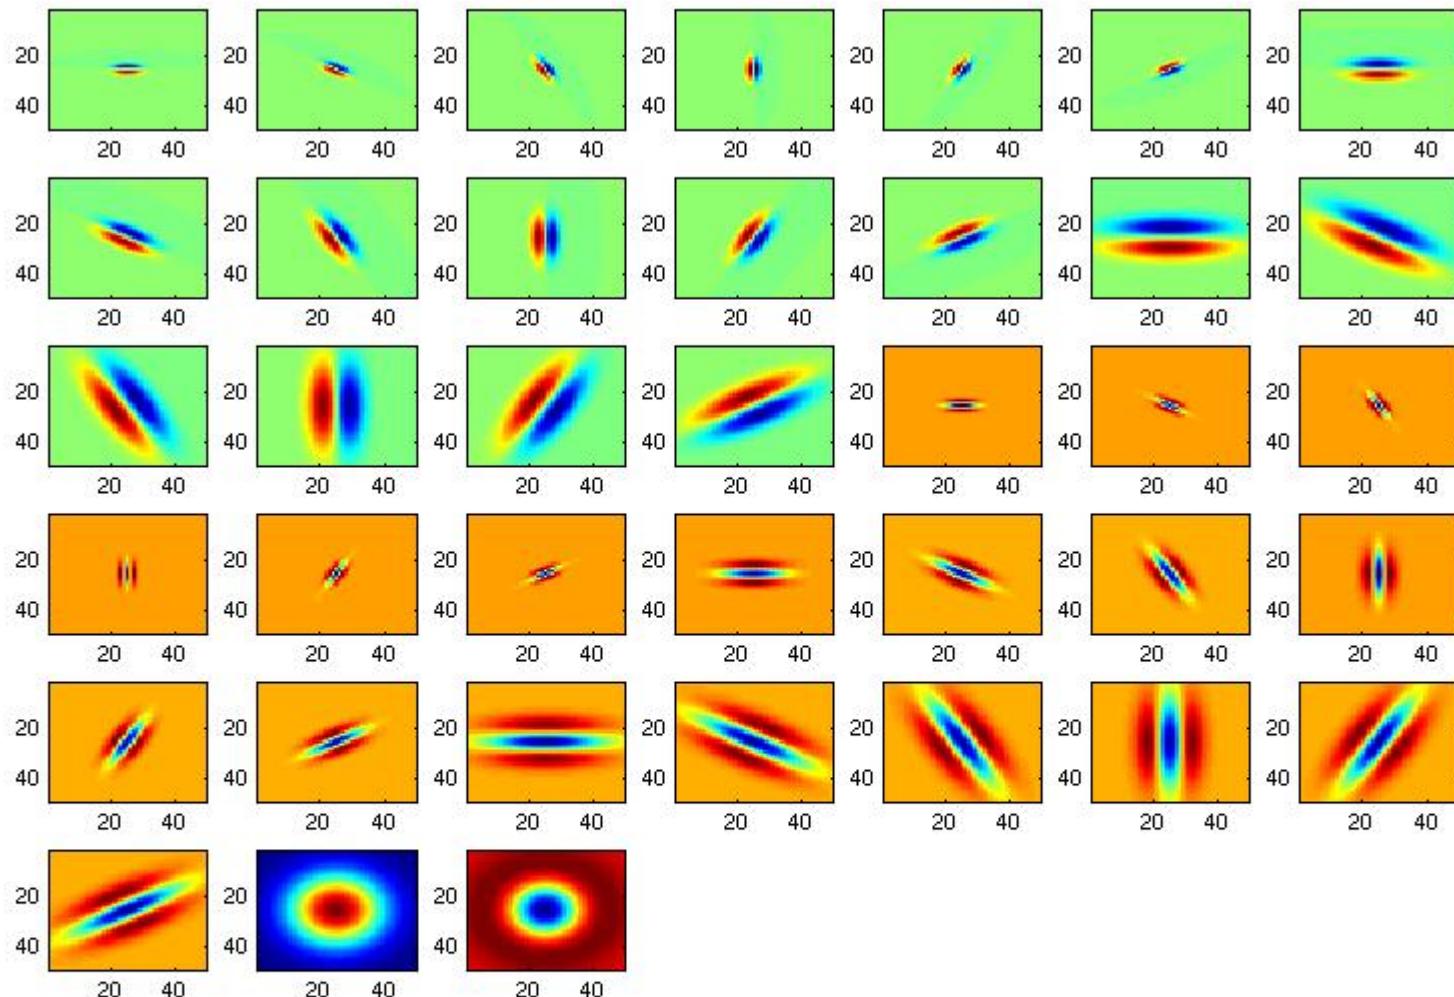
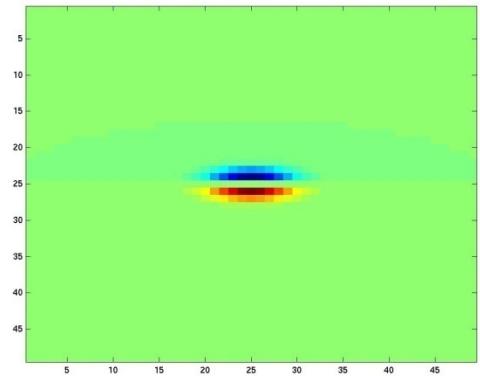
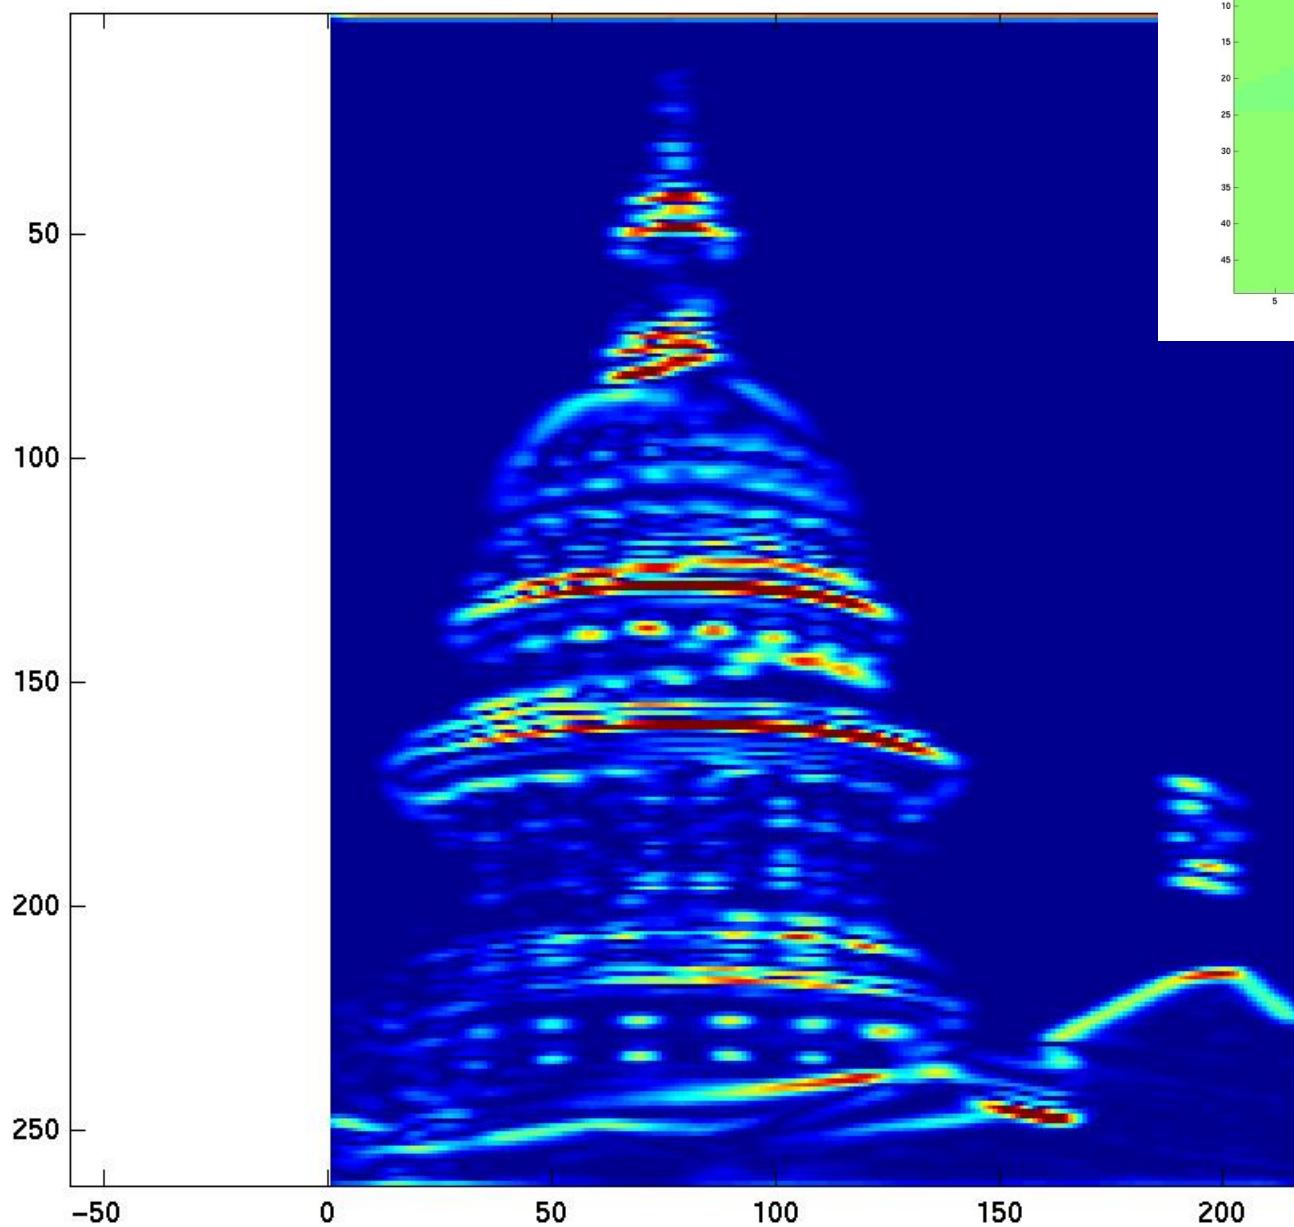


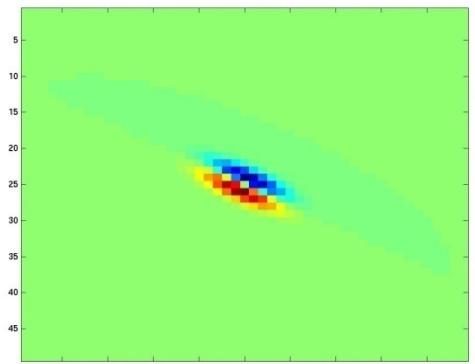
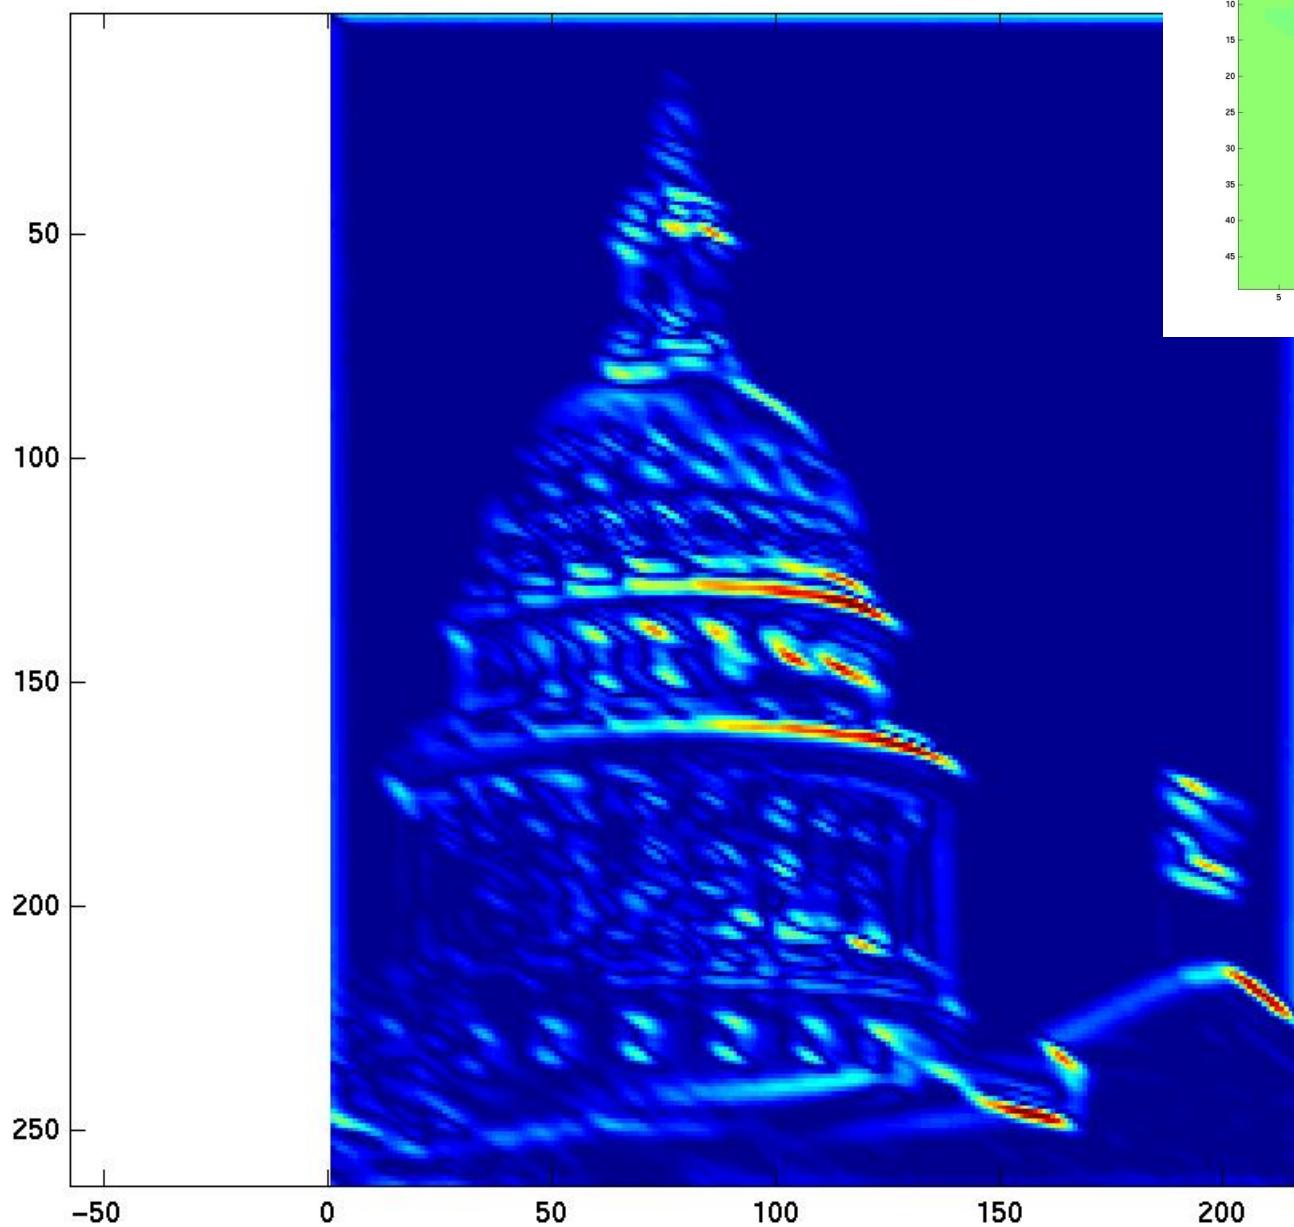
Image from <http://www.texasexplorer.com/austincap2.jpg>



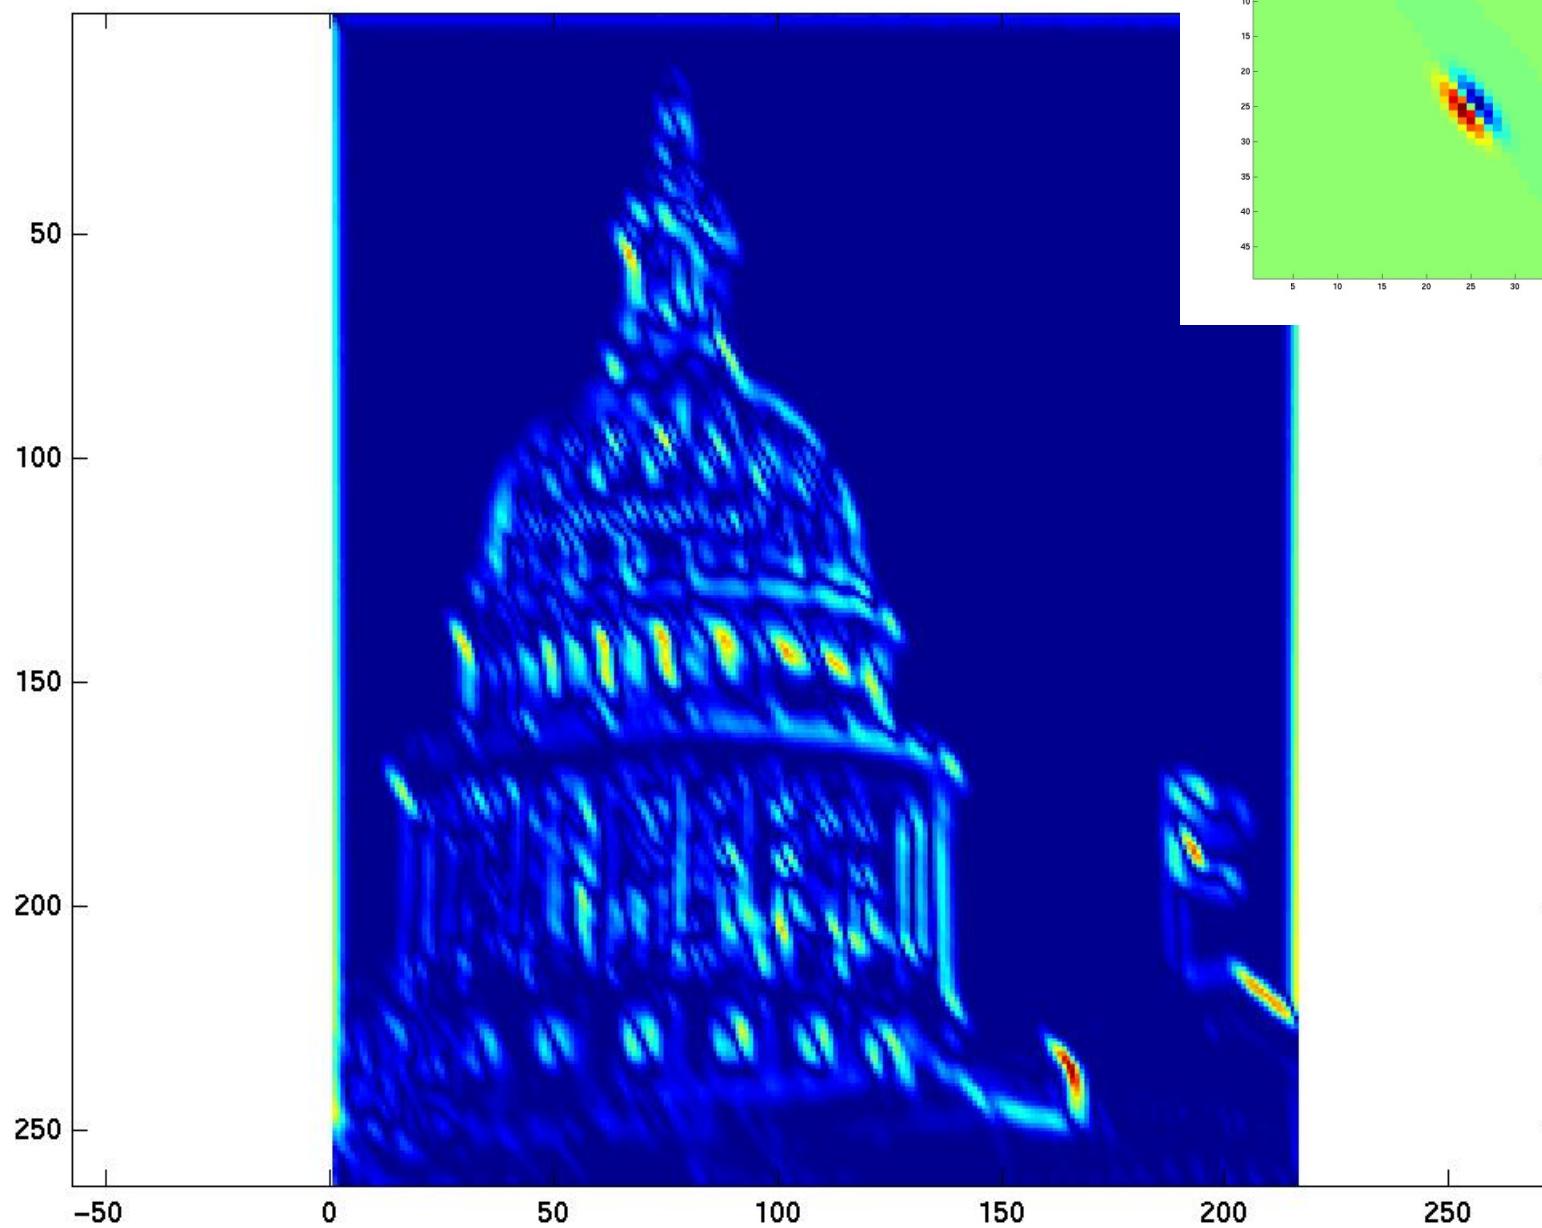
Slide credit: Kristen Grauman



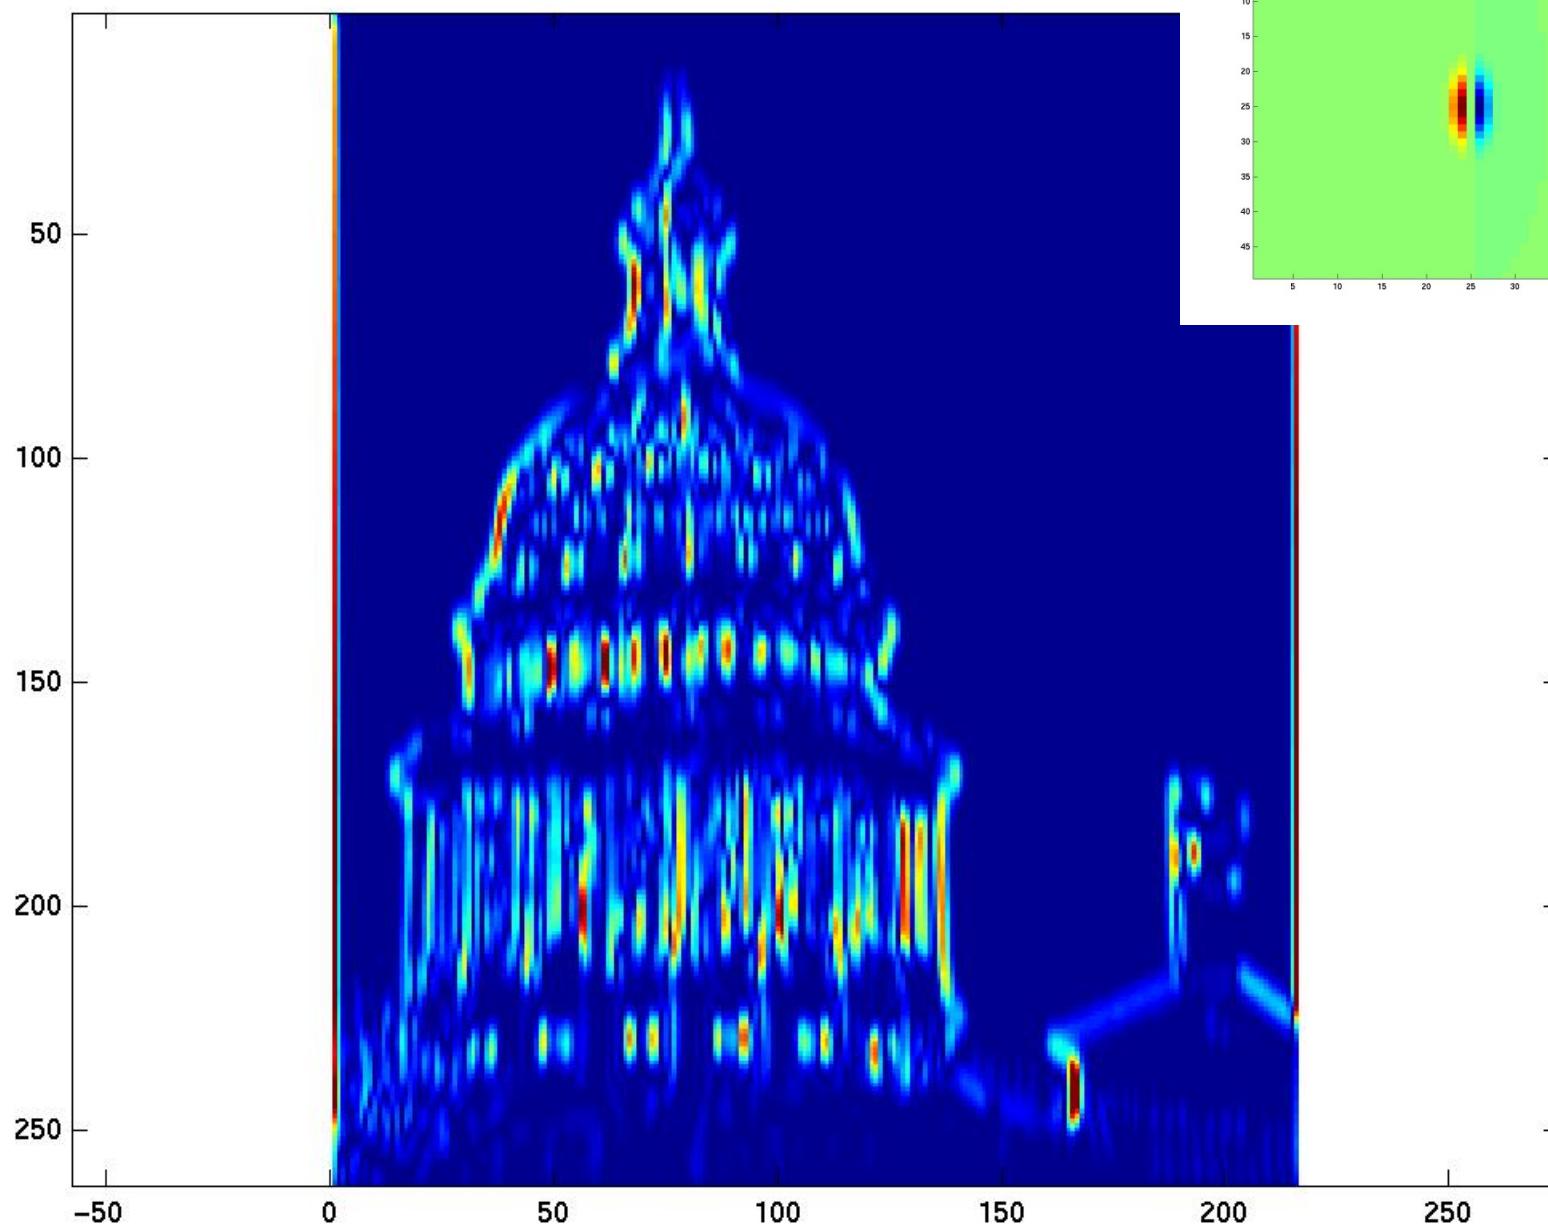
Slide credit: Kristen  
Grauman



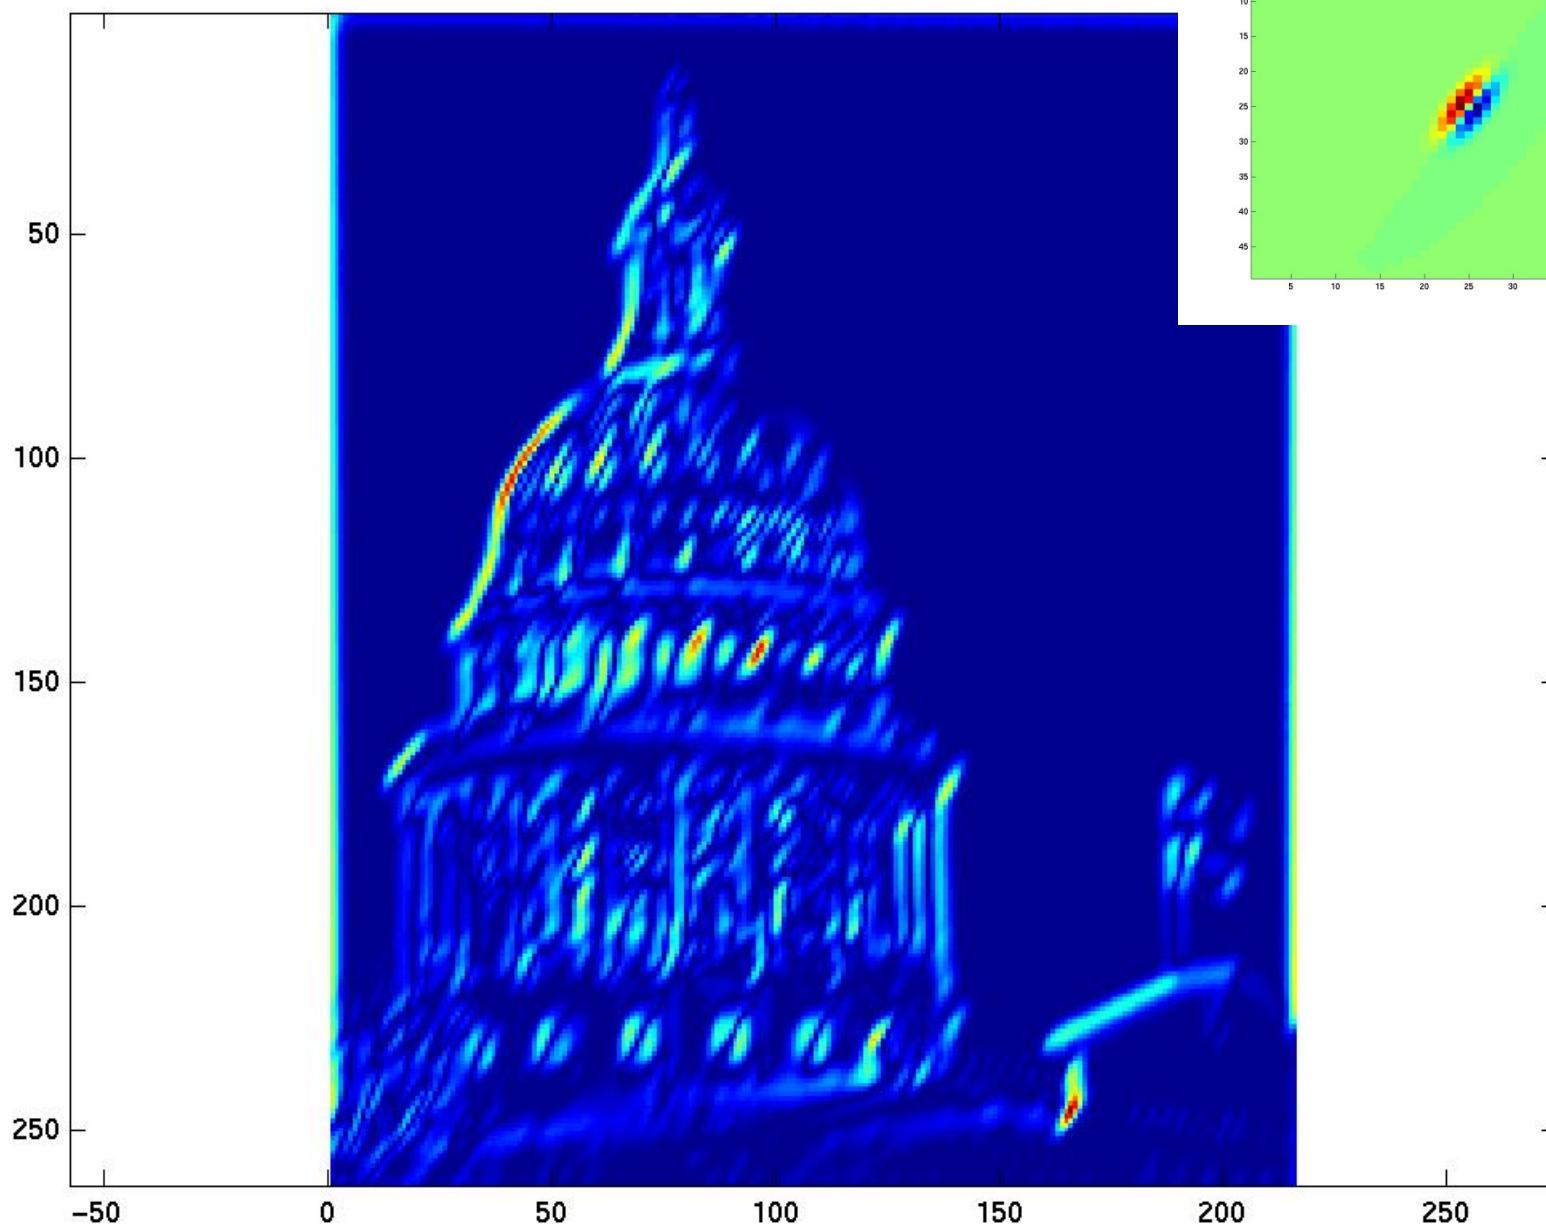
Slide credit: Kristen  
Grauman



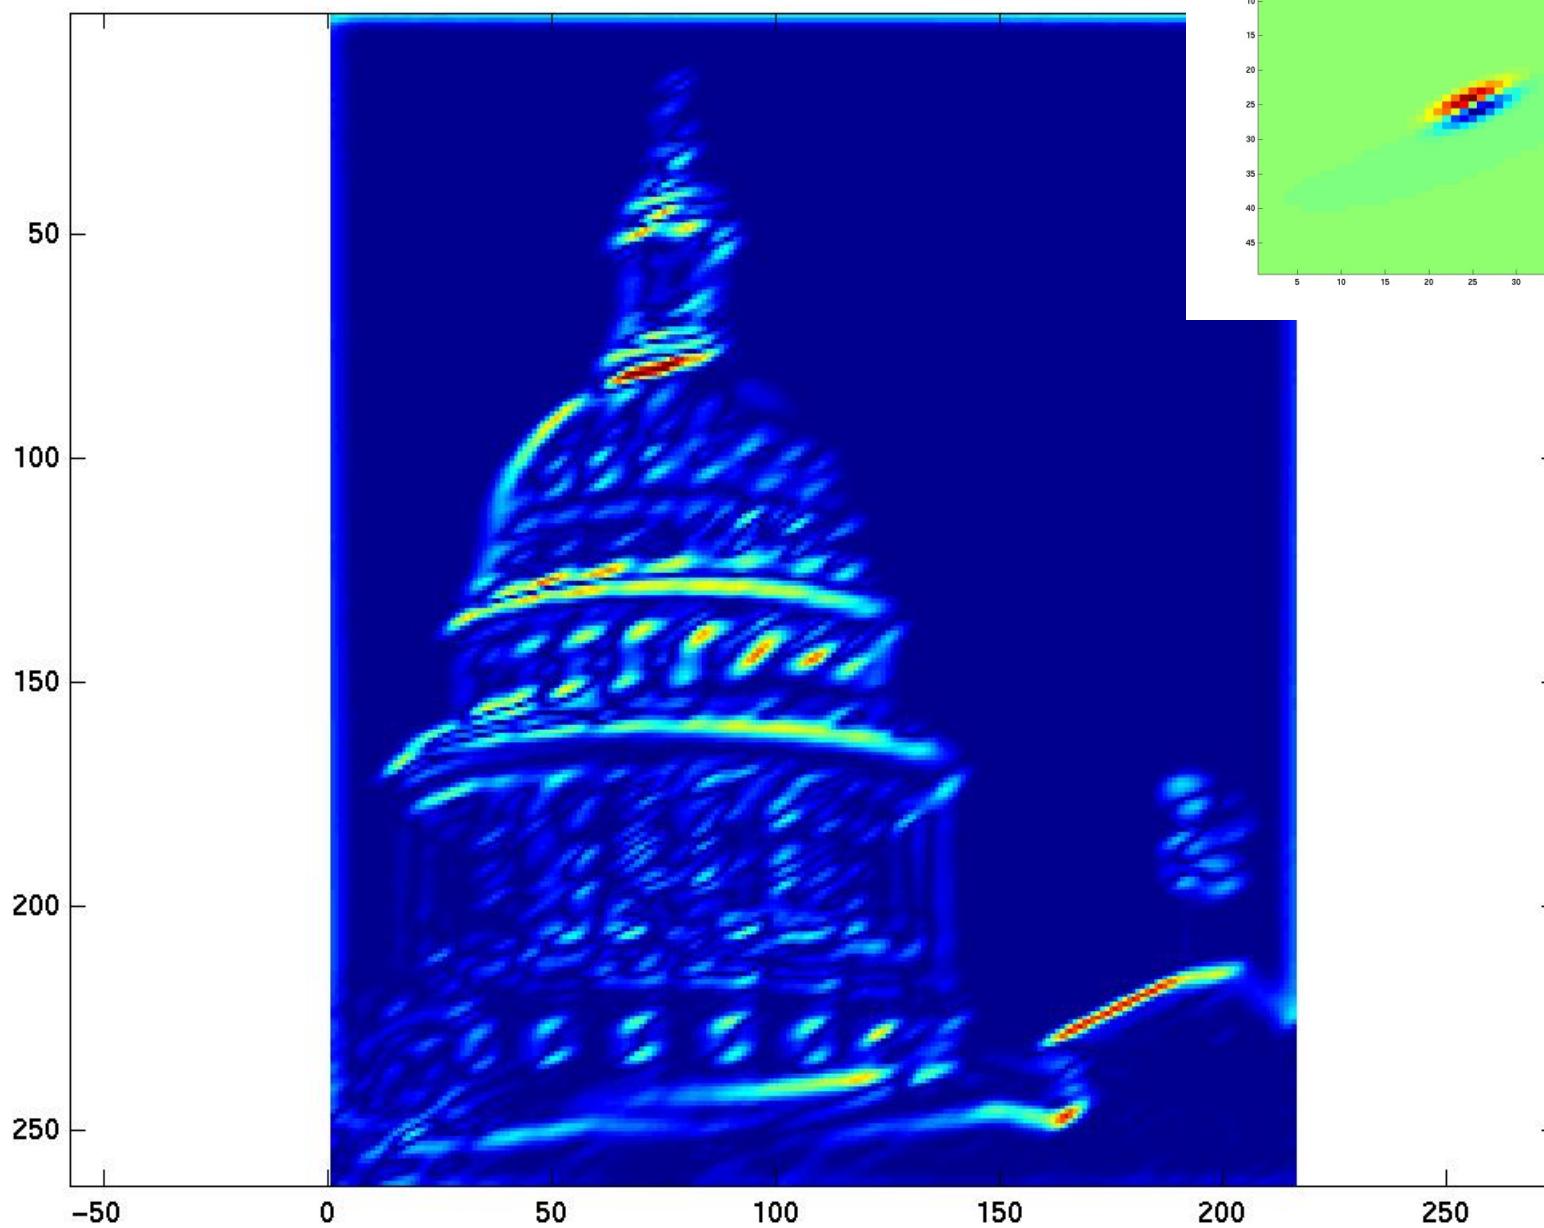
Slide credit: Kristen  
Grauman



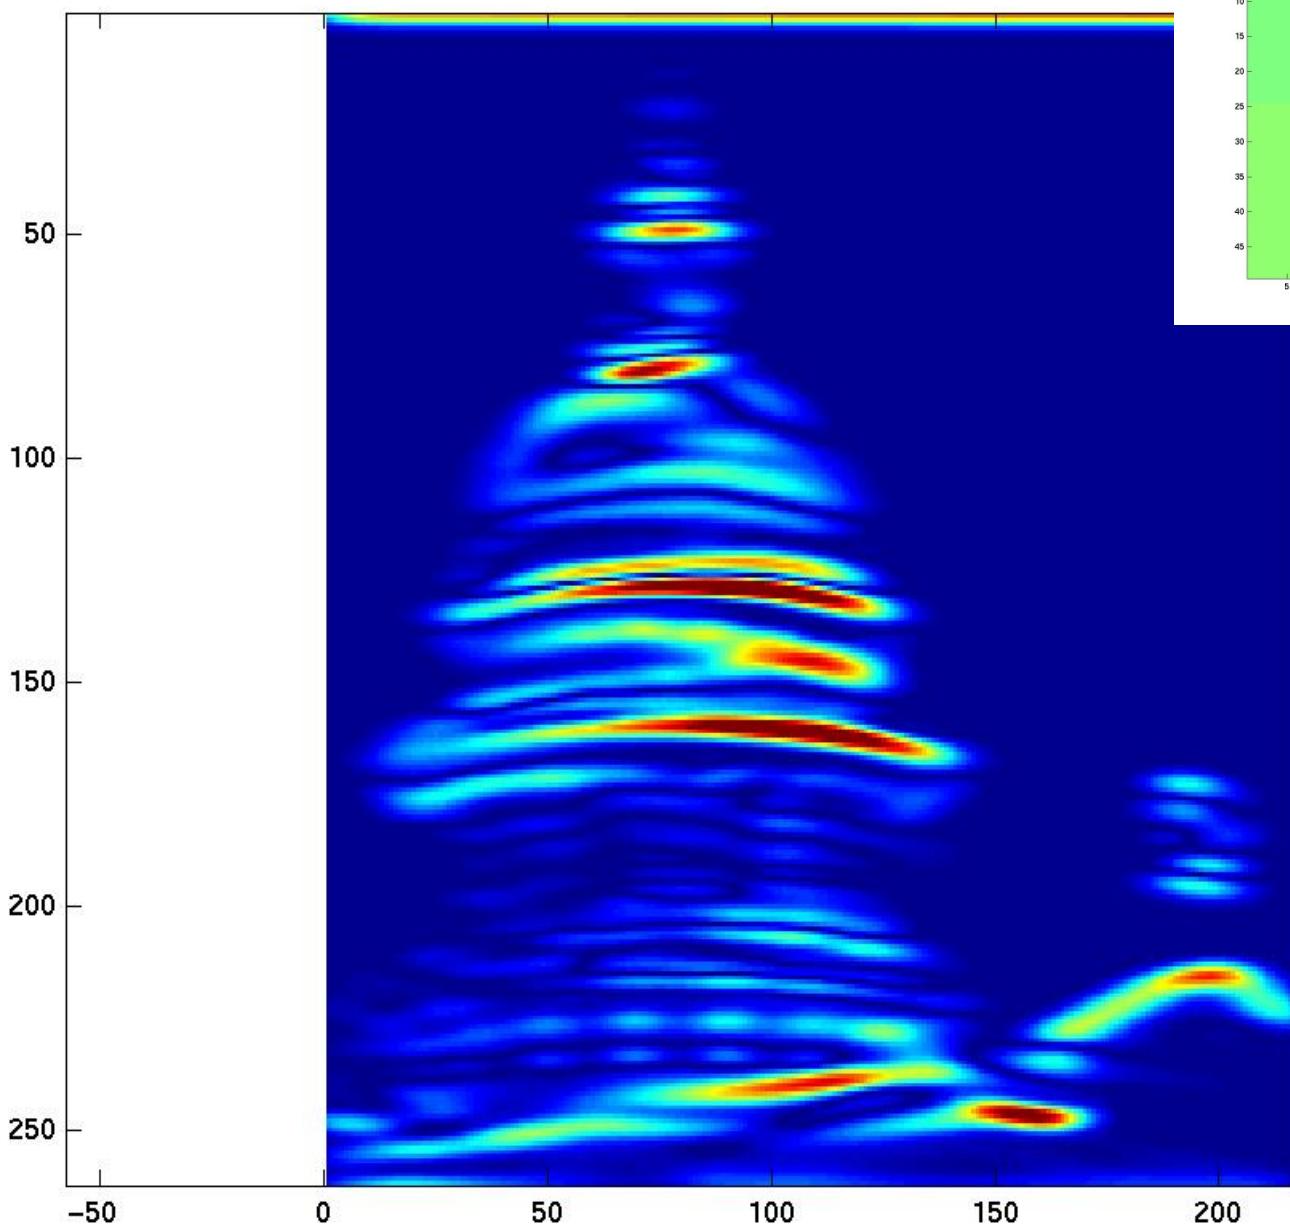
Slide credit: Kristen  
Grauman



Slide credit: Kristen  
Grauman

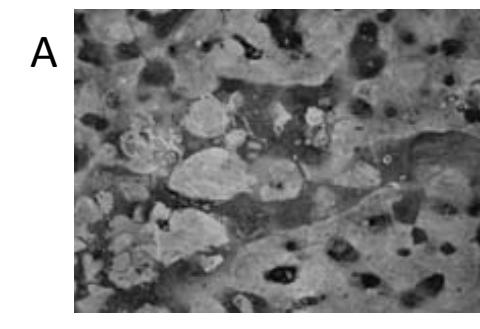
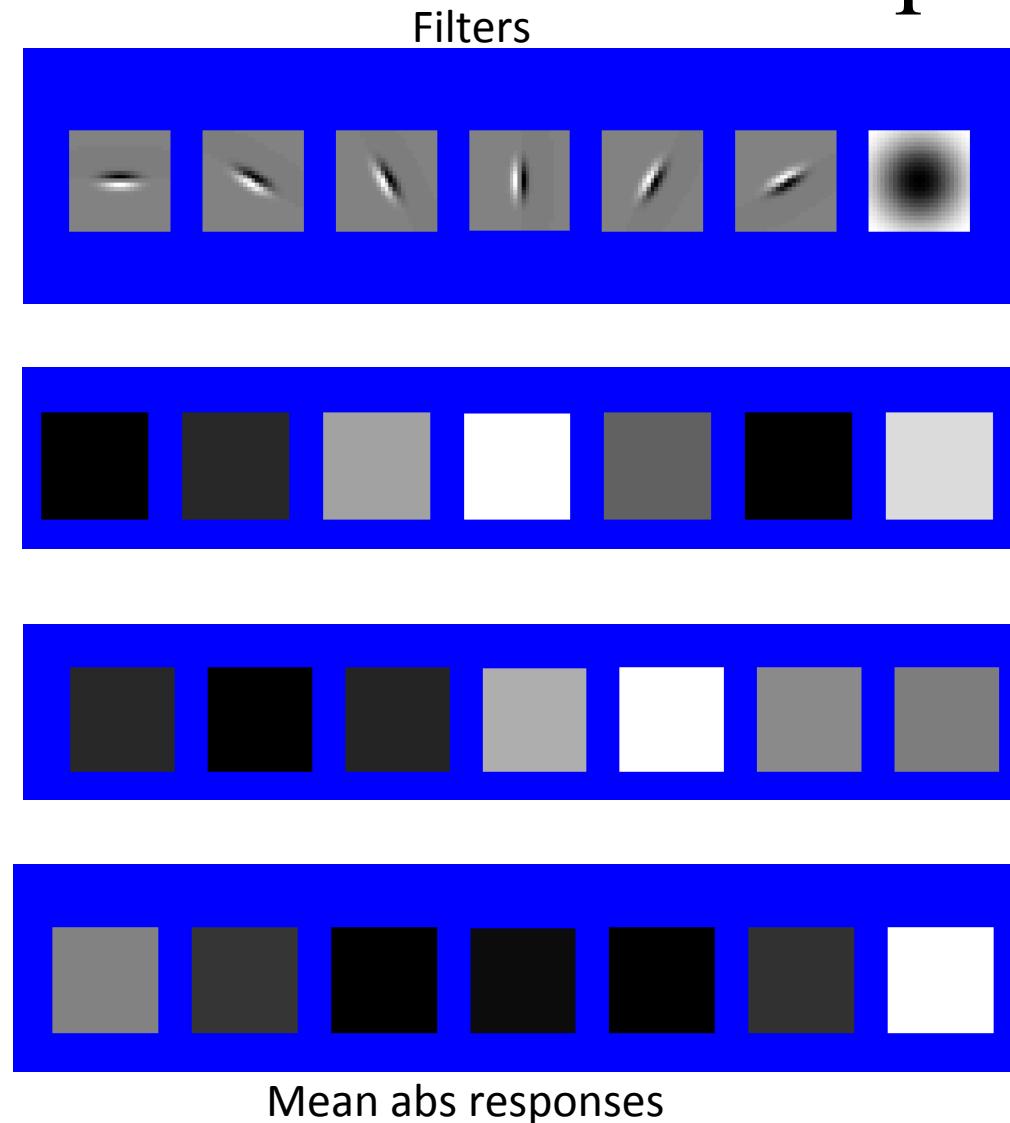


Slide credit: Kristen  
Grauman

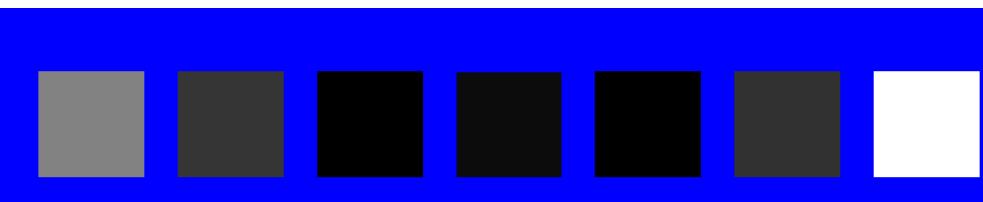
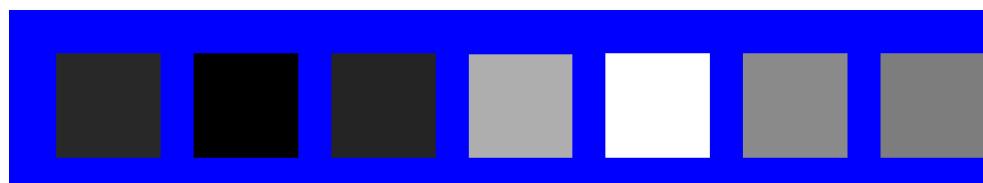
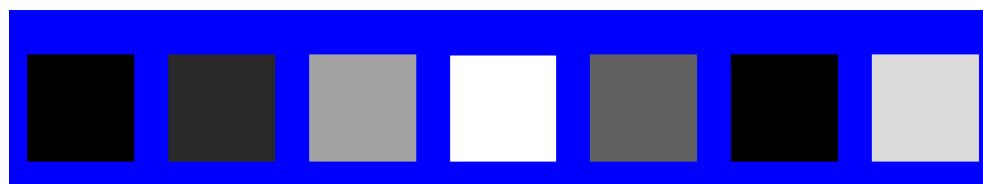
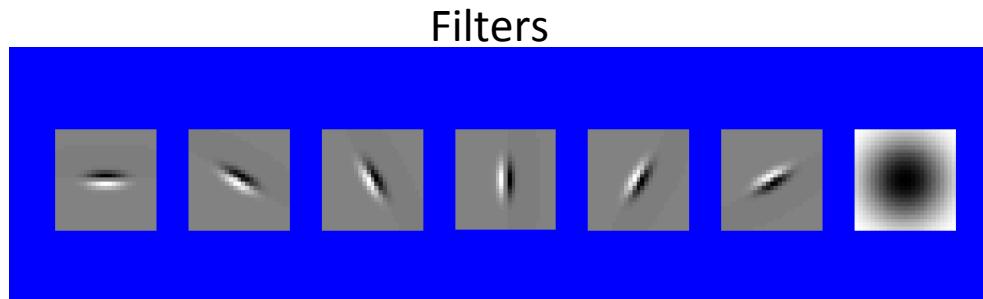
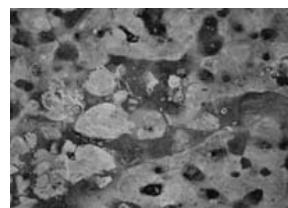


Slide credit: Kristen  
Grauman

# You try: Can you match the texture to the response?

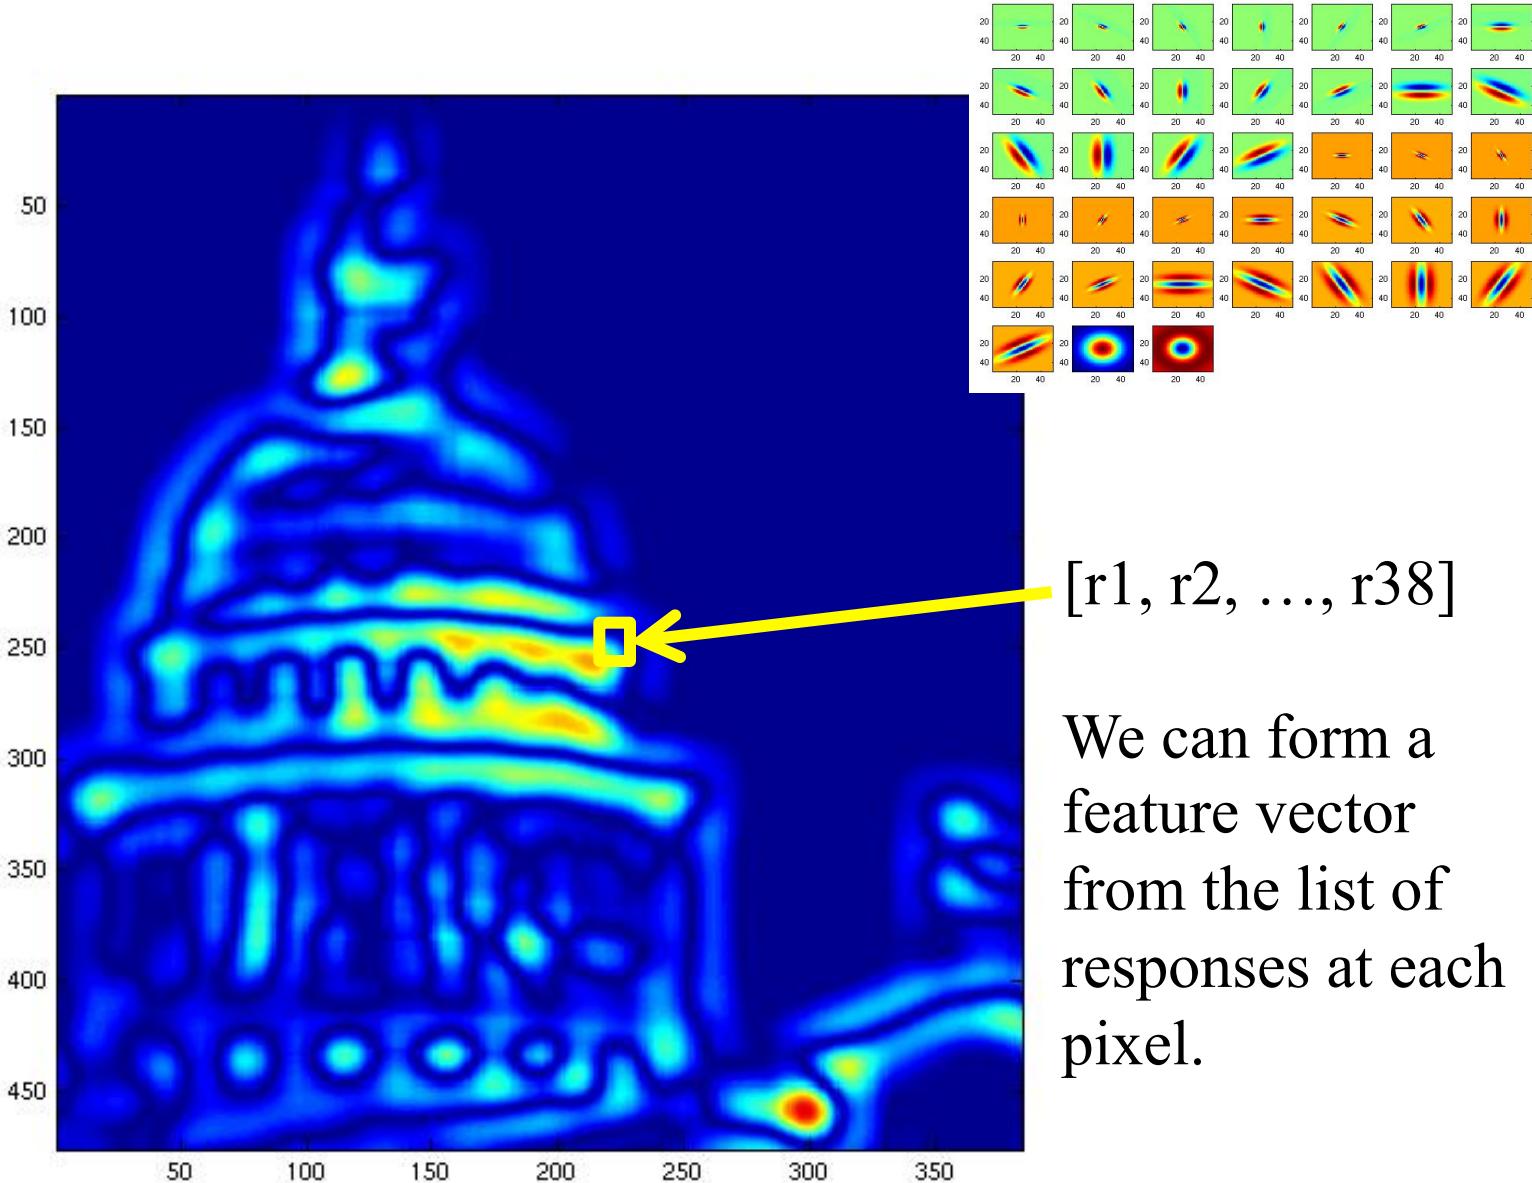


# Representing texture by mean abs response



Mean abs responses

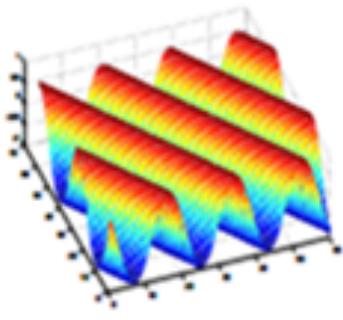
Slide credit Derek Hoiem



$[r_1, r_2, \dots, r_{38}]$

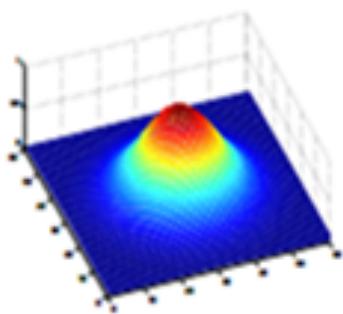
We can form a feature vector from the list of responses at each pixel.

# Gabor Filters



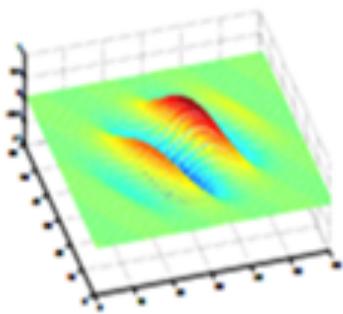
A Sinusoid oriented  $30^{\circ}$  with X-axis

(a)



A 2-D Gaussian

(b)

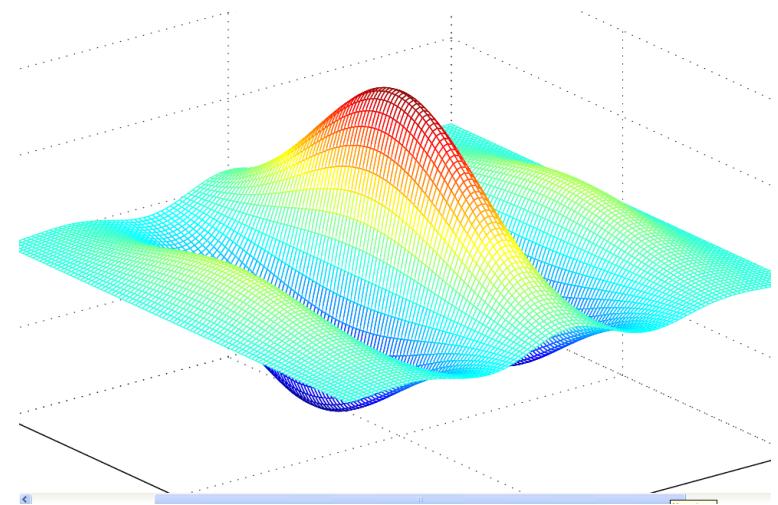
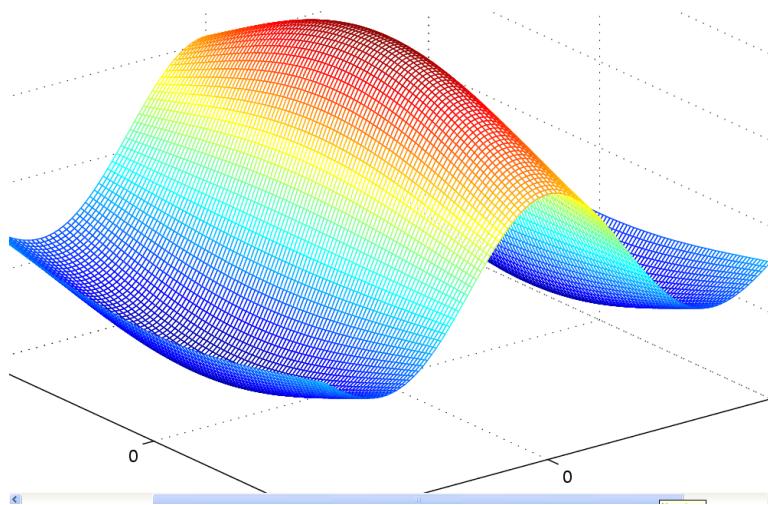


The corresponding 2-D Gabor filter

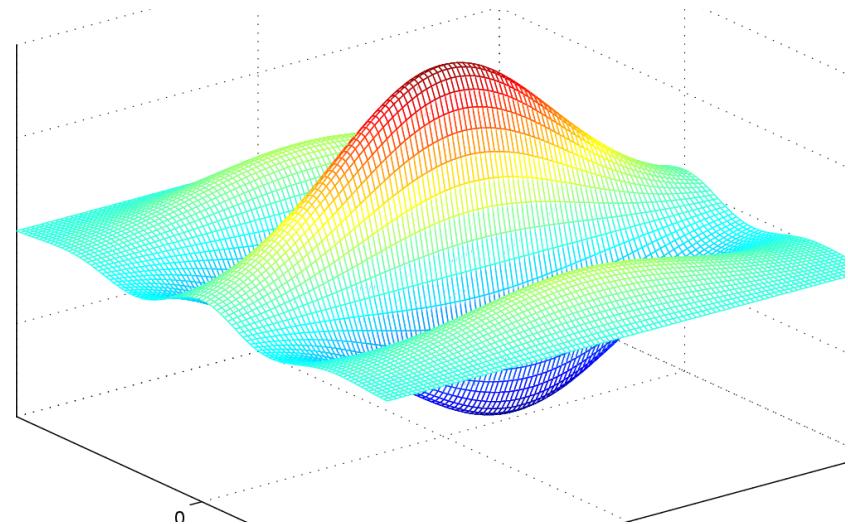
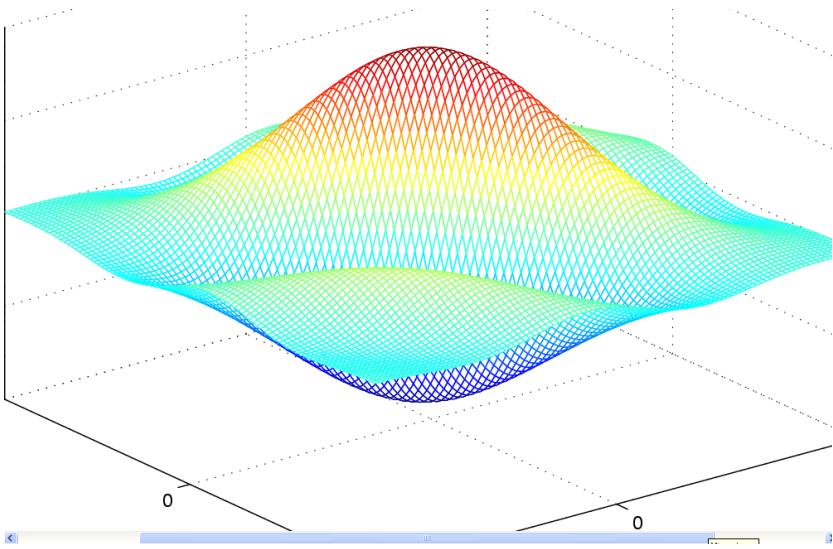
(c)

# Gabor Filters

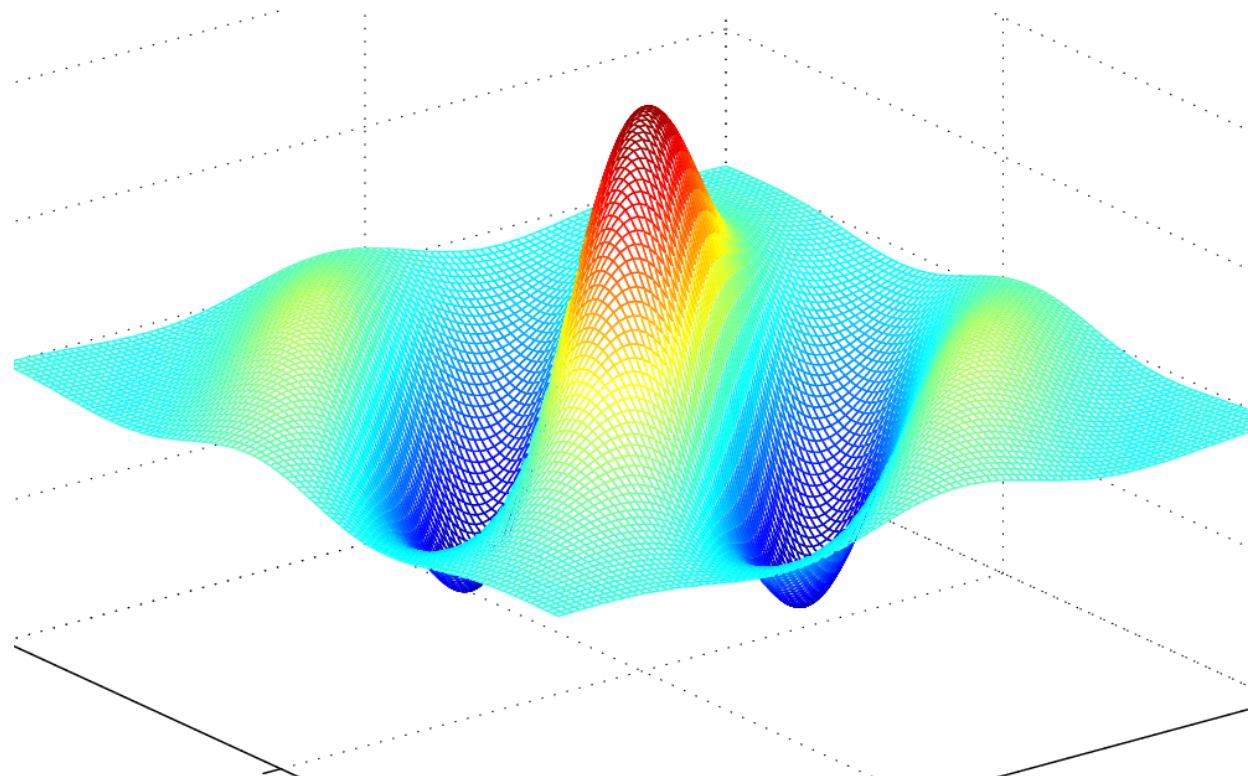
- Wavelets at different frequencies and different orientations



# Gabor Filters



# Gabor Filters



Gabor filter values visualization by Juergen Mueller.

variable naming conventions as used by [http://en.wikipedia.org/wiki/Gabor\\_filter](http://en.wikipedia.org/wiki/Gabor_filter).

Red: positive values. Blue: negative values

1. Wavelength is changed.

1. wavelength (lambda) : 5.00

2. orientation (theta) : 0.79

3. gaussvar (sigma) : 20.00

4. phaseoffset (phi) : 0.00

5. aspectratio (gamma) : 0.50

min gabor val : -0.91

max gabor val : 1.00

